


CS2040S – Data Structures and Algorithms

Lecture 8* – Bloom Filter

chongket@comp.nus.edu.sg



Bloom Filter – A Probabilistic Map



Motivation for using Bloom Filter

- When a hash table (specifically a **hash set**) is too large to fit into RAM, data will have to be read/written to and from hard disk which is much slower than RAM
- A possible solution is to use Bloom Filter instead of a classical hash table, as it is able to “compress” the amount of memory required to store the data

Introduction to Bloom Filter (1)

- Most important characteristic of a Bloom Filter is that retrieval of a key in the Bloom Filter is probabilistic
 - “No” means the key is definitely not in the filter
 - “Yes” has a certain probability that the key is actually not in the Bloom Filter (a false positive)

Introduction to Bloom Filter (2)

- Data Structure used – A bit array (in Java there is no bit primitive data type so have to use the **BitSet** class) which is initialized to 0
- Basic Operations
 - Insertion
 - Retrieval
- Can't do removal in a Bloom Filter !

Bloom Filter - Insert operation (1)

- Inserting an integer key k into the bloom filter requires setting up to M bits of the bit array to 1. To do this,
 - use M independent and uniformly distributed hash functions $H_1(k)$ to $H_M(k)$
 - Each hash function returns an (hopefully different) index of the bit array which is then set to 1
- M is usually bounded by a small value so can be considered a constant

Bloom Filter - Insert operation (2)

3 hash functions:

$$H_1(\text{key}) = (\text{key} \% 17) \% 10$$

$$H_2(\text{key}) = (\text{key} \% 29) \% 10$$

$$H_3(\text{key}) = (\text{key} \% 37) \% 10$$



Bit array of size 10

{123}

Inserting 123:

$$H_1(123) = 4$$

$$H_2(123) = 7$$

$$H_3(123) = 2$$



Bit array of size 10

{123, 306}

Inserting 306:

$$H_1(306) = 0$$

$$H_2(306) = 6$$

$$H_3(306) = 0$$



Bit array of size 10

Bloom Filter - Insert operation (3)

Inserting 7125:

$$H_1(7125) = 2$$

$$H_2(7125) = 0$$

$$H_3(7125) = 1$$

{123, 306, 7125}



Bit array of size 10

Overlap with bit representation for 123, 306 with 7125

Bloom Filter - Retrieval operation (1)

- To retrieve a key again use the M hash functions to get M indices in the filter
 - If all M indices are set to 1 then the key is possibly in the filter
 - If at least one index is 0 then the key is definitely not in the filter

Bloom Filter - Retrieval operation (2)

{123,306,7125}

Given



Bit array of size 10

Retrieving 7125:

$$H_1(7125) = 2$$

$$H_2(7125) = 0$$

$$H_3(7125) = 1$$

{123,306,7125}



Bit array of size 10



All 1's so return 7125 is in the Bloom Filter

Bloom Filter - Retrieval operation (3)

Retrieving 513:

$$H_1(513) = 3$$

$$H_2(513) = 0$$

$$H_3(513) = 2$$

{123,306,7125}



Bit array of size 10

Retrieving 74:

$$H_1(74) = 6$$

$$H_2(74) = 6$$

$$H_3(74) = 0$$

{123,306,7125}



Bit array of size 10

Considerations when designing a good Bloom Filter (1)

- In general, the bigger the bloom filter and the more hash functions used, the less likely the bits set for any key will coincide with any other key
- However this defeats the purpose of the bloom filter to “compress” the memory requirements
- Need to find a good compromise

Considerations when designing a good Bloom Filter (2)

- 4 variables for a bloom filter
 - N : the number of keys to be inserted
 - S : the size of the bit array
 - M : the number of hash functions used
 - P : the probability of a false positive
- If we fix the value for N or upper bound for N , and the required P , then S and M can be computed as follows

Considerations when designing a good Bloom Filter (3)

$$S = \text{Ceil} \left(-\frac{N * \ln P}{(\ln 2)^2} \right)$$

$$M = \text{round} \left(\frac{S}{N} * \ln 2 \right)$$

Compression Ratio :

Assuming keys = 32 bits or 4 byte integers

At N = 10M and P = 0.01

Bit array size = 100Mbit = 12.5 Mbyte

If use hashtable need at least 40Mbytes.

Compression ratio = 40/12.5 = **3.2 times**




Image Credit: [Abishek Bhat's article about bloom filter](#)

Considerations when designing a good Bloom Filter (4)

- In general the number of bits required to store each key in a bloom filter with false positive rate P is

$$\text{Number of bits per key} = 0.26 \ln \frac{1}{P}$$

- The number of bits required per key is independent of the size of the key itself, thus the greater the size of a key the greater the compression ratio

Bloom Filter Conclusion

- Pros
 - Useful when number of keys too large for classical hash table to fit into RAM
 - Useful when we can accept a small amount of errors (false positives) during retrieval
 - Use M bits to represent a key regardless of the size of the key (especially useful when key is a long string)
- Cons
 - Larger overhead to insert/retrieve a key since need to compute M hash functions instead of 1 hash function
 - Not useful for problems where no error is tolerated
 - Not cache-friendly since the M bits for a key may not be close together thus cannot fit into cache