

CS2040S

Data Structures and Algorithms

(e-learning edition)

Skip Lists!

Dictionary Interface


A collection of (key, value) pairs:

interface IDictionary

<code>void</code>	<code>insert(Key k, Value v)</code>	<i>insert (k,v) into table</i>
<code>Value</code>	<code>search(Key k)</code>	<i>get value paired with k</i>
<code>Key</code>	<code>successor(Key k)</code>	<i>find next key > k</i>
<code>Key</code>	<code>predecessor(Key k)</code>	<i>find next key < k</i>
<code>void</code>	<code>delete(Key k)</code>	<i>remove key k (and value)</i>
<code>boolean</code>	<code>contains(Key k)</code>	<i>is there a value for k?</i>
<code>int</code>	<code>size()</code>	<i>number of (k,v) pairs</i>

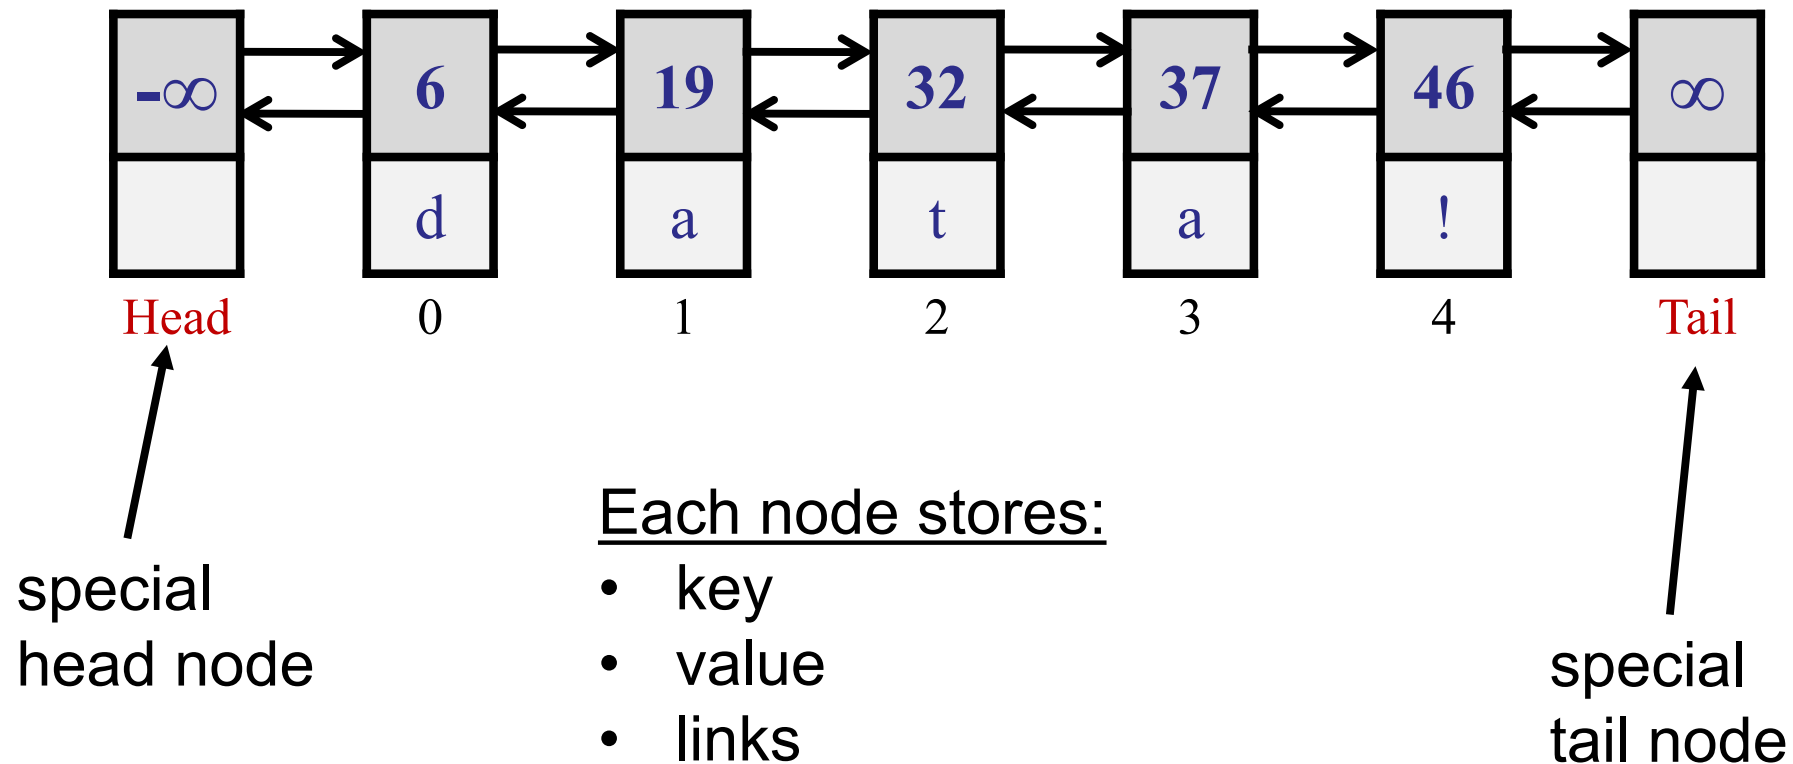
Implementing a dictionary, again...

Store keys in a sorted linked list:




Implementing a dictionary, again...

Store keys in a sorted linked list:



Implementing a dictionary, again...

Store keys in a sorted linked list:




Doubly linked:

Each node has a link to its predecessor and successor in the list.

Implementing a dictionary, again...

Store keys in a sorted linked list:




Time:

- Search: $O(n)$
- Insert: $O(n)$


Implementing a dictionary, again...

Why doesn't binary search work??



Implementing a dictionary, again...

Store keys in a sorted linked list:




Time:

- Search: $O(n)$
- Insert: $O(n)$

What if...

What if we use two lists?

- Express train
- Local train




`search(37)` takes only 3 steps!

What if...

Calculation:

If the “express” list skips 5 elements per “stop”, then search takes at most: $n/5 + 5$ steps




With two lists, how many elements should the express list skip per hop?

1. $O(1)$
2. $\log(n)$
- ✓ 3. \sqrt{n}
4. n/\sqrt{n}
5. $n/\log(n)$
6. Something else.

What if...

If the “express” list skips \sqrt{n} elements per “stop”, then search takes at most:

$$\frac{n}{\sqrt{n}} + \sqrt{n} = 2\sqrt{n} = O(\sqrt{n})$$



Why stop at two?

Add more lists:

– Two lists: $2\sqrt{n}$

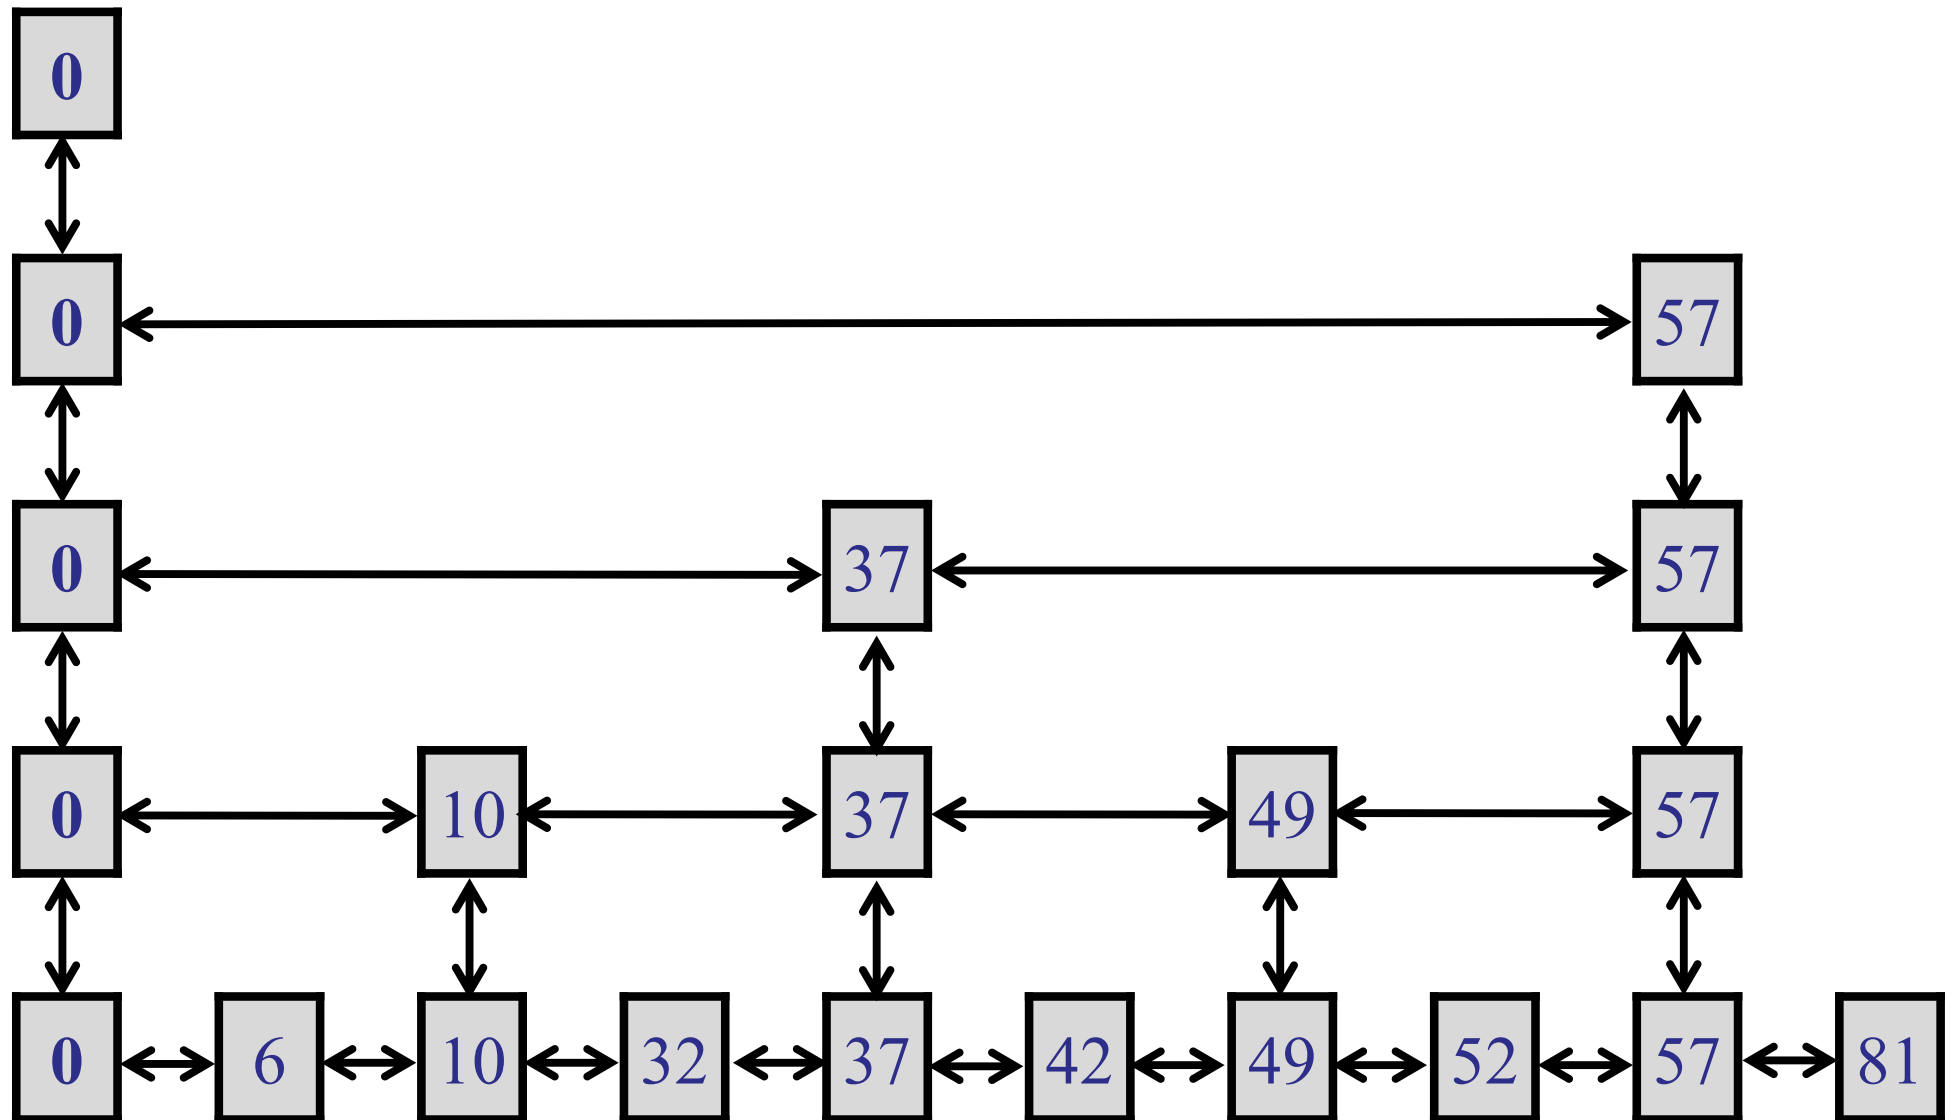
– Three lists: $3\sqrt[3]{n}$

...

– k lists: $k\sqrt[k]{n}$

– $\log(n)$ lists: $\log(n)\sqrt[\log(n)]{n} = \log(n)n^{1/\log(n)}$
 $= 2\log(n)$


Another way to think about it...



How many levels?

1. $O(1)$
- ✓ 2. $\log(n)$
3. $2\log(n)$
4. $\log^2(n)$
5. \sqrt{n}
6. None of the above.

Another way to think about it...



SkipList Background


Simple randomized, dynamic search structure

- Invented by William Pugh in 1989
- Easy to implement

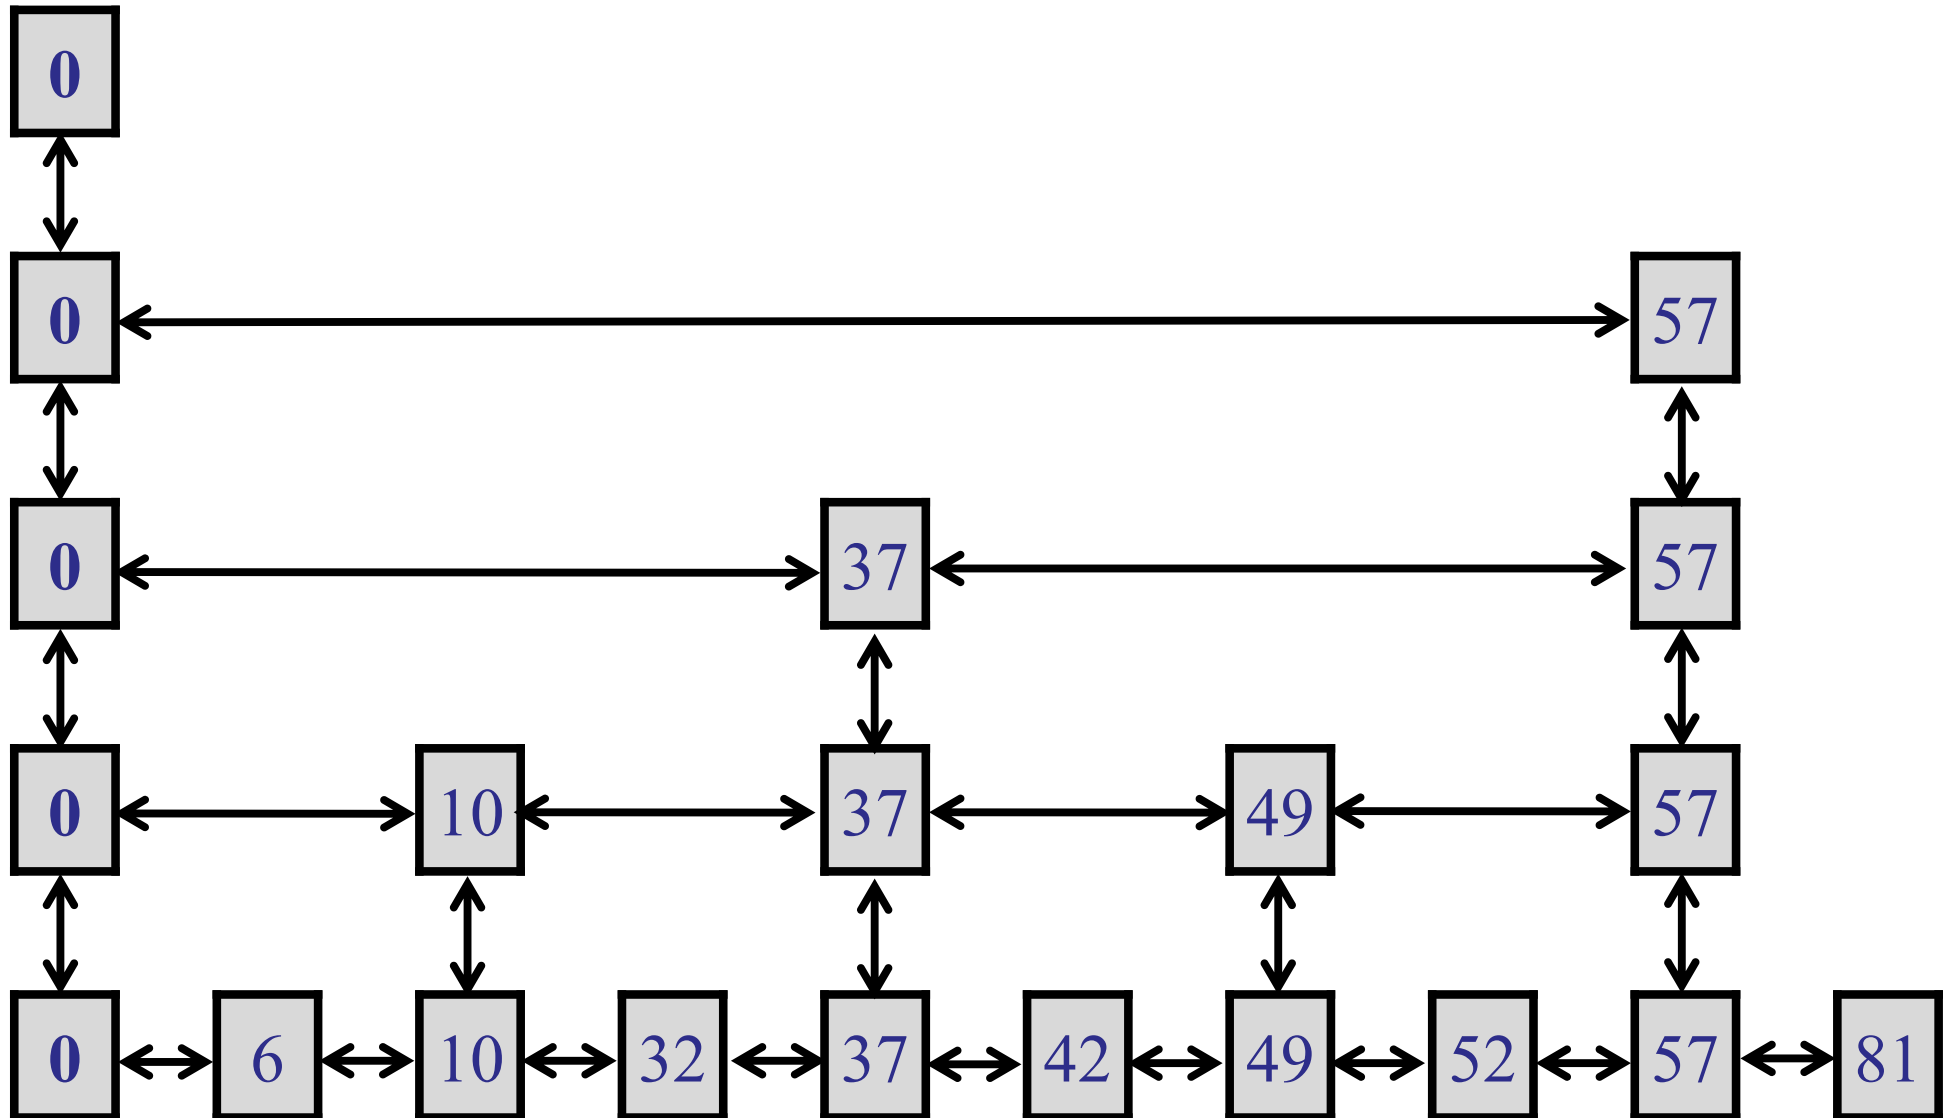
Maintains a set of n elements:

- search: $O(\log n)$ time
 - insert/delete: $O(\log n)$ time
- } *with high probability*

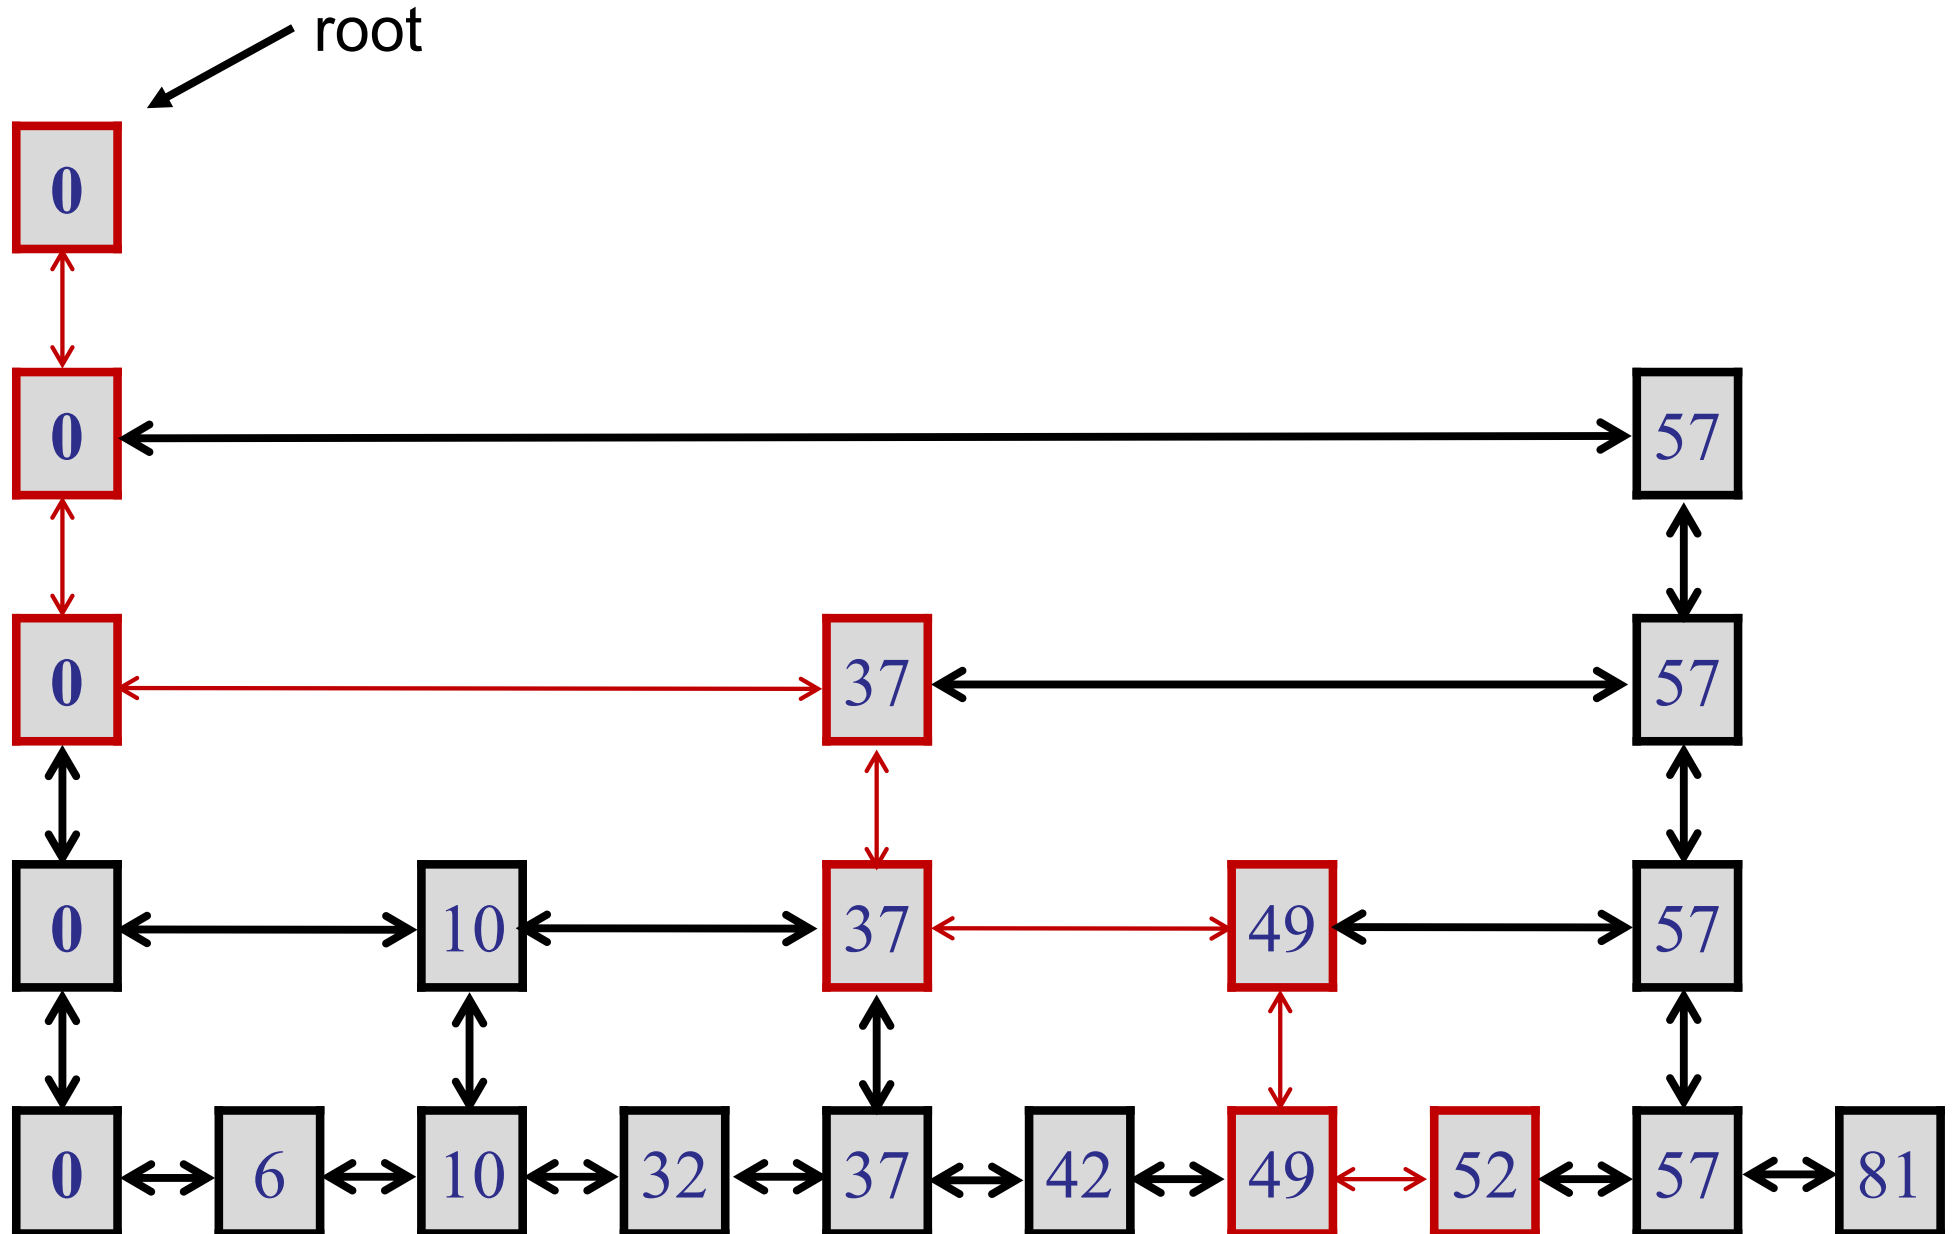
SkipList node



Example: search (52)

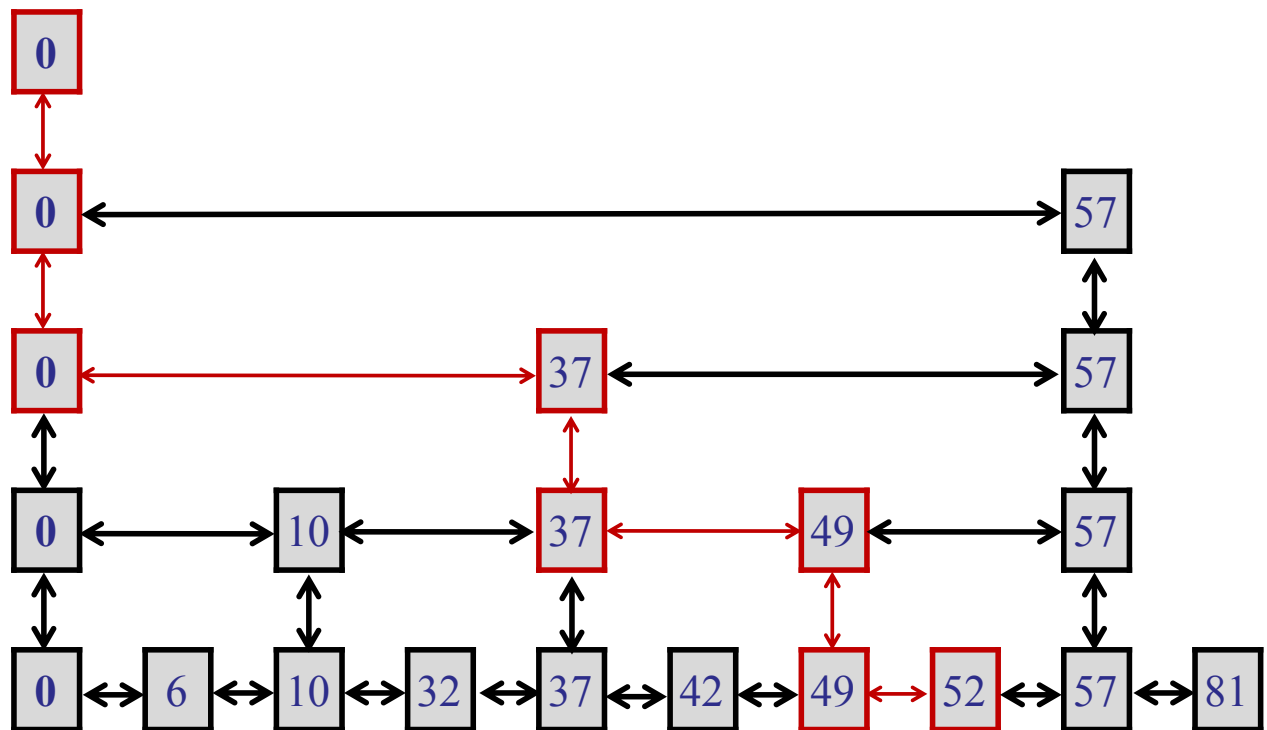


Example: search (52)



search(key)

1. **node = root;**
2. **while (node.key < key) and (node.level > 1):**
3. **while (node.next.key < key):**
4. **node = node.next**
5. **node = node.down**
6. **return node**



Insertions

To insert a new element:

1. Add element to bottom list.

(Invariant: bottom list contains every element.)

2. Add element to some other lists to maintain balance.

Goal: about half of elements at level j get promoted to level $j+1$.

Insertions

Key idea: flip a coin

1. $k = 0$;
2. `while (!done) {`
3. Insert element into level k list.
4. Flip a fair coin:
5. with probability $\frac{1}{2}$: `done = true`;
6. with probability $\frac{1}{2}$: $k = k + 1$;
7. `}`

Insertions

To insert a new element:

1. Add element to bottom list.

(Invariant: bottom list contains every element.)

2. Flip coins to decide how many levels to promote.

– On average: **Level 0:** n

Level 1: $n/2$

Level 2: $n/4$

...

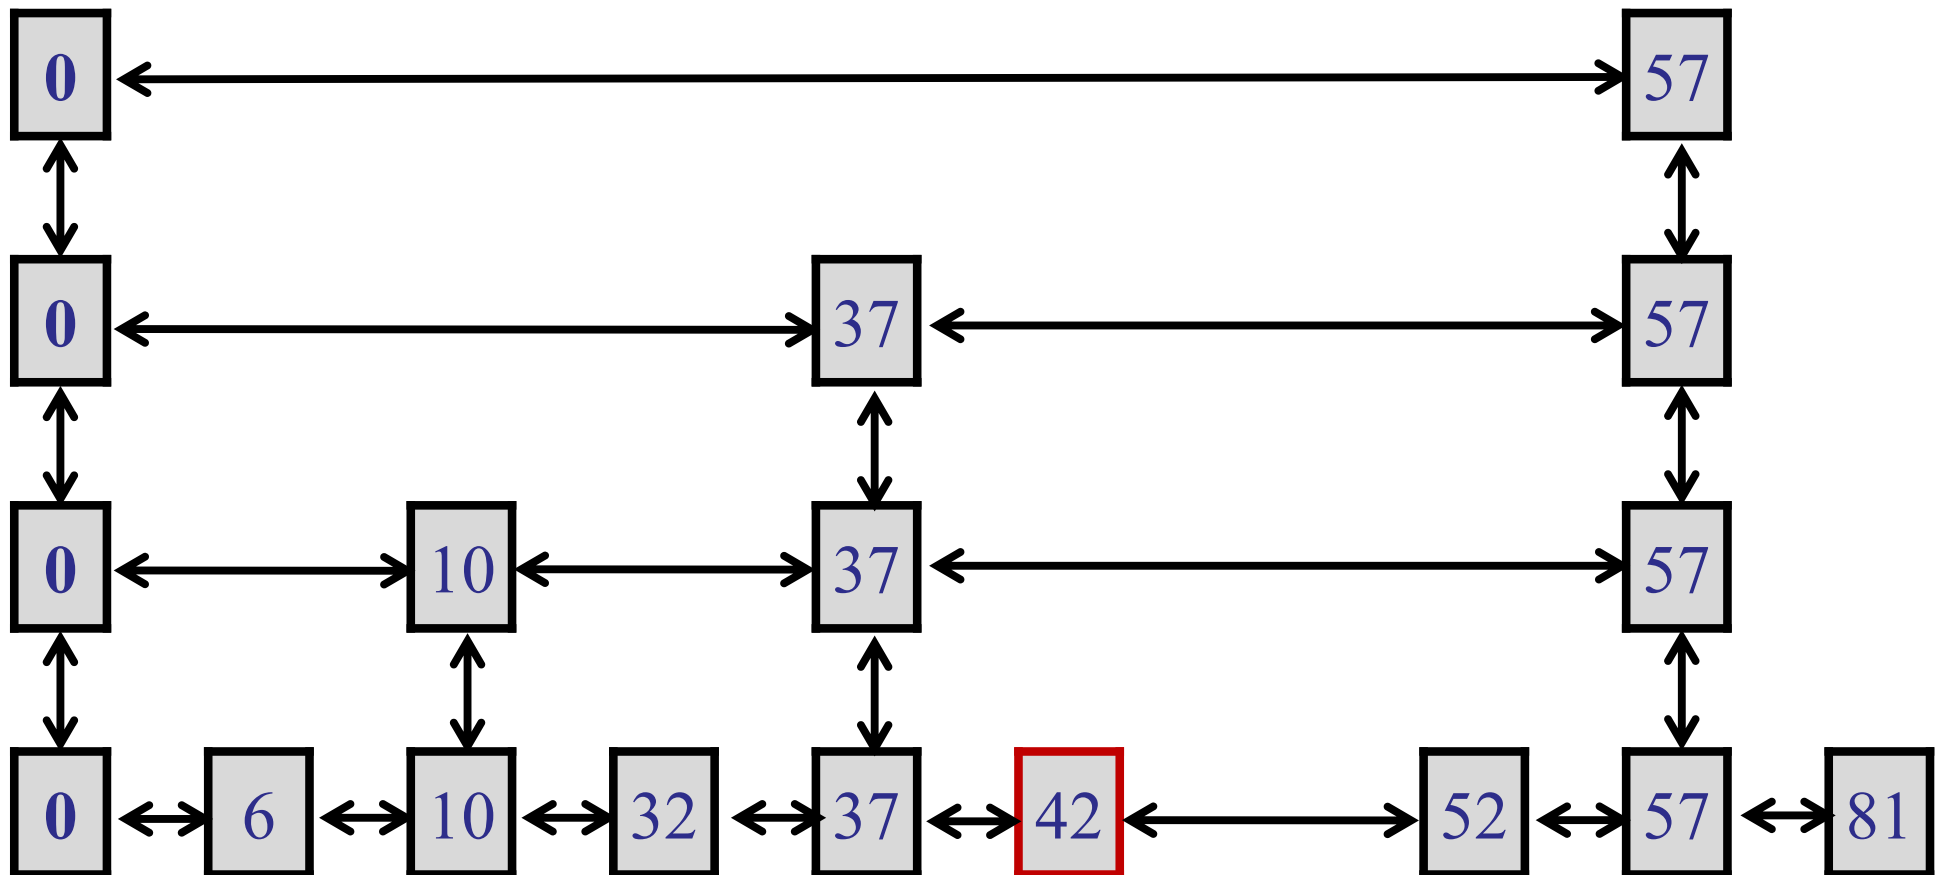
Level $\log(n)$: $O(1)$

SkipList

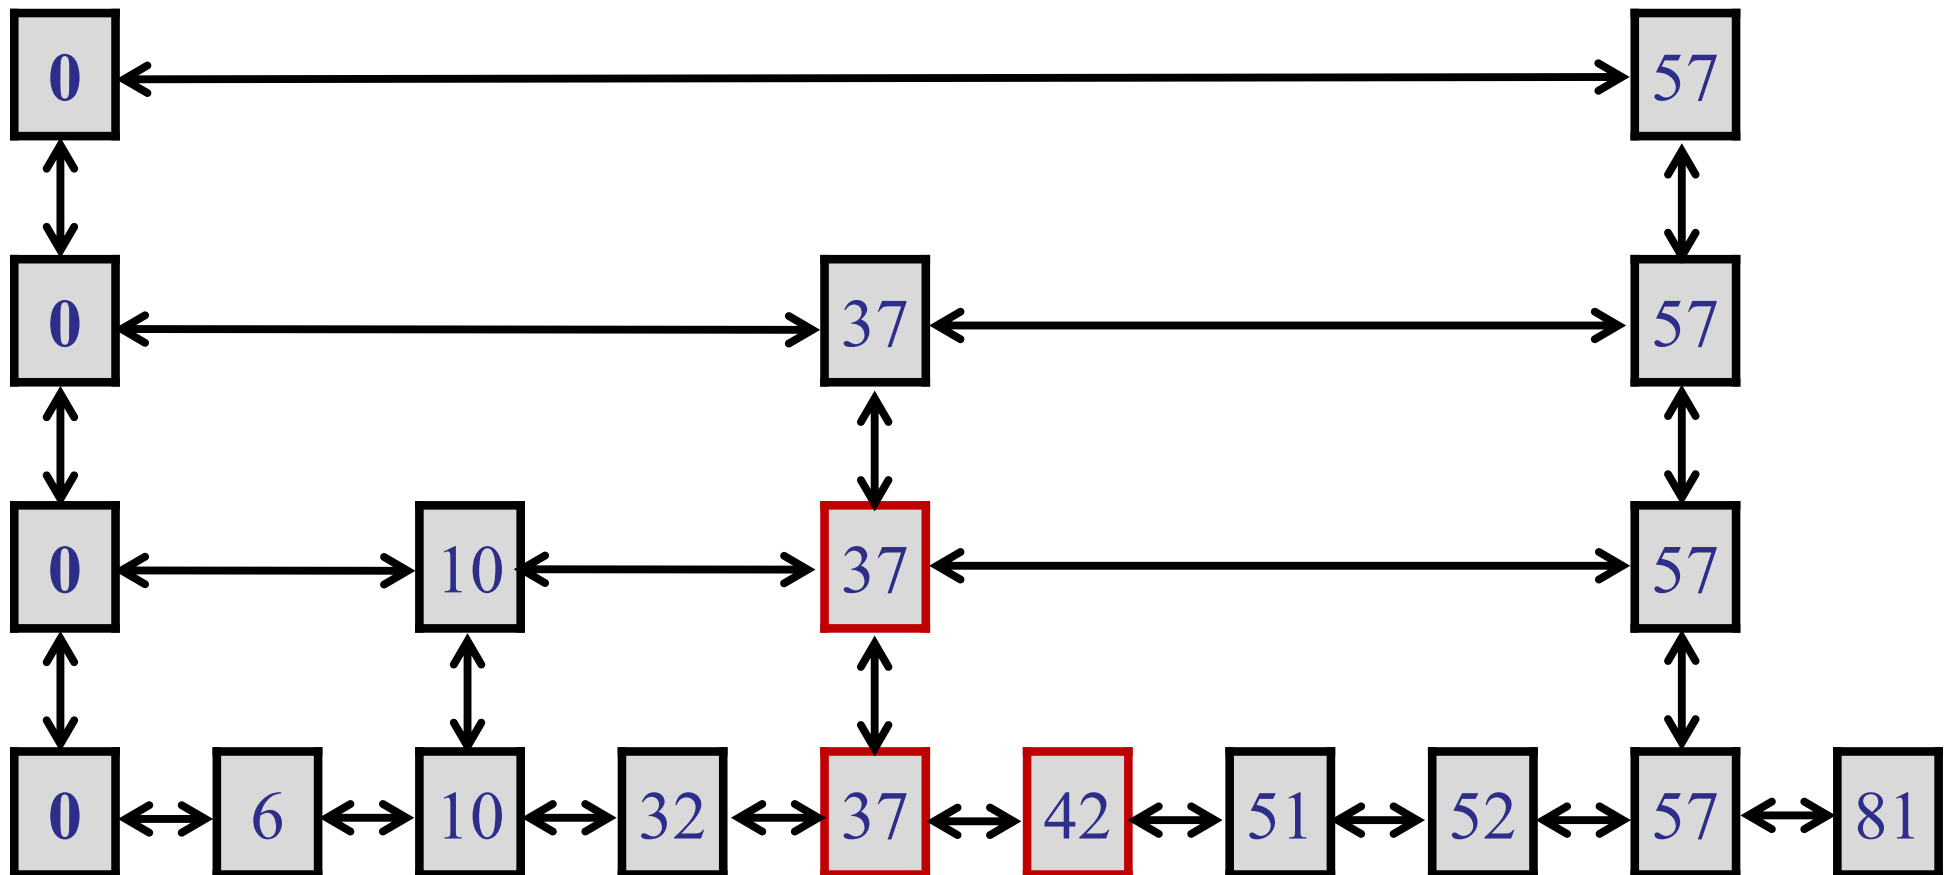
Randomized process:

- Not a perfect distribution.
- Good, on average.
- Really good, *almost always*.
- As usual with randomized algorithms, easy to implement, harder to analyze.

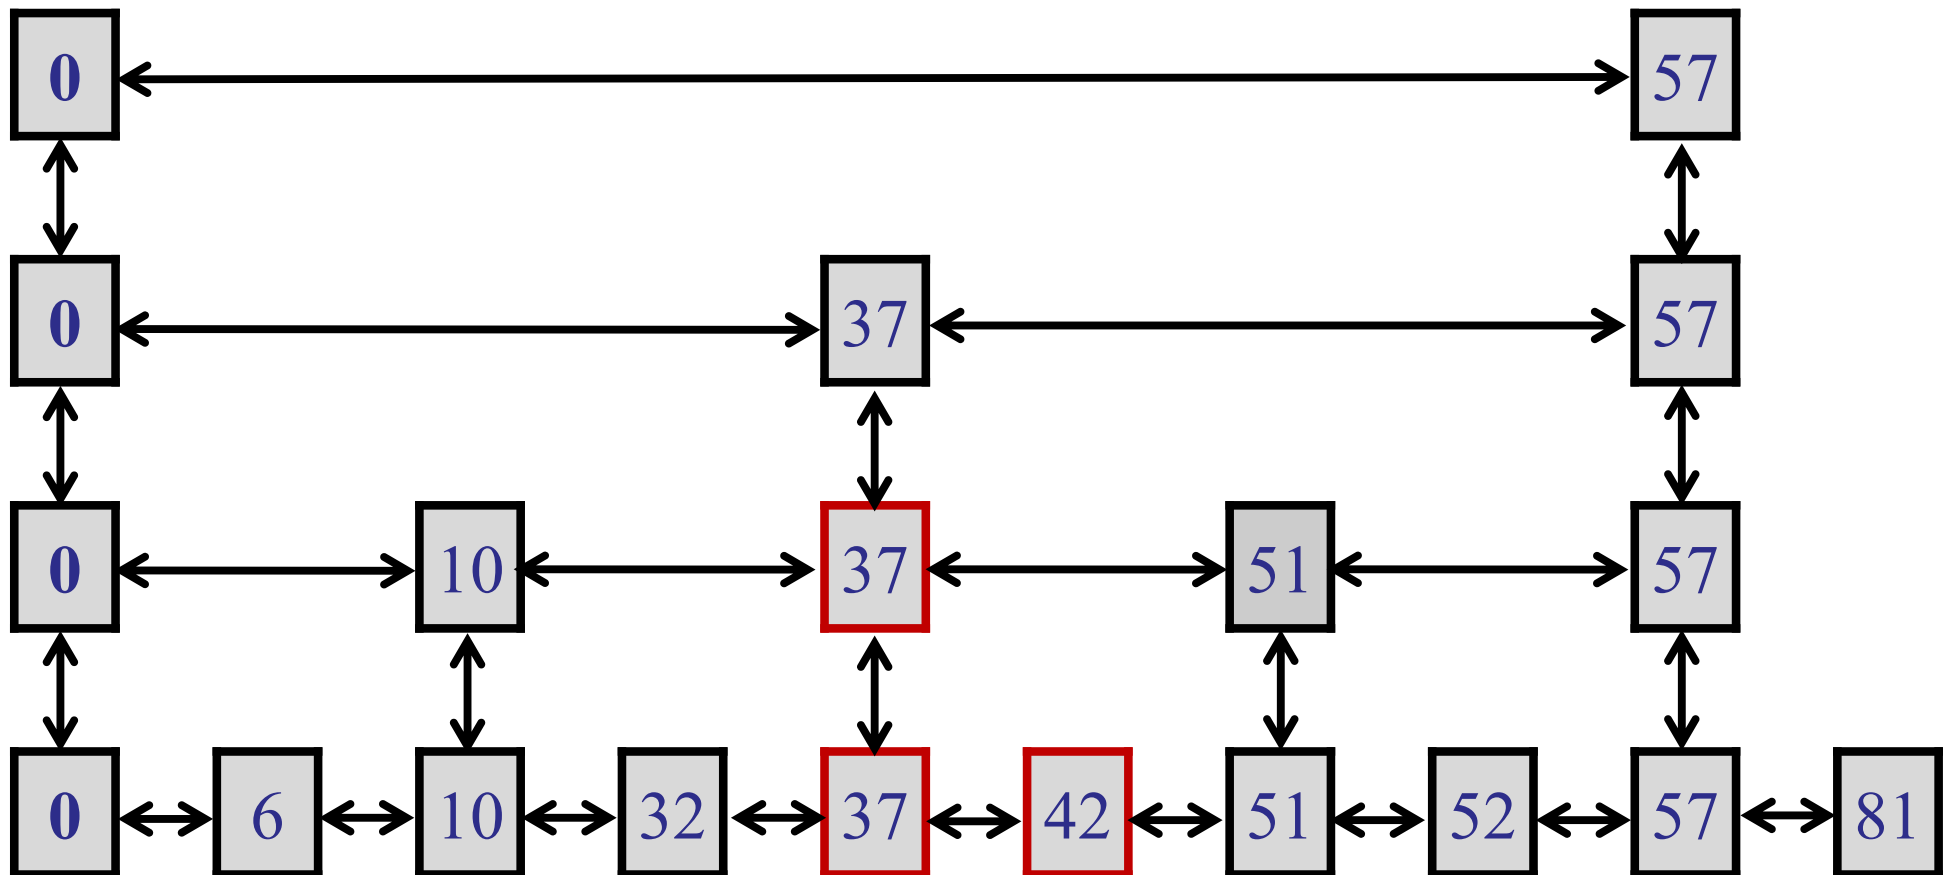
Example: insert(51)



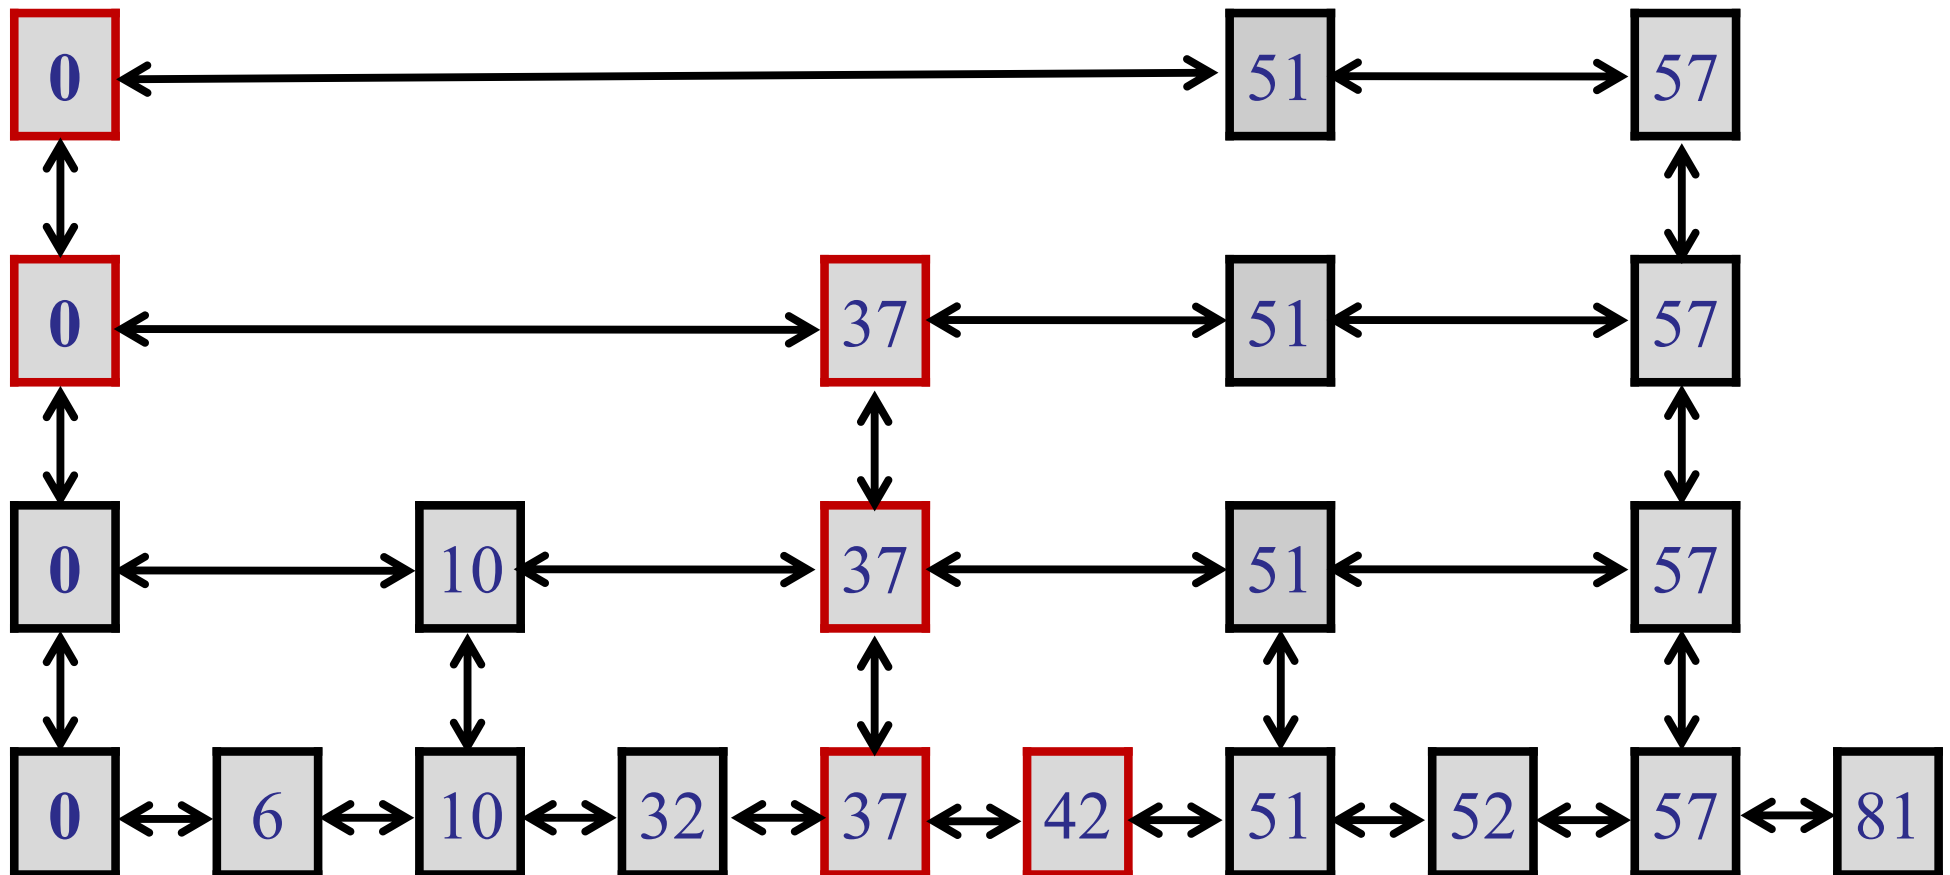
Example: insert (51)



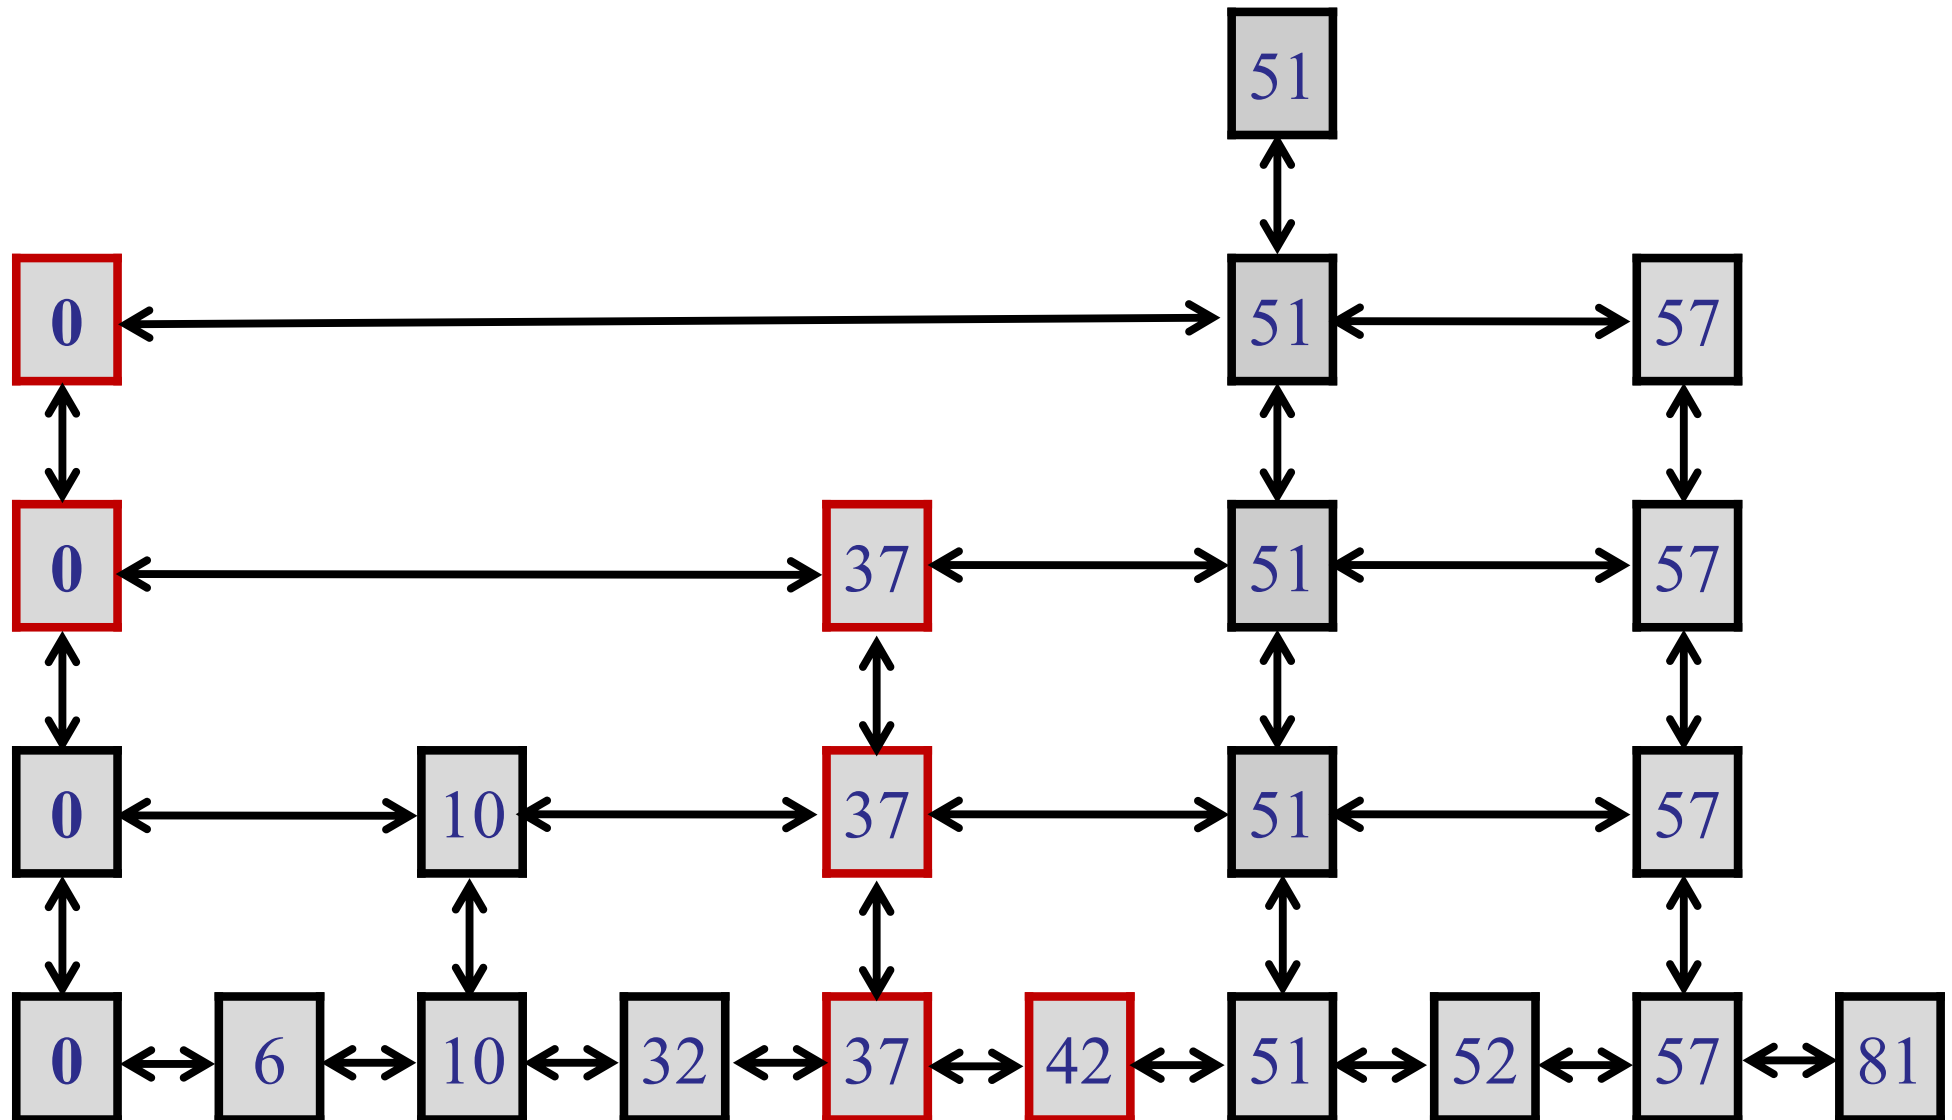
Example: insert (51)



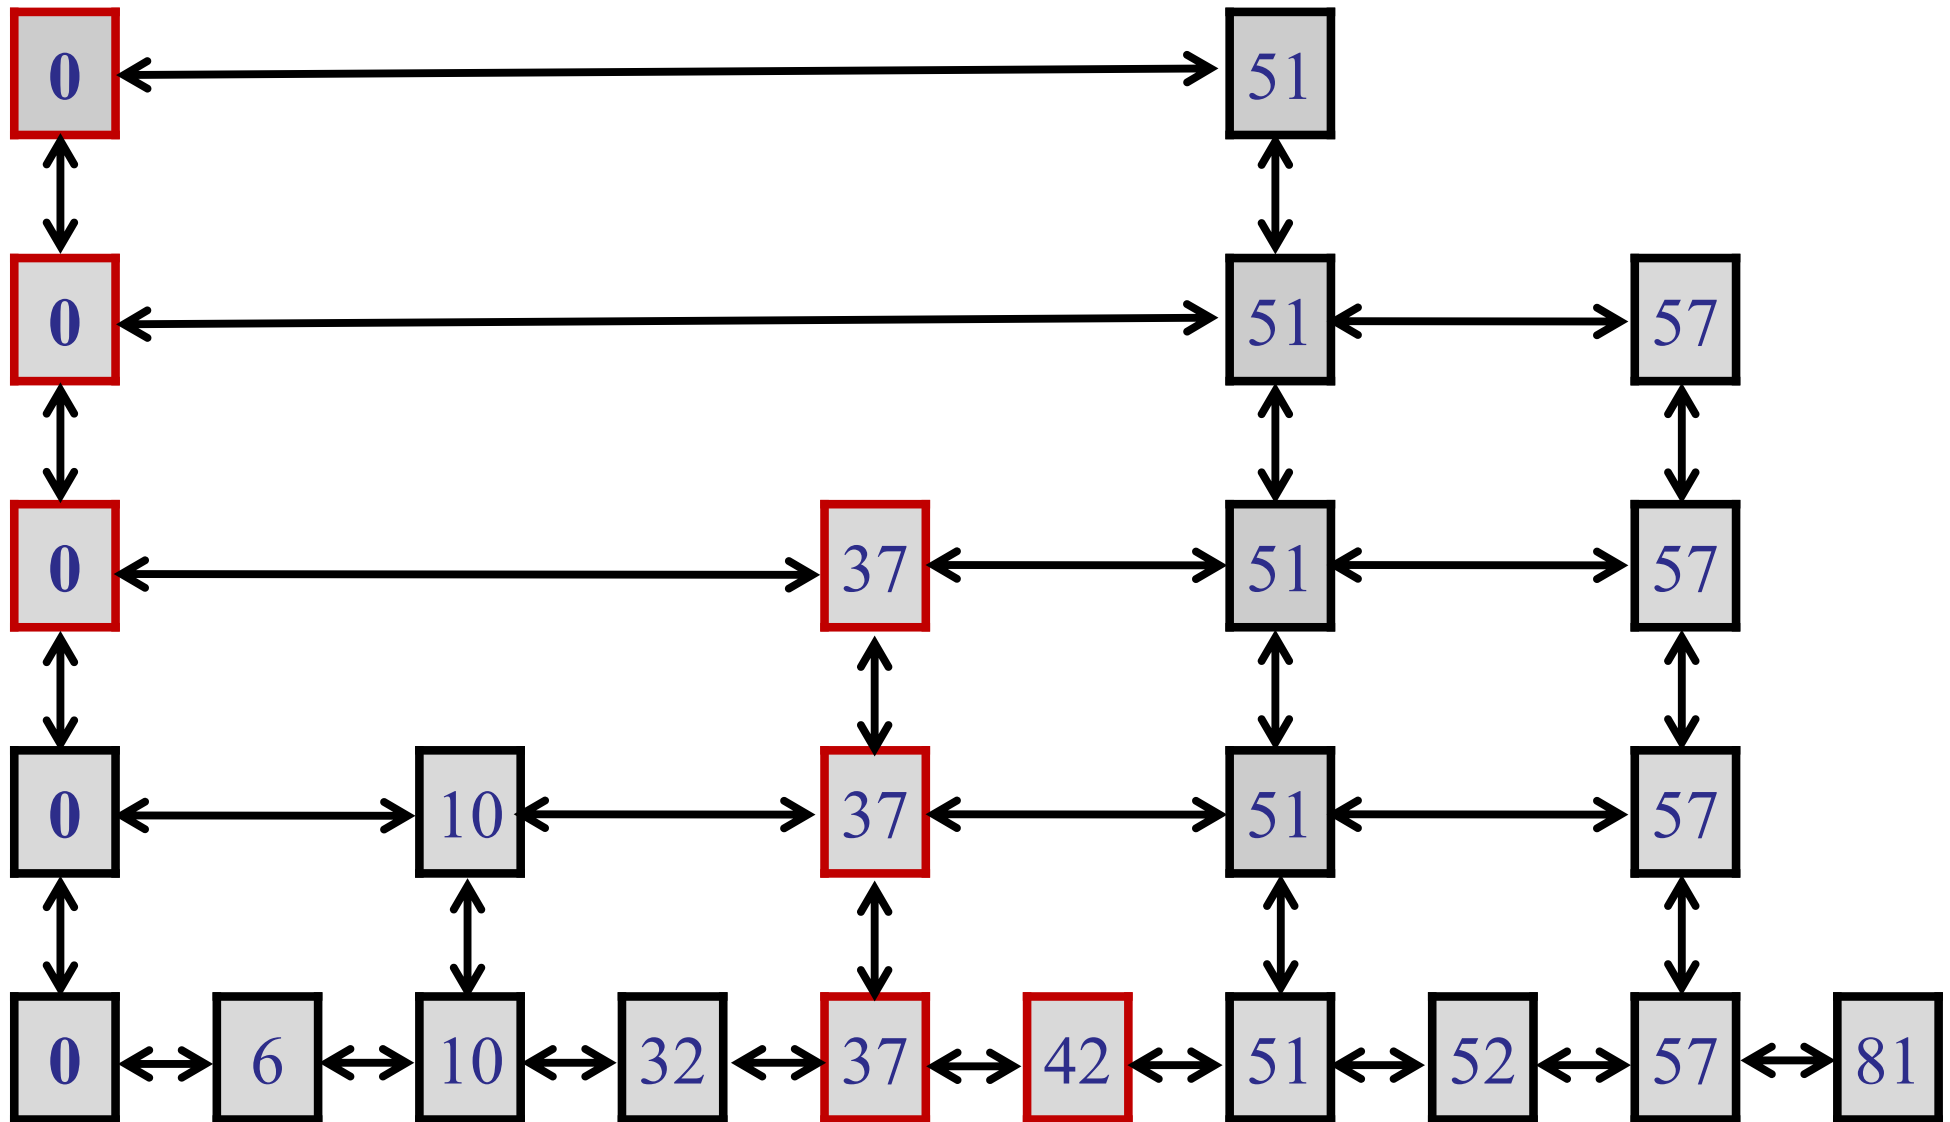
Example: insert(51)



Example: insert(51)



Example: insert(51)



SkipList Analysis

SkipList Analysis

Claim: Every **search** and **insert** operation completes in **$O(\log n)$** time *with high probability* (i.e., at least **$1 - 1/n$**).

Key steps:

- Analyze number of levels in a SkipList.
- Look at distribution of promotions:

SkipList is efficient when each jump skips about the same number of elements.

SkipList Analysis

Claim: With high probability, a SkipList with n elements has $O(\log n)$ levels.

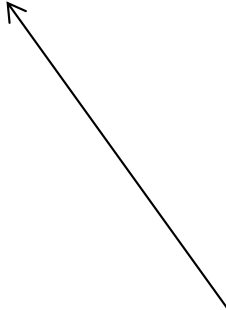
SkipList Analysis

Claim: With high probability, a SkipList with n elements has $O(\log n)$ levels.

Proof:

Fix an element x .

$$\Pr[x \text{ is higher than } c \log(n)] \leq \frac{1}{2^{c \log n}} \leq \frac{1}{n^c}$$



Probability of flipping
more than $c \log(n)$ heads
in a row!

SkipList Analysis

Proof: Fix an element x .

$$\Pr[x \text{ is higher than } c \log(n)] \leq \frac{1}{2^{c \log n}} \leq \frac{1}{n^c}$$

Define:

- e_1 = probability first element is too high $< 1/n^c$
- e_2 = probability second element is too high $< 1/n^c$
- ...
- e_n = probability n^{th} element is too high $< 1/n^c$

$$\Pr(\text{any element is too high}) \leq \frac{n}{n^c} \leq \frac{1}{n^{c-1}}$$

SkipList Analysis

Claim: With high probability, a SkipList with n elements has $O(\log n)$ levels.

SkipList Analysis

Done!



Claim: Every **search** and (**insert**) operation completes in $O(\log n)$ time *with high probability*.

Done!



Key steps:

- Analyze number of levels in a SkipList.
- Look at distribution of promotions:

SkipList is efficient when each jump skips about the same number of elements.

SkipList Analysis

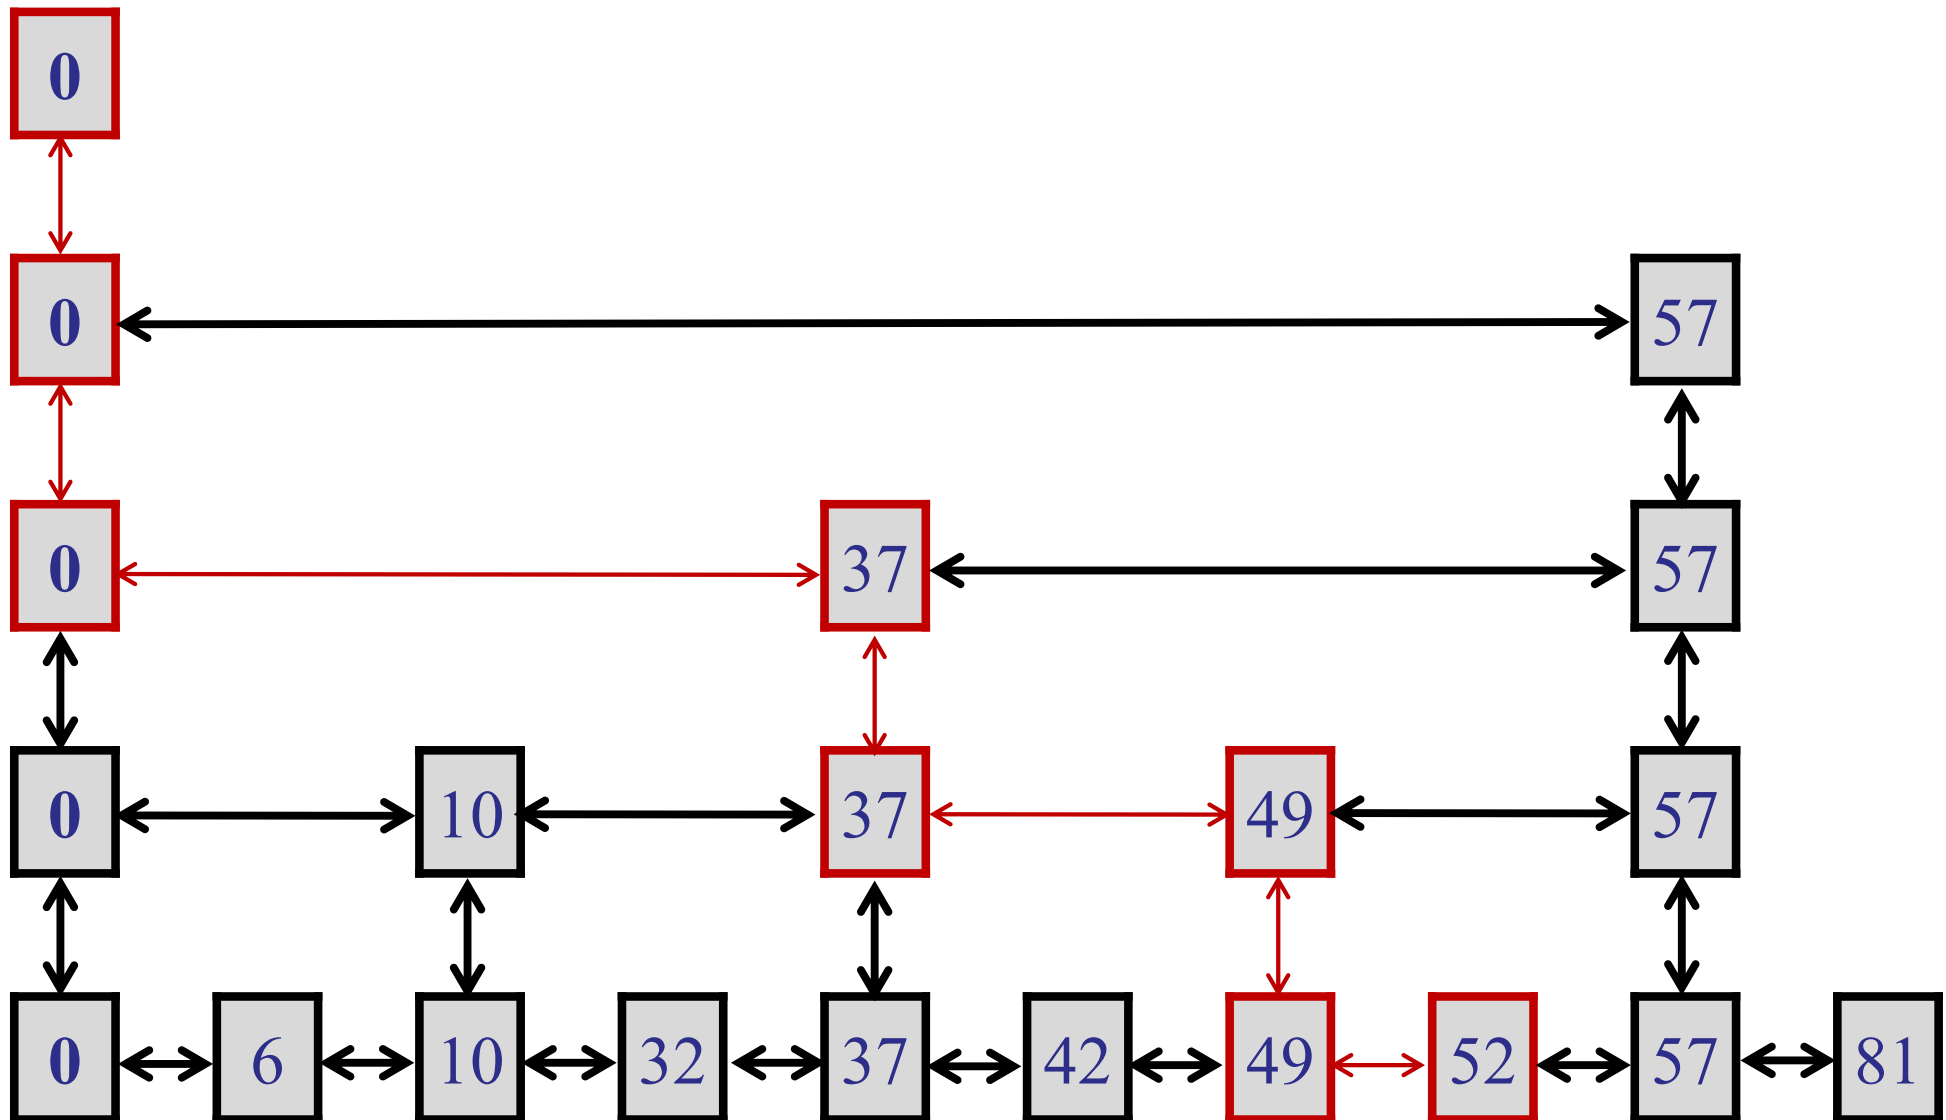
Analyzing a search:

Neat idea: analyze the search backwards.

- Start at leaf.
- For each node visited:
 - If node was not promoted (TAILS), go left.
 - If node was promoted (HEADS), go up.
- Stop at root of tree.

Exact same
traversal as
search, in reverse!

Example: search (52)



SkipList Analysis

Analyzing a search:

Neat idea: analyze the search backwards.

- Start at leaf.
- For each node visited:
 - If node was not promoted (TAILS), go left.
 - If node was promoted (HEADS), go up.
- Stop at root of tree.

Occurs at
most $O(\log n)$
times!



SkipList Analysis

Analyzing a search:

Neat idea: analyze the search backwards.

- Start at leaf.
 - For each node visited:
 - If node was not promoted (TAILS), go left.
 - If node was promoted (HEADS), go up.
 - Stop at root of tree.
- At most $O(\log n)$
-

New question: How many times to flip a coin until we get $c \log(n)$ heads?

SkipList Analysis

Claim: With high probability, after $10 c \log n$ coin flips, you get $c \log n$ heads.

SkipList Analysis

Proof:

- Say we flip $10 c \log(n)$ coins.
- $\Pr[\text{exactly } c \log(n) \text{ heads}] =$

$$\binom{10c \log n}{c \log n} \left(\frac{1}{2}\right)^{c \log n} \left(\frac{1}{2}\right)^{9c \log n}$$

Number of ways to choose
 $c \log(n)$ heads out of all the flips:

TTTTHH

TTHTH

...

SkipList Analysis

Proof:

- Say we flip $10c \log(n)$ coins.
- $\Pr[\text{exactly } c \log(n) \text{ heads}] =$

$$\binom{10c \log n}{c \log n} \left(\frac{1}{2}\right)^{c \log n} \left(\frac{1}{2}\right)^{9c \log n}$$

Probability each of the
H comes up heads.

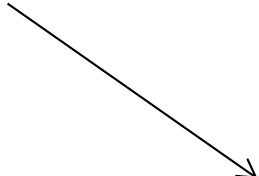
Probability each of the
T comes up tails.

SkipList Analysis

Proof:

- Say we flip $10 c \log(n)$ coins.
- $\Pr[\text{exactly } c \log(n) \text{ heads}] =$


bad case! $\binom{10c \log n}{c \log n} \left(\frac{1}{2}\right)^{c \log n} \left(\frac{1}{2}\right)^{9c \log n}$



- $\Pr[\text{at most } c \log(n) \text{ heads}] \leq$

$$\binom{10c \log n}{c \log n} \left(\frac{1}{2}\right)^{9c \log n}$$

If all $9c \log(n)$
are tails, then not
enough heads!



SkipList Analysis

Bounding binomials:

$$\left(\frac{y}{x}\right)^x \leq \binom{y}{x} \leq \left(\frac{ey}{x}\right)^x$$

SkipList Analysis

Bounding binomials:

$$\left(\frac{y}{x}\right)^x \leq \binom{y}{x} \leq \left(\frac{ey}{x}\right)^x$$

$$\begin{aligned} \binom{10c \log n}{\log n} &\leq \left(\frac{e10c \log n}{\log n}\right)^{\log n} \leq (10e)^{\log n} \\ &\leq n^{\log(10e)} \end{aligned}$$

SkipList Analysis

Proof:

- Say we flip $10 c \log(n)$ coins.
- $\Pr[\text{at most } c \log(n) \text{ heads}] \leq$

$$\binom{10c \log n}{c \log n} \left(\frac{1}{2}\right)^{9c \log n}$$
$$\leq n^{c \log(10e)} \frac{1}{n^{9c}}$$

$$\alpha = c(9 - \log(10) - \log(e)) \leq \frac{1}{n^\alpha}$$

Generalize for other values of 10...

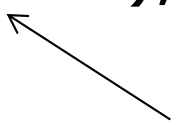
SkipList Analysis

Claim: With high probability, after $O(\log n)$ coin flips, you get $c \log n$ heads.

SkipList Analysis

Analyzing a search:

Neat idea: analyze the search backwards.

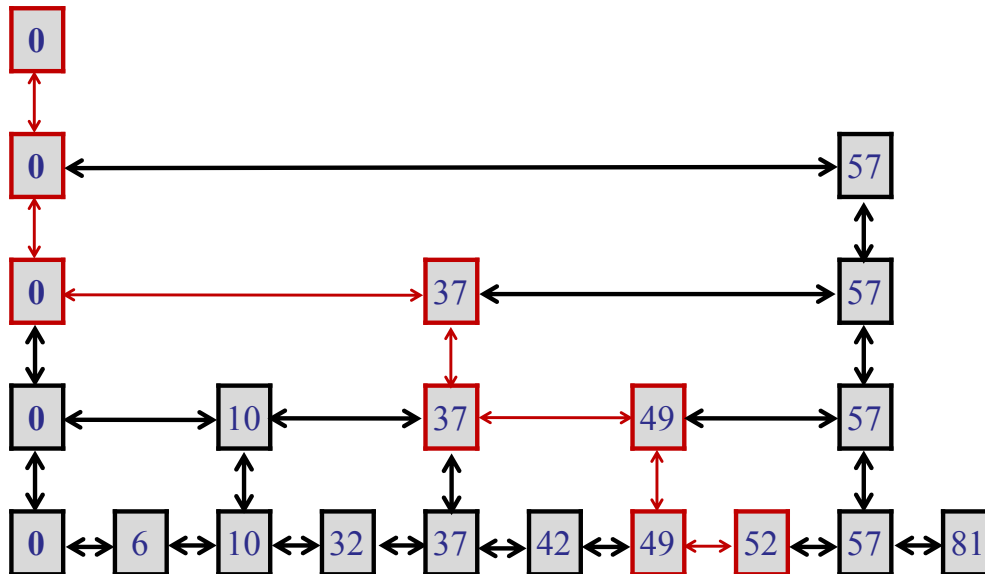
- Start at leaf.
 - For each node visited:
 - If node was not promoted (TAILS), go left.
 - If node was promoted (HEADS), go up.
 - Stop at root of tree.
- At most $O(\log n)$
- 

After $10 c \log(n)$ coin flips, we will get $c \log(n)$ heads (with high probability).

SkipList Analysis

Claim: With high probability, after $O(\log n)$ coin flips, you get $c \log n$ heads.

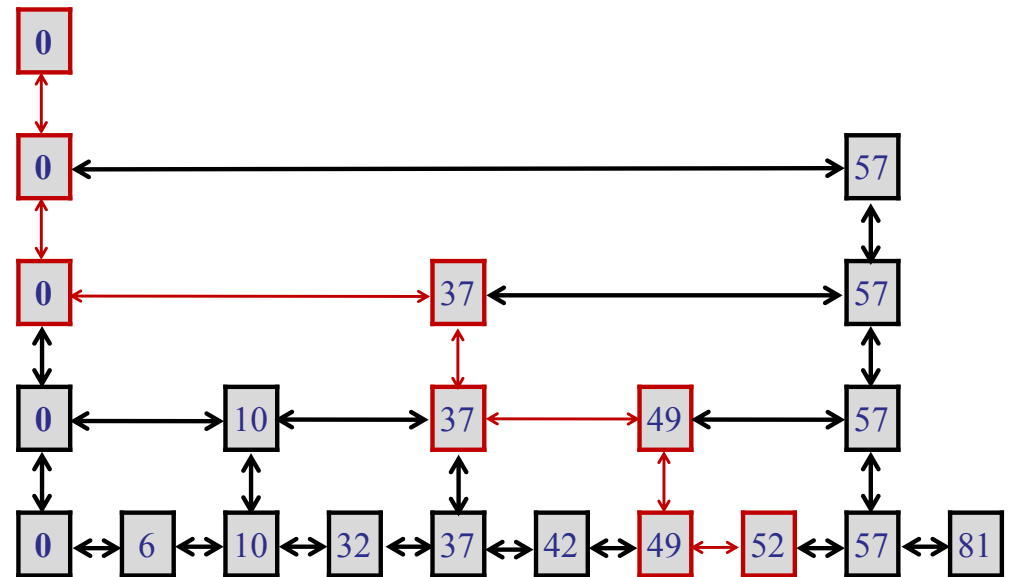
Conclusion: Each search takes $O(\log n)$ steps with high probability.



Conclusions

SkipLists

- Simple, efficient, randomized search structure.
- Easy to implement.



Analysis:

- Nice randomized calculations.
- Key idea: analyze backwards!
- Reduce to the problem of flipping coins.