

CS2040S

Data Structures and Algorithms

(e-learning edition)

What comes next?

Algorithms Everywhere

Web browser:

Parsing

Substring manipulation (Week 7)

XML trees (Week 5)



Internet

Internet routing:

TCP (congestion control)

IP routing

BGP (Bellman-Ford, Week 9)

Content caching (Week 7)

DNS



Web servers:

Load balancing (PS 2)

Scheduling (Week 8)

Memory allocation

Google:

PageRank

(Week 10)

String matching

(Week 7)

Database:

B-trees (Week 5)

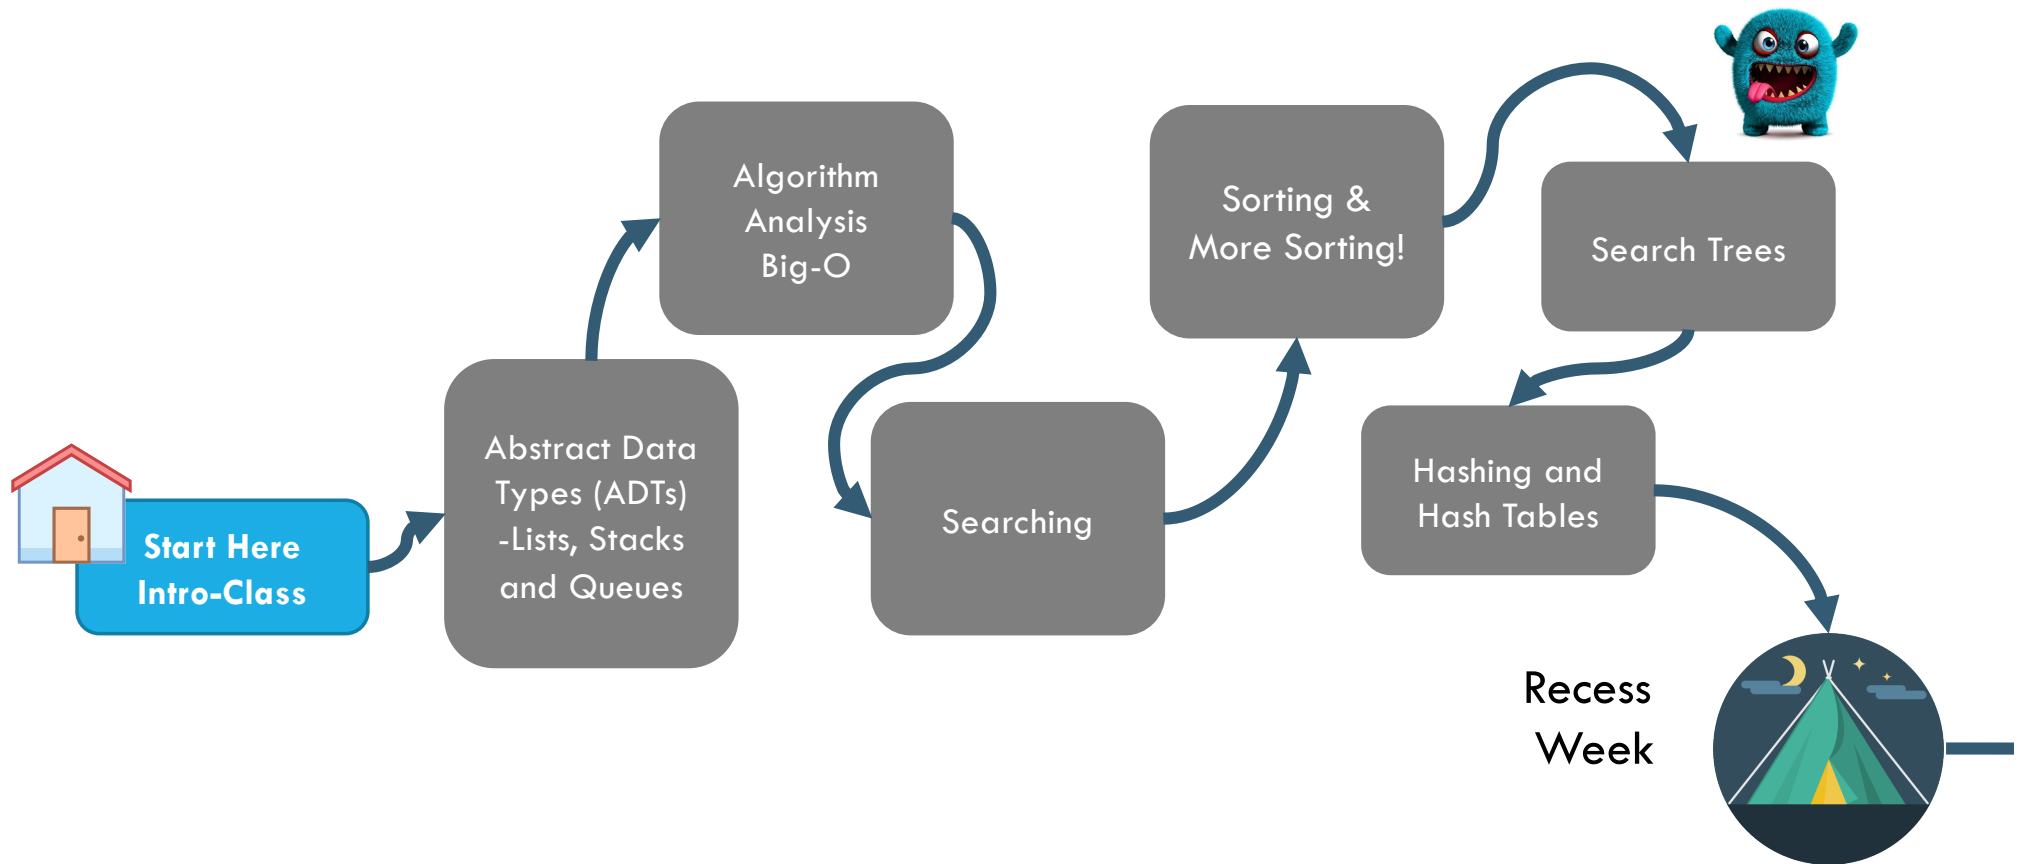
Search (Week 2)

Sorting (Week 3)



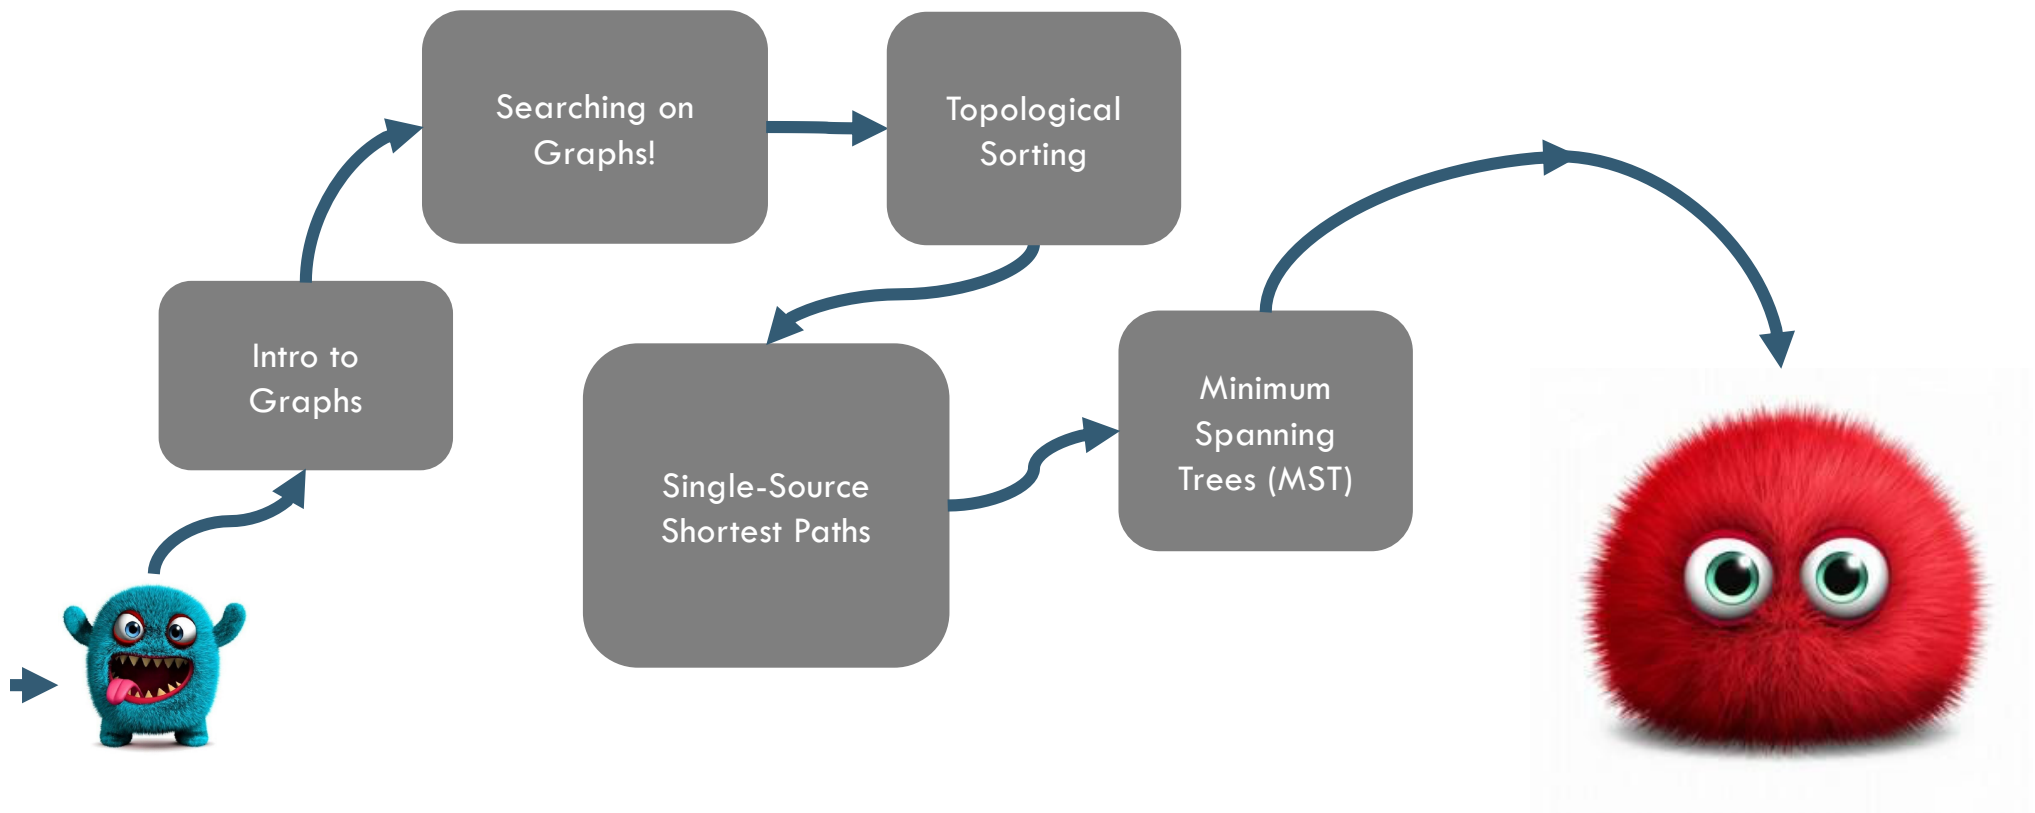
What is CS2040S about?

COURSE STRUCTURE



PART I: ORGANIZING YOUR DATA

COURSE STRUCTURE



PART II: MODELLING AND SOLVING PROBLEMS

Data Structures

Problem Solving

Algorithms

Desirable features of your algorithm:



How do you choose the right algorithm for the right problem?

How do you design new algorithms for new problems?

FRAMEWORK: ALGORITHM & DATA STRUCTURE

What problem does it solve?

How does it work?

How to implement it?

What is its asymptotic performance?

What is its real world performance?

What are the trade-offs?



GOALS

By the end of this course, you should be able to:

- **Apply algorithmic thinking and techniques** for solving computational problems.
- **Describe the structure and operation** of different **data structures and algorithms** under the standard computational model.
- **Assess the suitability** of different data-structures and algorithms for a specific computational problem.
- **Adapt existing data-structures and algorithms** to solve specific computational problems.

What comes next?

What's next?

Topic 1: Searching and Optimization

Topic 2: Sorting

Topic 3: Trees

Topic 4: Hashing

Topic 5: Shortest Paths

Topic 6: Minimum Spanning Trees

Topic 1: Searching and Optimization

Binary Search:

- Fast way to search monotonic data.
- Fast way to find max/min for convex data.

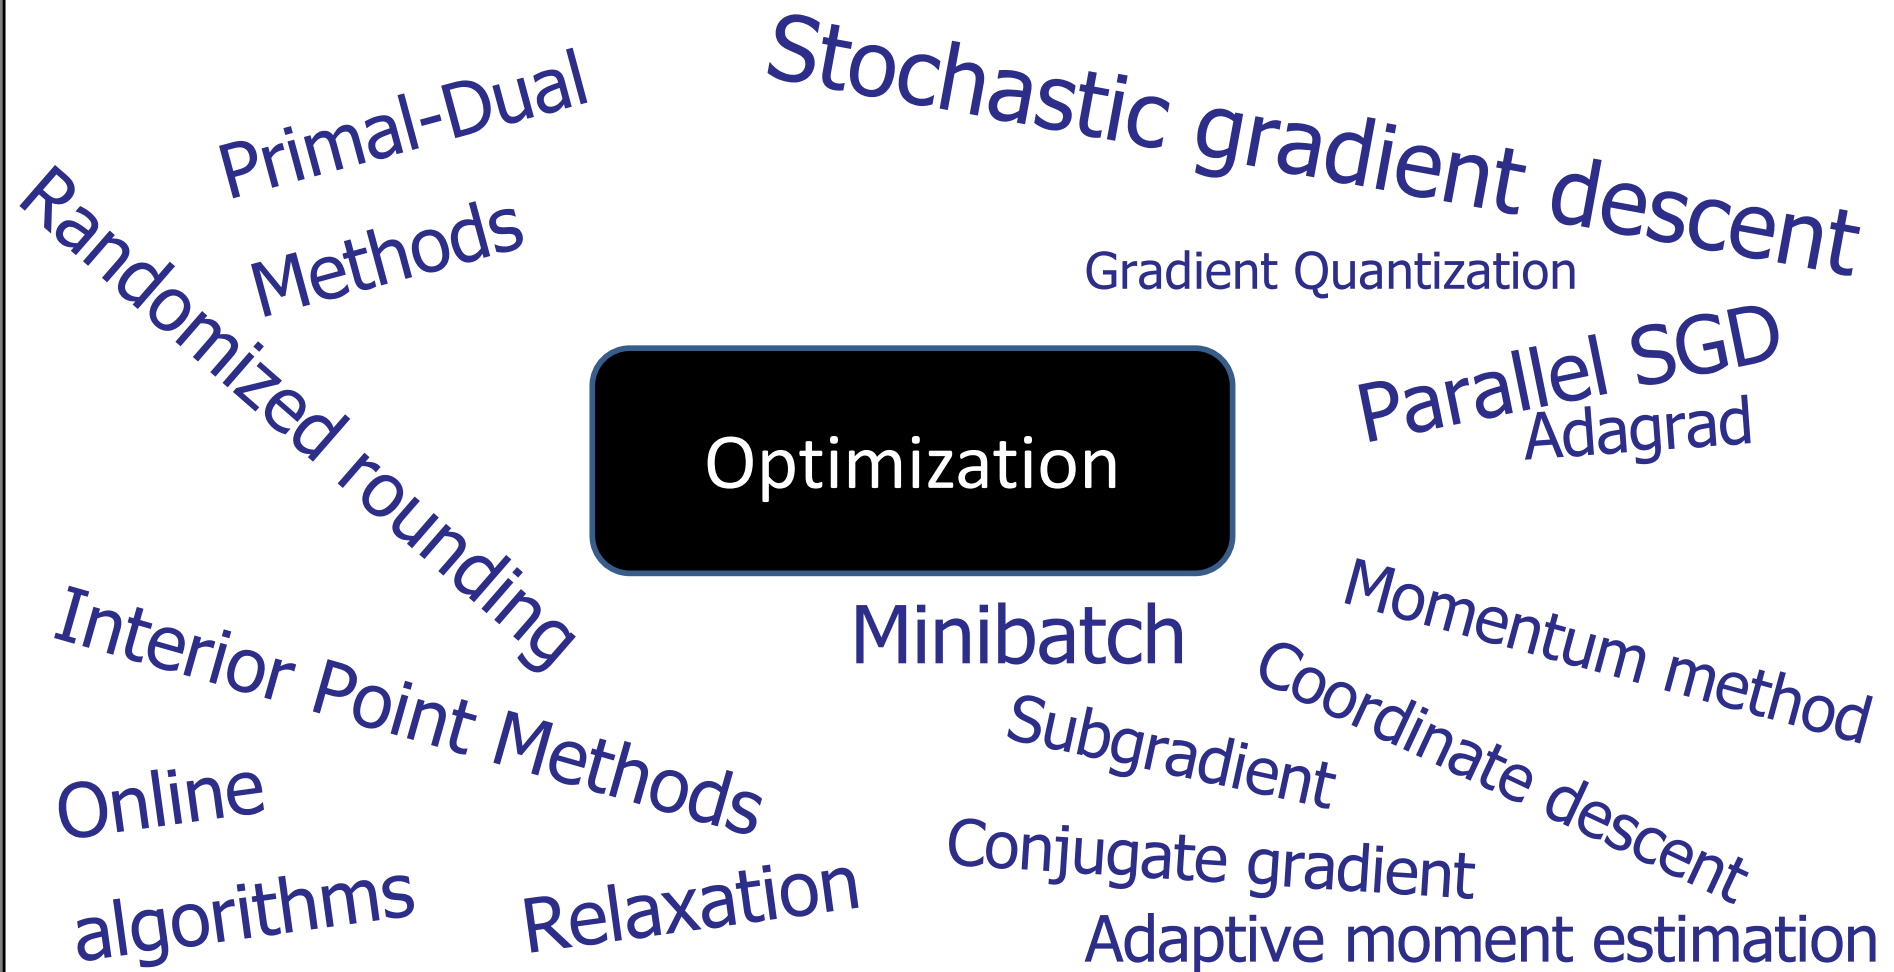
Newton's Method:

- Fast way to search/max/min for convex, twice-differentiable data.

Gradient Descent:

- Fast way to search/max/min for convex data.
- Fast way to optimize wide variety of functions.

Huge area of research and development:



Optimization algorithms:
General techniques.

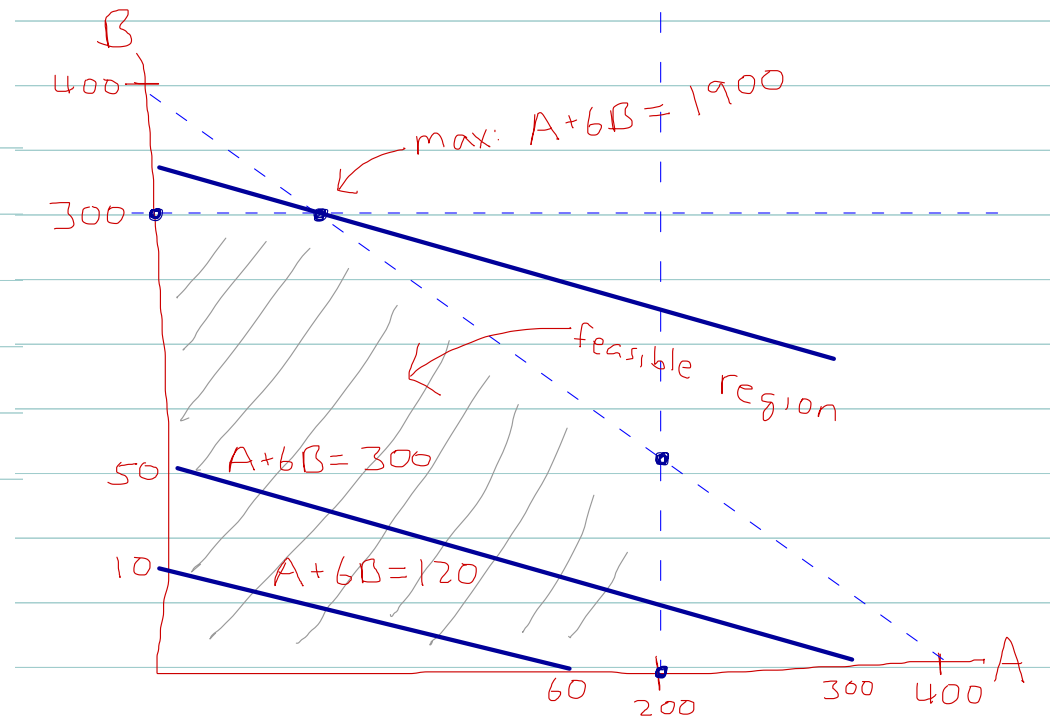
Machine learning:
How to train a model?

Optimization Algorithms

Linear Programming:

- How to optimize a linear function subject to linear constraints.
- E.g., simplex method

$$\begin{aligned} &\max (A+6B) \\ &\text{where: } A \leq 200 \\ &\quad B \leq 300 \\ &\quad A+B \leq 400 \\ &\quad A \geq 0 \\ &\quad B \geq 0 \end{aligned}$$



Optimization Algorithms

Linear Programming:

- How to optimize a linear function subject to linear constraints.
- E.g., simplex method

Applications:

- Graph algorithms (e.g., max-flow)
- Approximation algorithms (e.g., weighted vertex cover, weighted set cover, multicommodity flow)
- And many, many real-world problems.

Optimization Algorithms

Linear Programming:

- How to optimize a linear function subject to linear constraints.
- E.g., simplex method

And more...

- Semidefinite programming
- Integer linear programming (NP-hard)
- Quadratic programming (NP-hard)
- Constraint satisfaction (NP-hard)

Topic 2: Sorting

Fast Sorting Algorithms:

- QuickSort
- MergeSort
- HeapSort

Key properties:

- Running time
- Space usage
- In-place

Sorting Faster!

Multi-pivot QuickSort...

Sorting Faster!

Yahoo TeraSort:

- Each node has:
 - 8 cores: 2GHz
 - 8 GB RAM
 - 4 disks: 4TB each
- 40 nodes / rack (interconnect: 1GB/s switch)
- 25-100 racks (interconnect: 8GB/s switch)

➔ ~ 16,000 cores

Sorting Faster!

Yahoo TeraSort:

- Each node has:
 - 8 cores: 2GHz
 - 8 GB RAM
 - 4 disks: 4TB each
- 40 nodes / rack (interconnect: 1GB/s switch)
- more racks (interconnect: 8GB/s switch)

➔ 50,400 cores

2013:

Yahoo (Hadoop) sorts
100TB of data in 72
minutes.

Sorting Faster!

DataBricks TeraSort:

- 206 nodes
- 6,592 cores

Record (2014):

DataBricks (Spark) sorts
100TB of data in 23
minutes.

DataBricks PetaSort:

- 190 nodes
- 6,080 cores

Record (2014):

DataBricks (Spark) sorts
1PB of data in 234
minutes.

Sorting Faster!

Tencent Sort:

- 512 nodes
- 5,024 cores

Record (2016):
98.8 seconds

Sorting Faster!

It's a race:

- High performance clusters.
- Hardware interconnect.
- Parallel performance.

Topic 3: Search Trees

Many types of search trees:

- **AVL trees** and Red-Black trees
- **B-trees**
- **Skip Lists**
- **Tries**
- **Interval Trees, Order Statistics Trees**
- **Range Trees, kd-Trees**
- Splay trees

Key tree-related questions:

High dimensional data:

- How should we store higher dimensional data?
- How should we index higher dimensional data?
- How should we search higher dimensional data?

Examples:

R-trees, spatial indices, quadtrees, Z-order curve, ...

Key tree-related questions:

Performance?

Predicting Performance

Example: 100 TB of data

1) Store data sorted in an array

- ⇒ Scan all the data: $O(n)$
- ⇒ (Binary) search: $O(\log n)$

2) Store data in a linked list

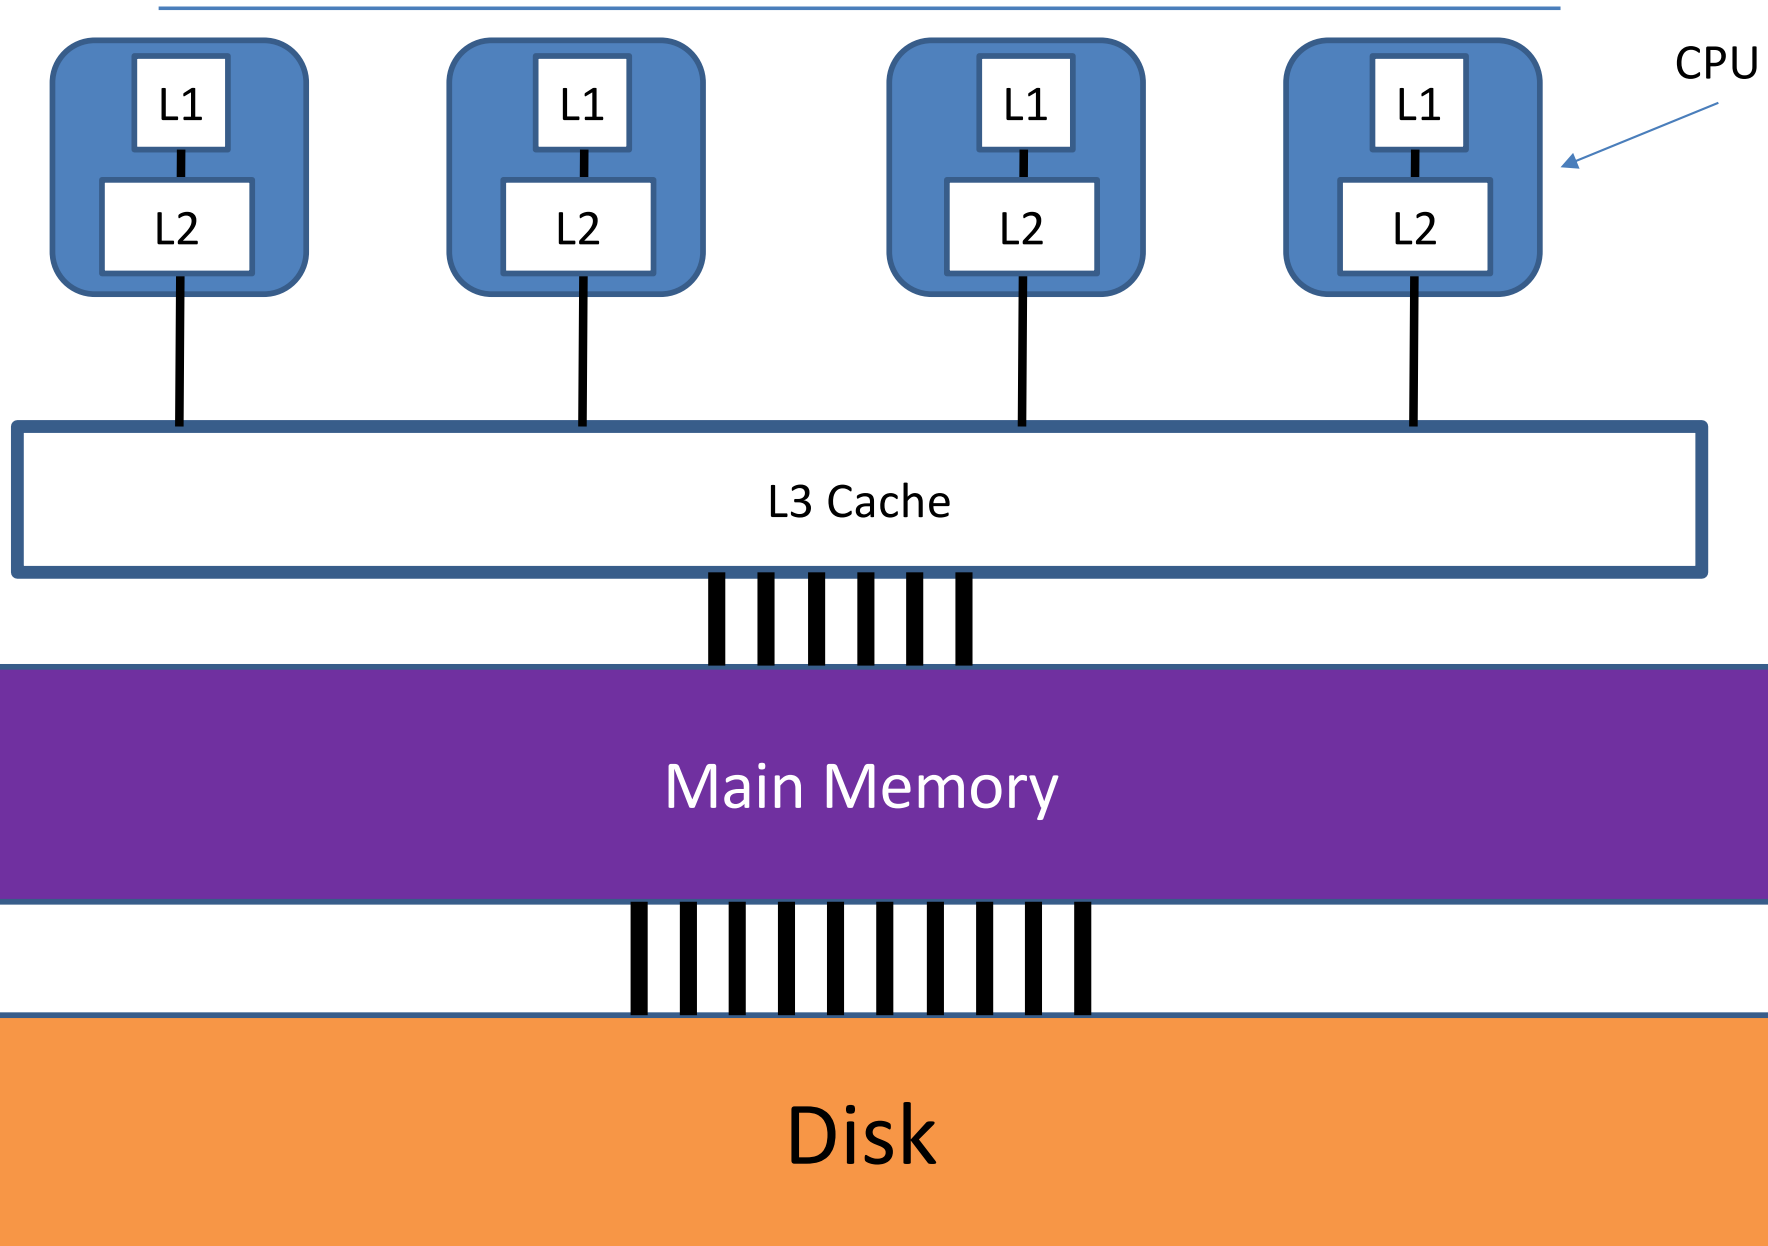
- ⇒ Scan all the data: $O(n)$
- ⇒ Search: $O(n)$

3) Store data in a red-black tree

- ⇒ Scan all the data: $O(n)$
- ⇒ Search: $O(\log n)$

Analysis is not predicting performance very well!

A Real Computer (?)



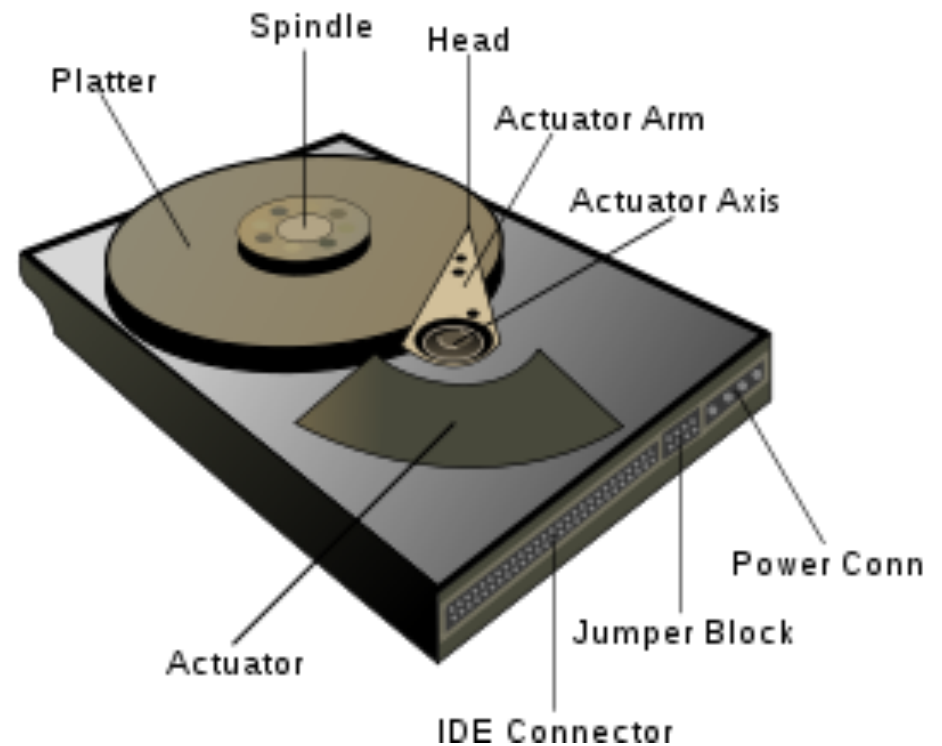
Disks

Where is most data stored? **Hard disk!**

- Magnetic
- Mechanical
- Slow (6000rpm = 10ms)

Two step access:

1. *seek (find right track)*
2. *read track*

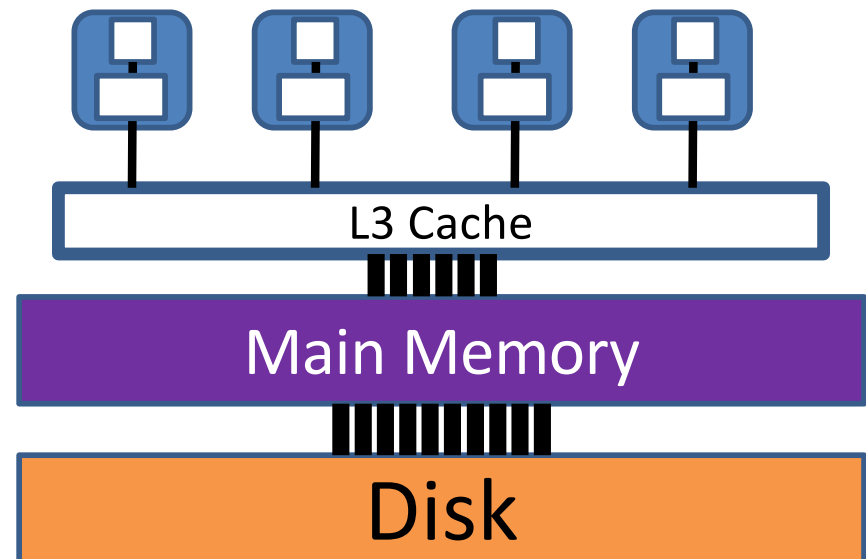


Haswell Architecture (2-18 cores)

Memory Type	size	line size	clock cycles
L1 cache	64 KB	64 B	~4
L2 cache	256 KB	64 B	~10
L3 cache	2-40 MB	64 B	40-74
L4 (optional)	128 MB		
Main Memory	< 128 GB	16 KB	~200-350
SSD Disk	BIG	Variable (e.g., 16KB)	~20,000
Disk	BIGGER	Variable (e.g., 16KB)	~20,000,000

Notes:

- Several other "caches" e.g., TLB, micro-op cache, instruction cache, etc.
- L1/L2 caches are per core.
- L3/L4 cache are shared per socket.
- Main memory shared cross socket.



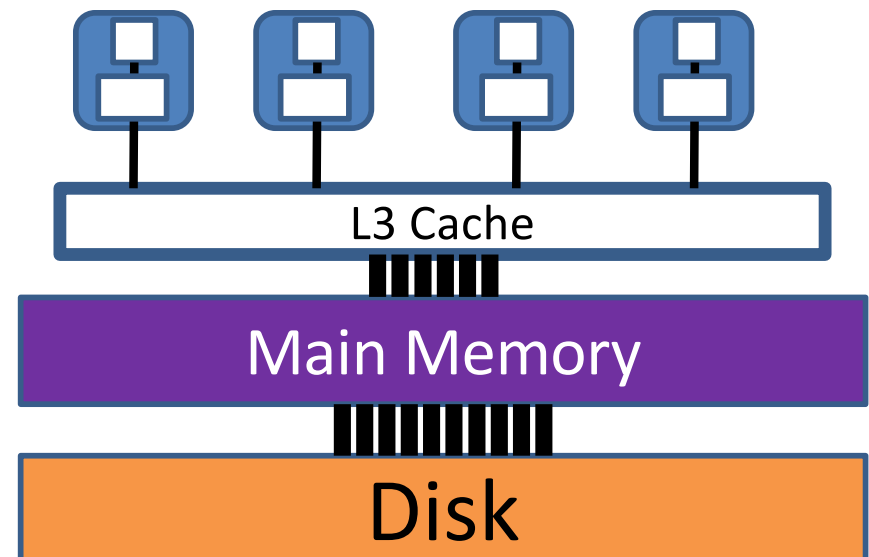
Haswell Architecture

A simple example calculation:

What fraction of operations "hit" each cache?

- ⇒ 90% L1 hit rate (4 cycles)
- ⇒ 8% L2 hit rate (10 cycles)
- ⇒ 2% main memory (300 cycles)

← *Just an example..*



Haswell Architecture

A simple example calculation:

What fraction of operations "hit" each cache?

- ⇒ 90% L1 hit rate (4 cycles)
- ⇒ 8% L2 hit rate (10 cycles)
- ⇒ 2% main memory (300 cycles)

← *Just an example..*

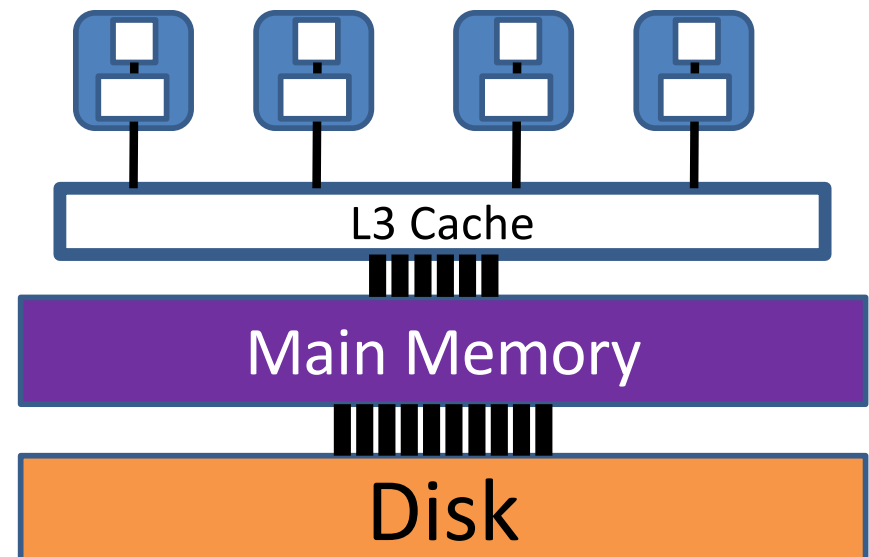
What fraction of time for each cache?

- ⇒ 35% waiting for L1
- ⇒ 8% waiting for L2
- ⇒ 57% waiting for main memory

Conclusion:

98% cache hit →

57% waiting on main memory



Haswell Architecture

A simple example calculation:

What fraction of operations "hit" each cache?

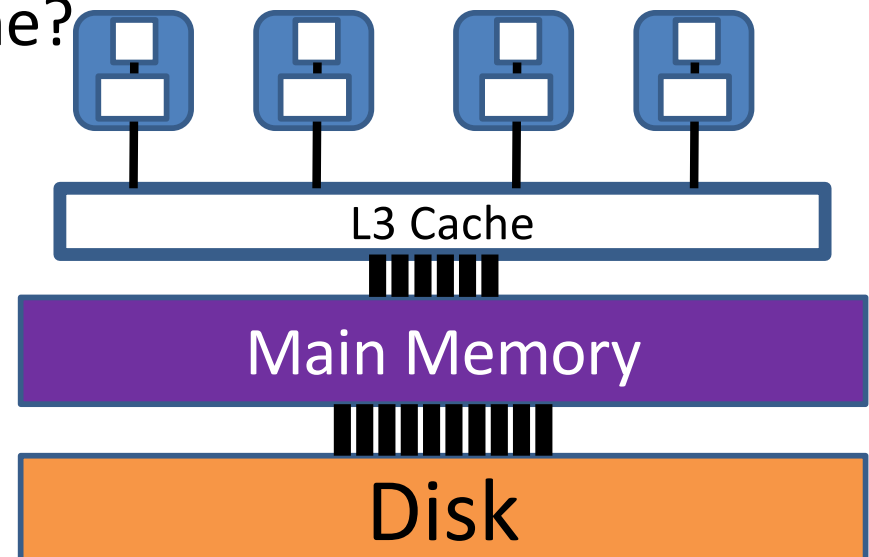
- ⇒ 90% L1 hit rate (4 cycles)
- ⇒ 8% L2 hit rate (10 cycles)
- ⇒ 1.8% main memory (300 cycles)
- ⇒ 0.2% disk (20,000,000 cycles)

← *Just an example..*

What fraction of time for each cache?

- ⇒ 99.98% waiting for disk

Disk is much, much worse!



Where is the bottleneck?

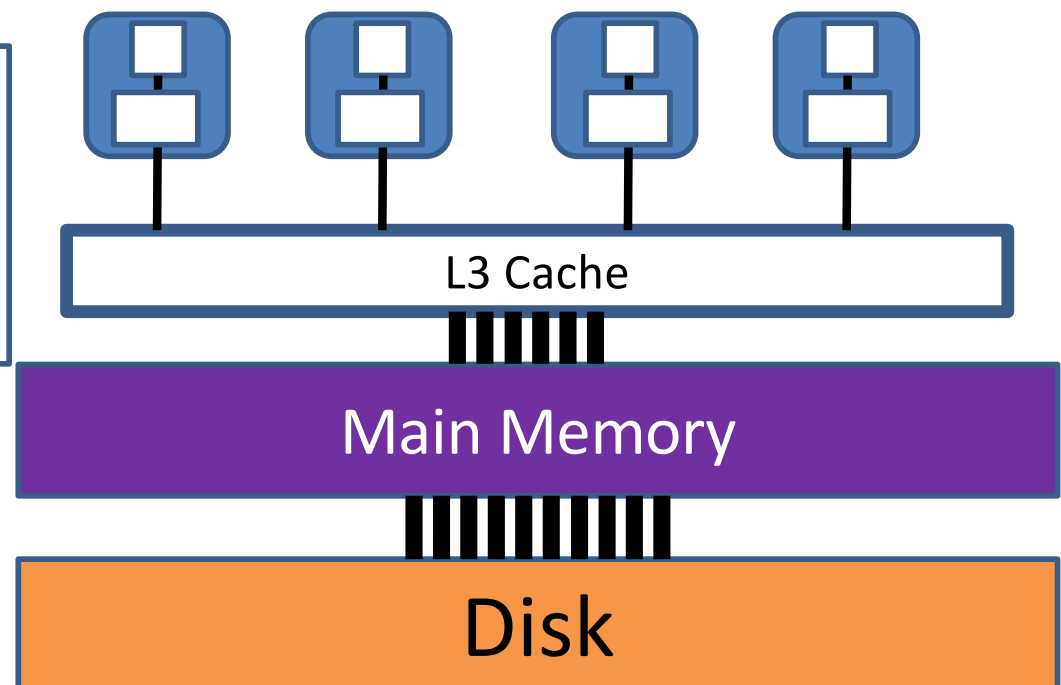
The bottleneck depends on the application:

- Small working set data lives in L1/L2 cache → **fast**.
- Medium working set data lives in main memory → **bottleneck is memory latency**.
- Big data lives on disk → **bottleneck is disk latency / bandwidth**.

For most applications,
one level dominates the
cost.

(Costs grow fast!

Largest level dominates.)



B-trees

Basic facts

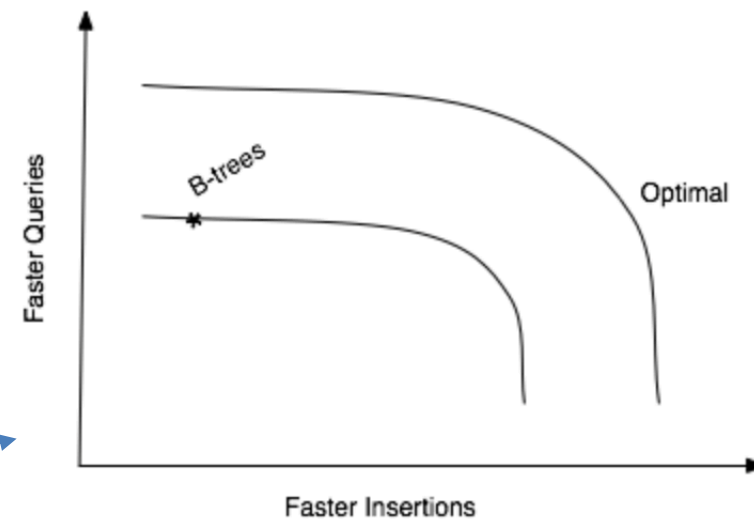
- One of the most important data structures out there today. (Variants used in all major databases.)
- Very fast. (Not just asymptotic analysis, but in practice nearly impossible to beat a well-implemented B-tree.)
- Benefit comes both from good cache performance, low overhead, good parallelization, etc.

Faster Trees?

Goal:

A external memory data structure with fast searches, and super-fast insertions/deletions.

“Write-optimized data structure.”



Percona advertisement graph
(not technical)

B-trees are not on the optimal insert/query tradeoff curve

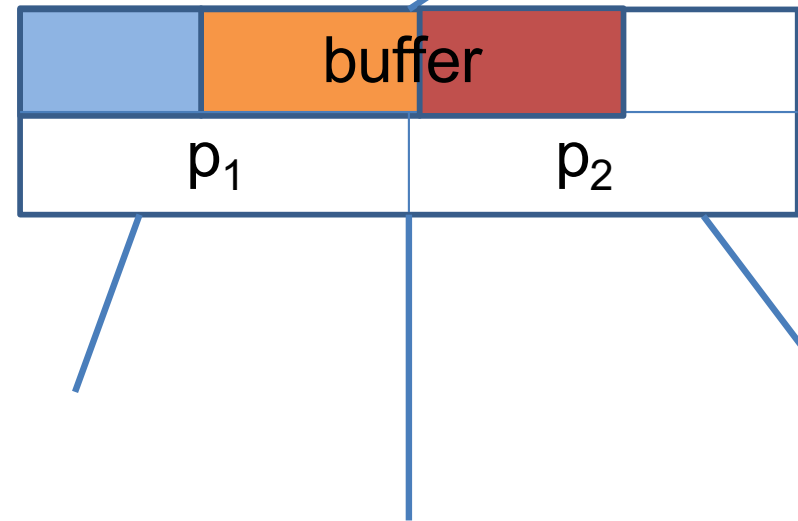
Buffer Tree

Summary

Cost of operations:

insert/delete: $O\left(\frac{1}{B} \log n\right)$

search: $O(\log n)$



Faster Trees?

Goal:

A external memory data structure with fast searches, and super-fast insertions/deletions.

“Write-optimized data structure.”

Examples:

- LSM: log-structured merge trees
- COLA: cost-oblivious lookahead array

Hot area of DBS research today...

See:
BetrFS

Topic 4: Hash Tables

Two key types of hash tables:

- Chaining
- Open Addressing

Hash Sets / Filters:

- Fingerprint Hash Tables
- Bloom Filters

Better Collision Resolution


Chaining:

- $O(1)$ *expected* search
- $O(1)$ *worst-case* insertion

Cuckoo Hashing:

- $O(1)$ *worst-case* search
- $O(1)$ *expected* insertion

Neat, newer
hashing method!



More realistic hash functions

How well does linear probing really work?

Does open addressing work with realistic hash functions?

- Example: limited independence hash functions
- Example: tabular hashing

YES: it works really well!

Better hash functions

Faster hash functions:

- Example: xxHash, etc.
- Example: tabular hashing

*Fastest today??
For GPUs??*

Cryptographic hash functions:

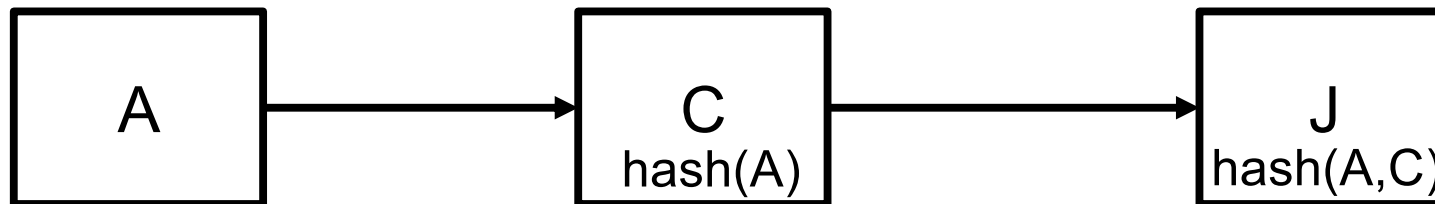
- Example: MD6,
- Example: SHA-512

Are these secure?

Blockchains

Blockchain = Hashchain

- proof-of-work == inverting hash function
- hash summarizes chain (see: Merkle trees!)



Topic 4: Hash Tables

Two key types of hash tables:

- Chaining
- Open Addressing

Hash Sets / Filters:

- Fingerprint Hash Tables
- Bloom Filters

Better filters

Quotient Filters:

- Optimal trade-off error vs. space.
- Practical and easy to implement

Cuckoo Filters:

- Interesting alternative to a Bloom Filter

Optimizations: Bloomier filters, compact approximators...

Key question: minimize space, minimize error

Applications of Filters

Caches:

- What is in your cache?
- Check bloom filters before accessing cache.

Learned Bloom Filters:

- Use machine learning to choose most useful items to store in filter.
- Then we need less space!

See: <https://papers.nips.cc/paper/7328-a-model-for-learned-bloom-filters-and-optimizing-by-sandwiching.pdf>

Topic 5: Shortest Paths

Several situations:

- Unweighted graphs
- Directed acyclic graphs
- Graphs with negative weights
- Graph with positive weights
- All-pairs shortest paths

Topic 5: Shortest Paths

Key algorithms:

- BFS
- Topological sort
- Bellman-Ford
- Dijkstra
- Floyd-Warshall

Topic 6: Minimum Spanning Trees

Key algorithms:

- Prim's
- Kruskal's
- Boruvka's

Sub-components:

- Union-Find
- Priority Queues

Key Questions

Big Data

Parallelism

Distributed Network Applications

Parallel Algorithms

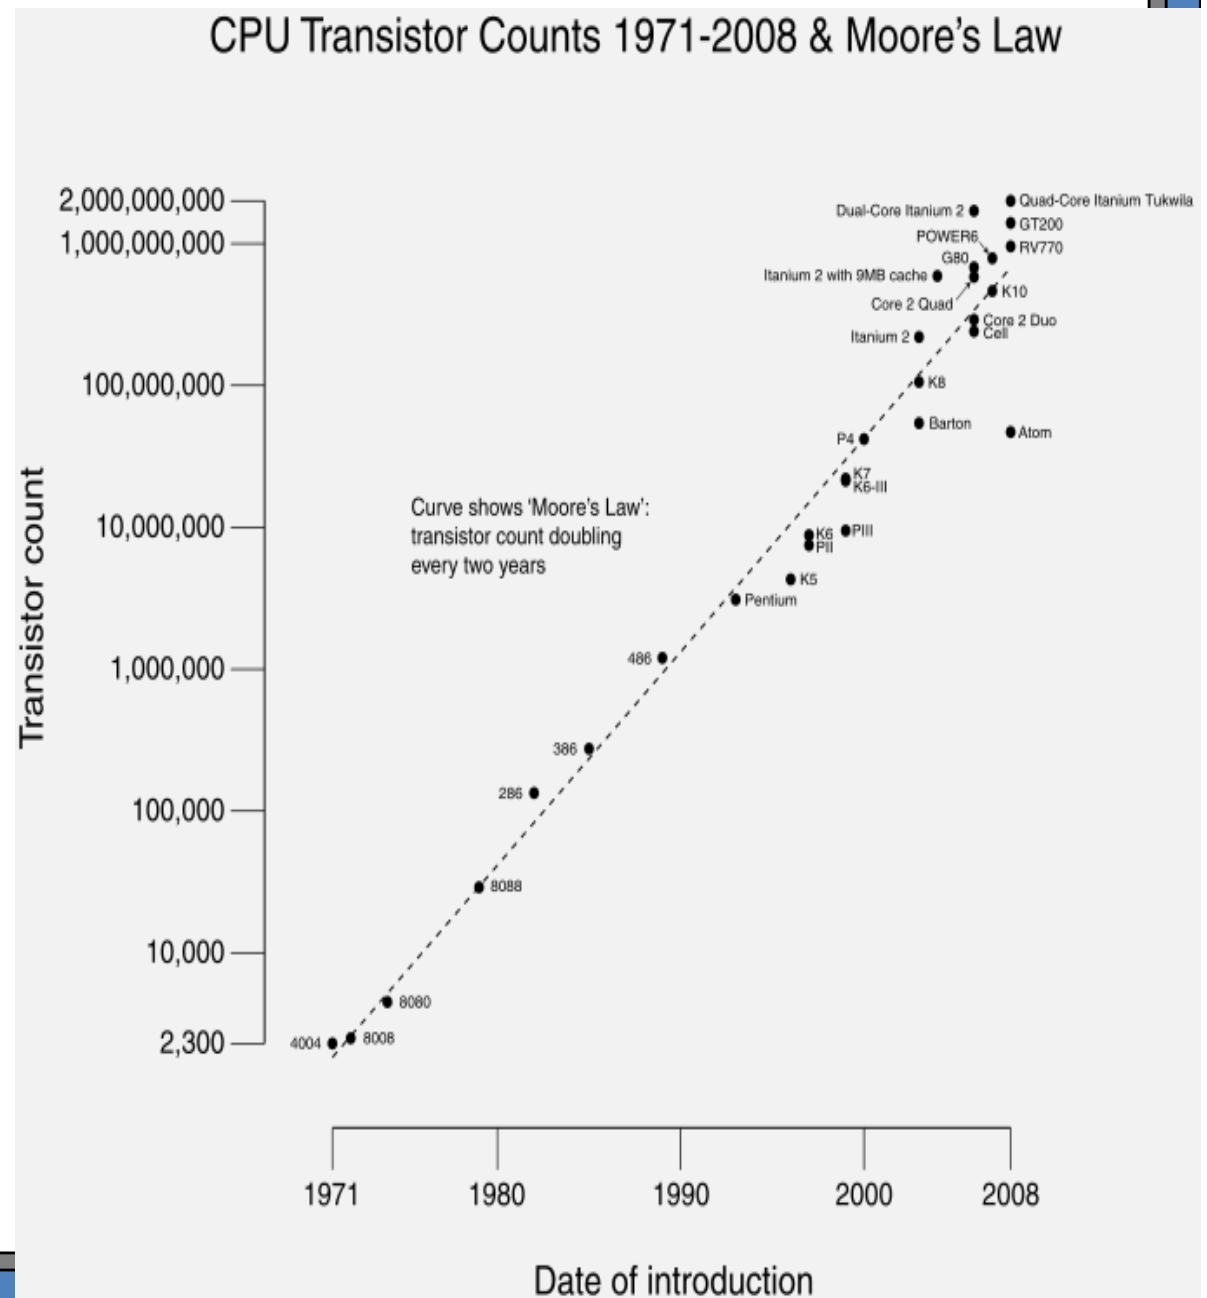
Moore's Law

Number of transistors
doubles every 2 years!

“The complexity for minimum component costs has increased at a rate of roughly a factor of two per year... Certainly over the short term this rate can be expected to continue, if not to increase.” Gordon Moore, 1965

Limits will be reached
in 10-20 years...maybe.

Source: Wikipedia



Parallel Algorithms

More transistors == faster computers?

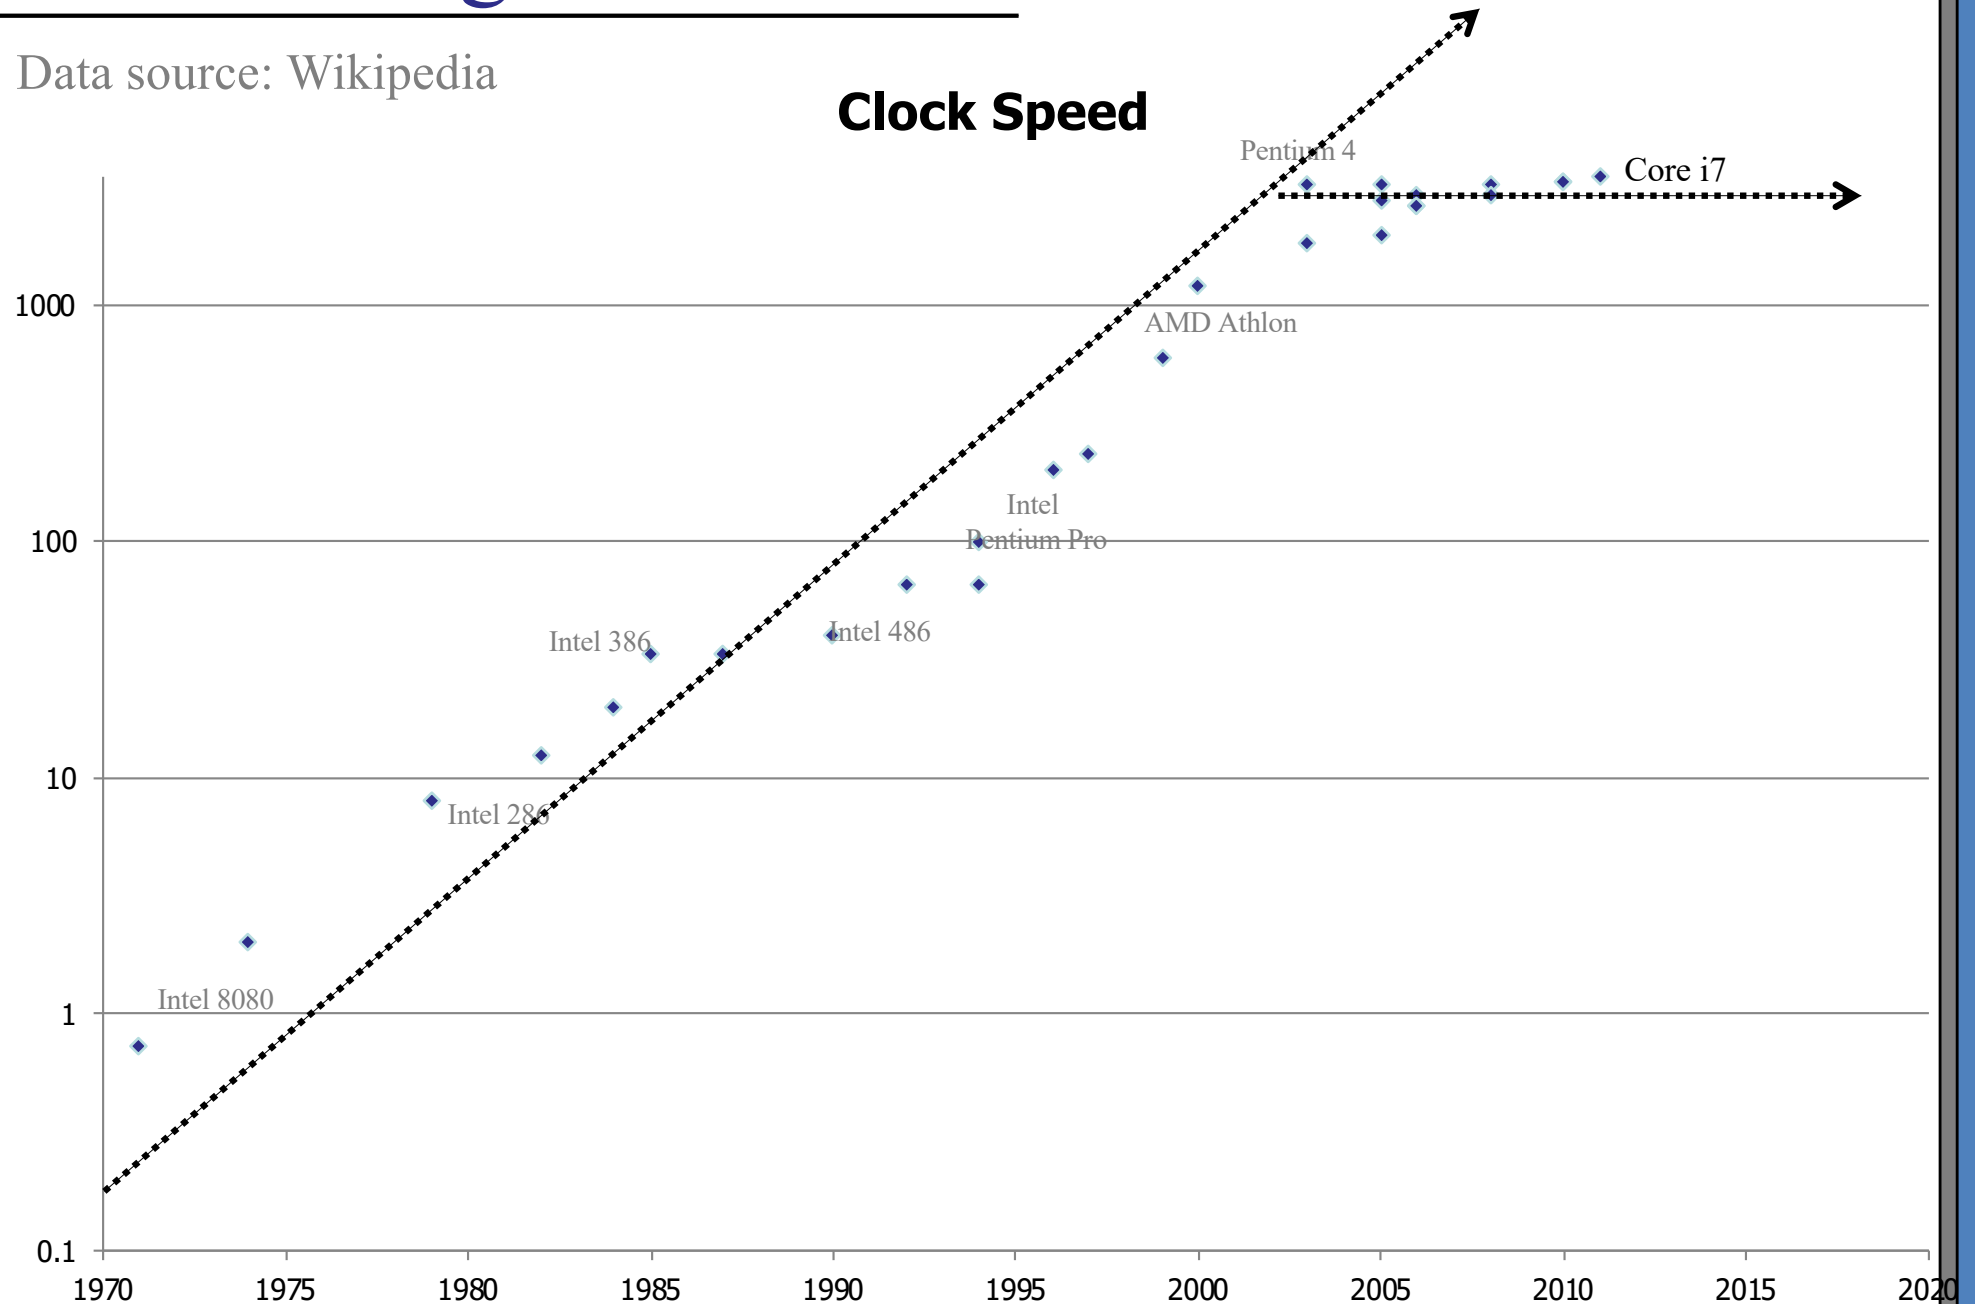
- More transistors per chip → smaller transistors.
- Smaller transistors → faster
- Conclusion:

Clock speed doubles every two years, also.

Parallel Algorithms

Data source: Wikipedia

Clock Speed



Parallel Algorithms

What to do with more transistors?

- More functionality
 - GPUs, FPUs, specialized crypto hardware, etc.
- Deeper pipelines
- More clever instruction issue (out-of-order issue, scoreboarding, etc.)
- More on chip memory (cache)

Limits for making faster processors?

Parallel Algorithms

Problems with faster clock speeds:

- Heat

- Faster switching creates more heat.

- Wires

- Adding more components takes more wires to connect.
- Wires don't scale well!

- Clock synchronization

- How do you keep the entire chip synchronized?
- If the clock is too fast, then the time it takes to propagate a clock signal from one edge to the other matters!

Parallel Algorithms

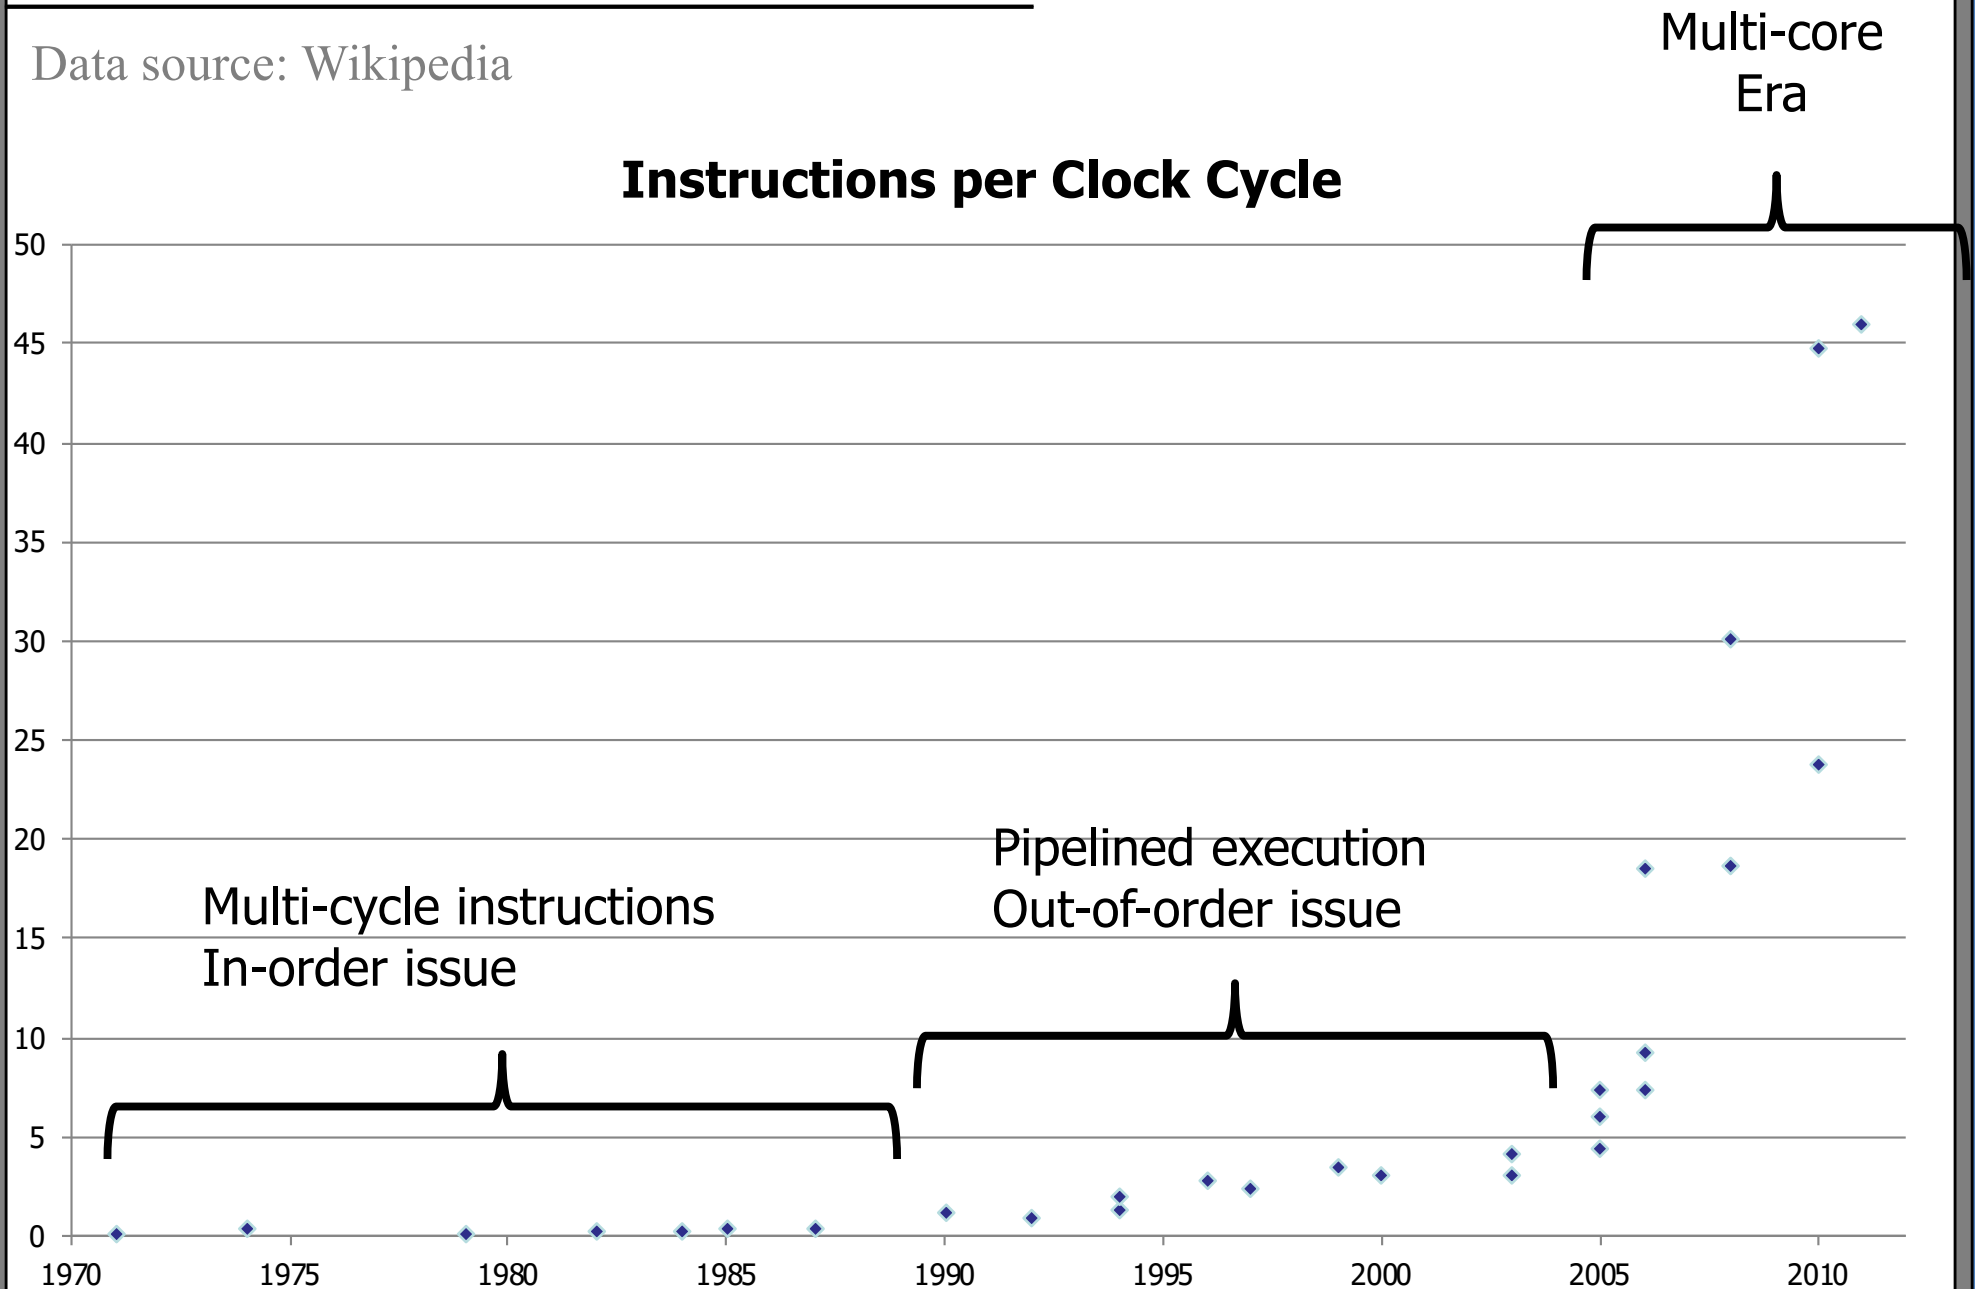
Conclusion:

- We have lots of new transistors to use.
- We can't use them to make the CPU faster.

What do we do?

Parallel Algorithms

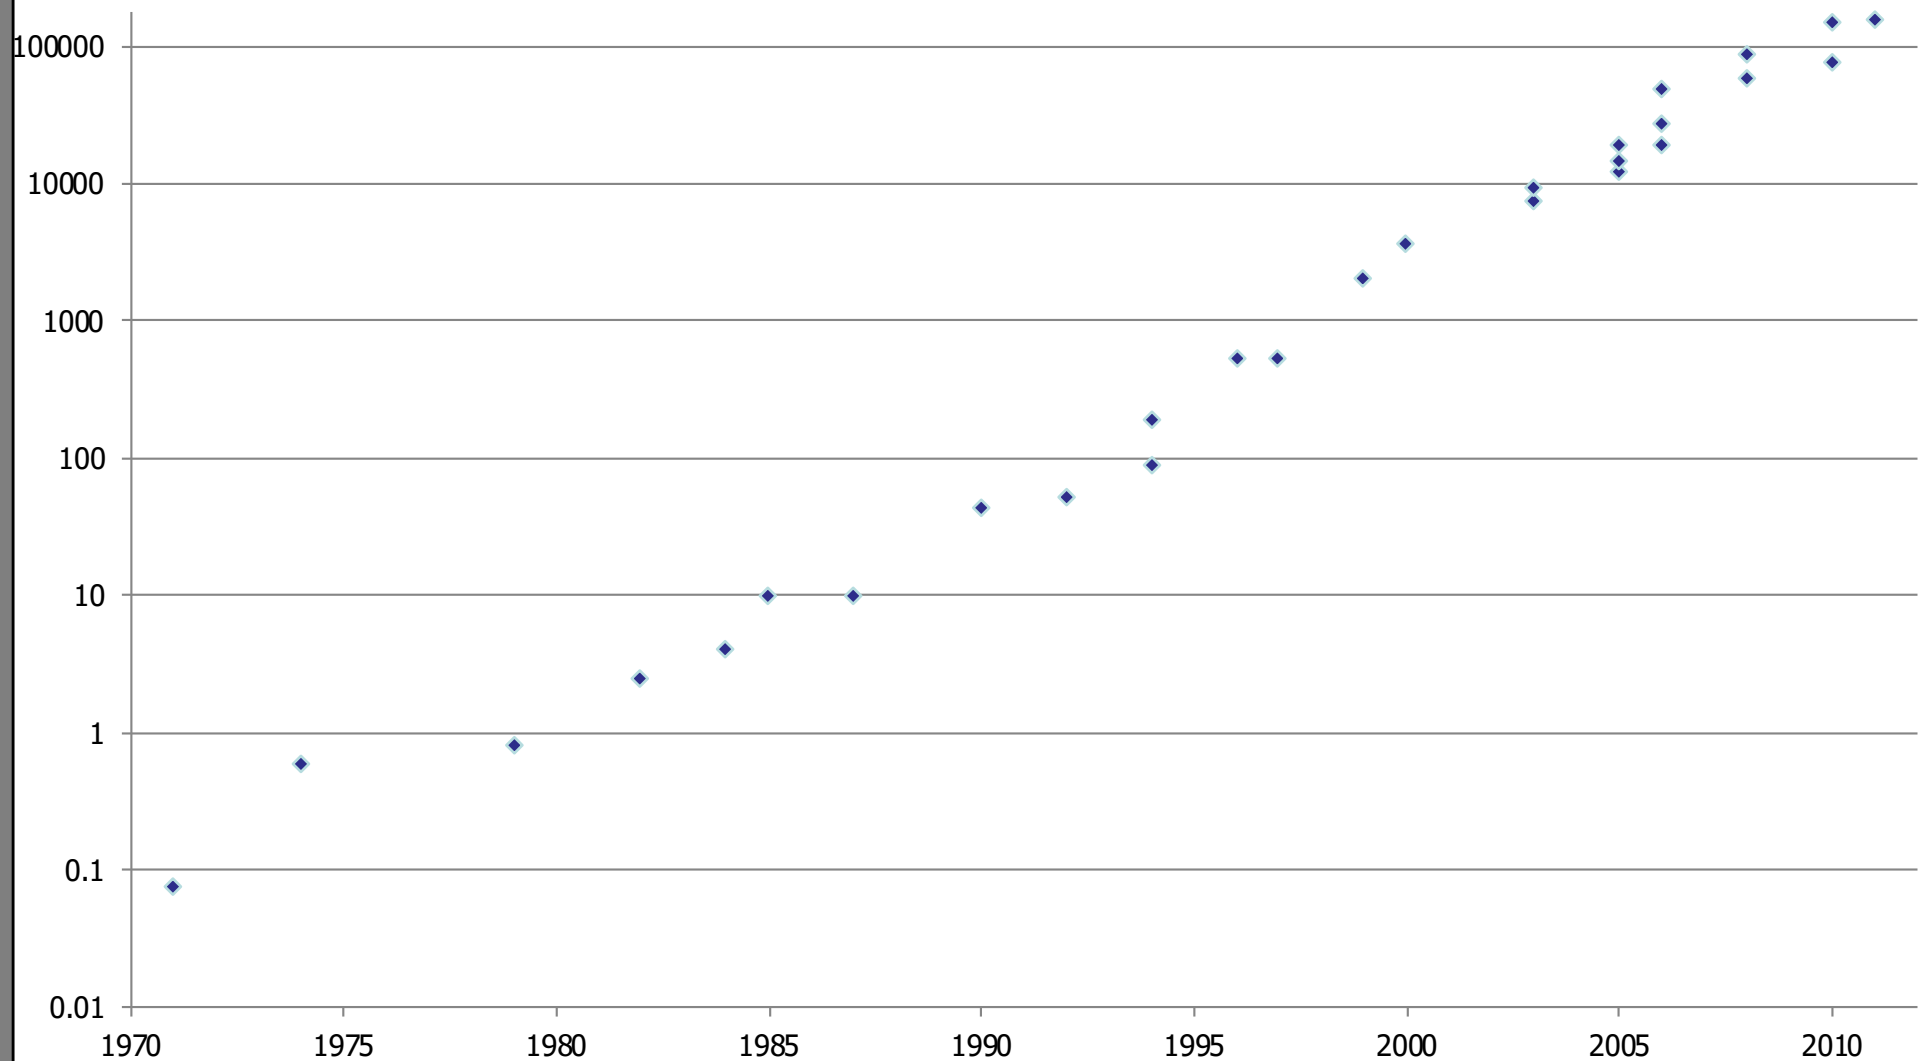
Data source: Wikipedia



Parallel Algorithms

Data source: Wikipedia

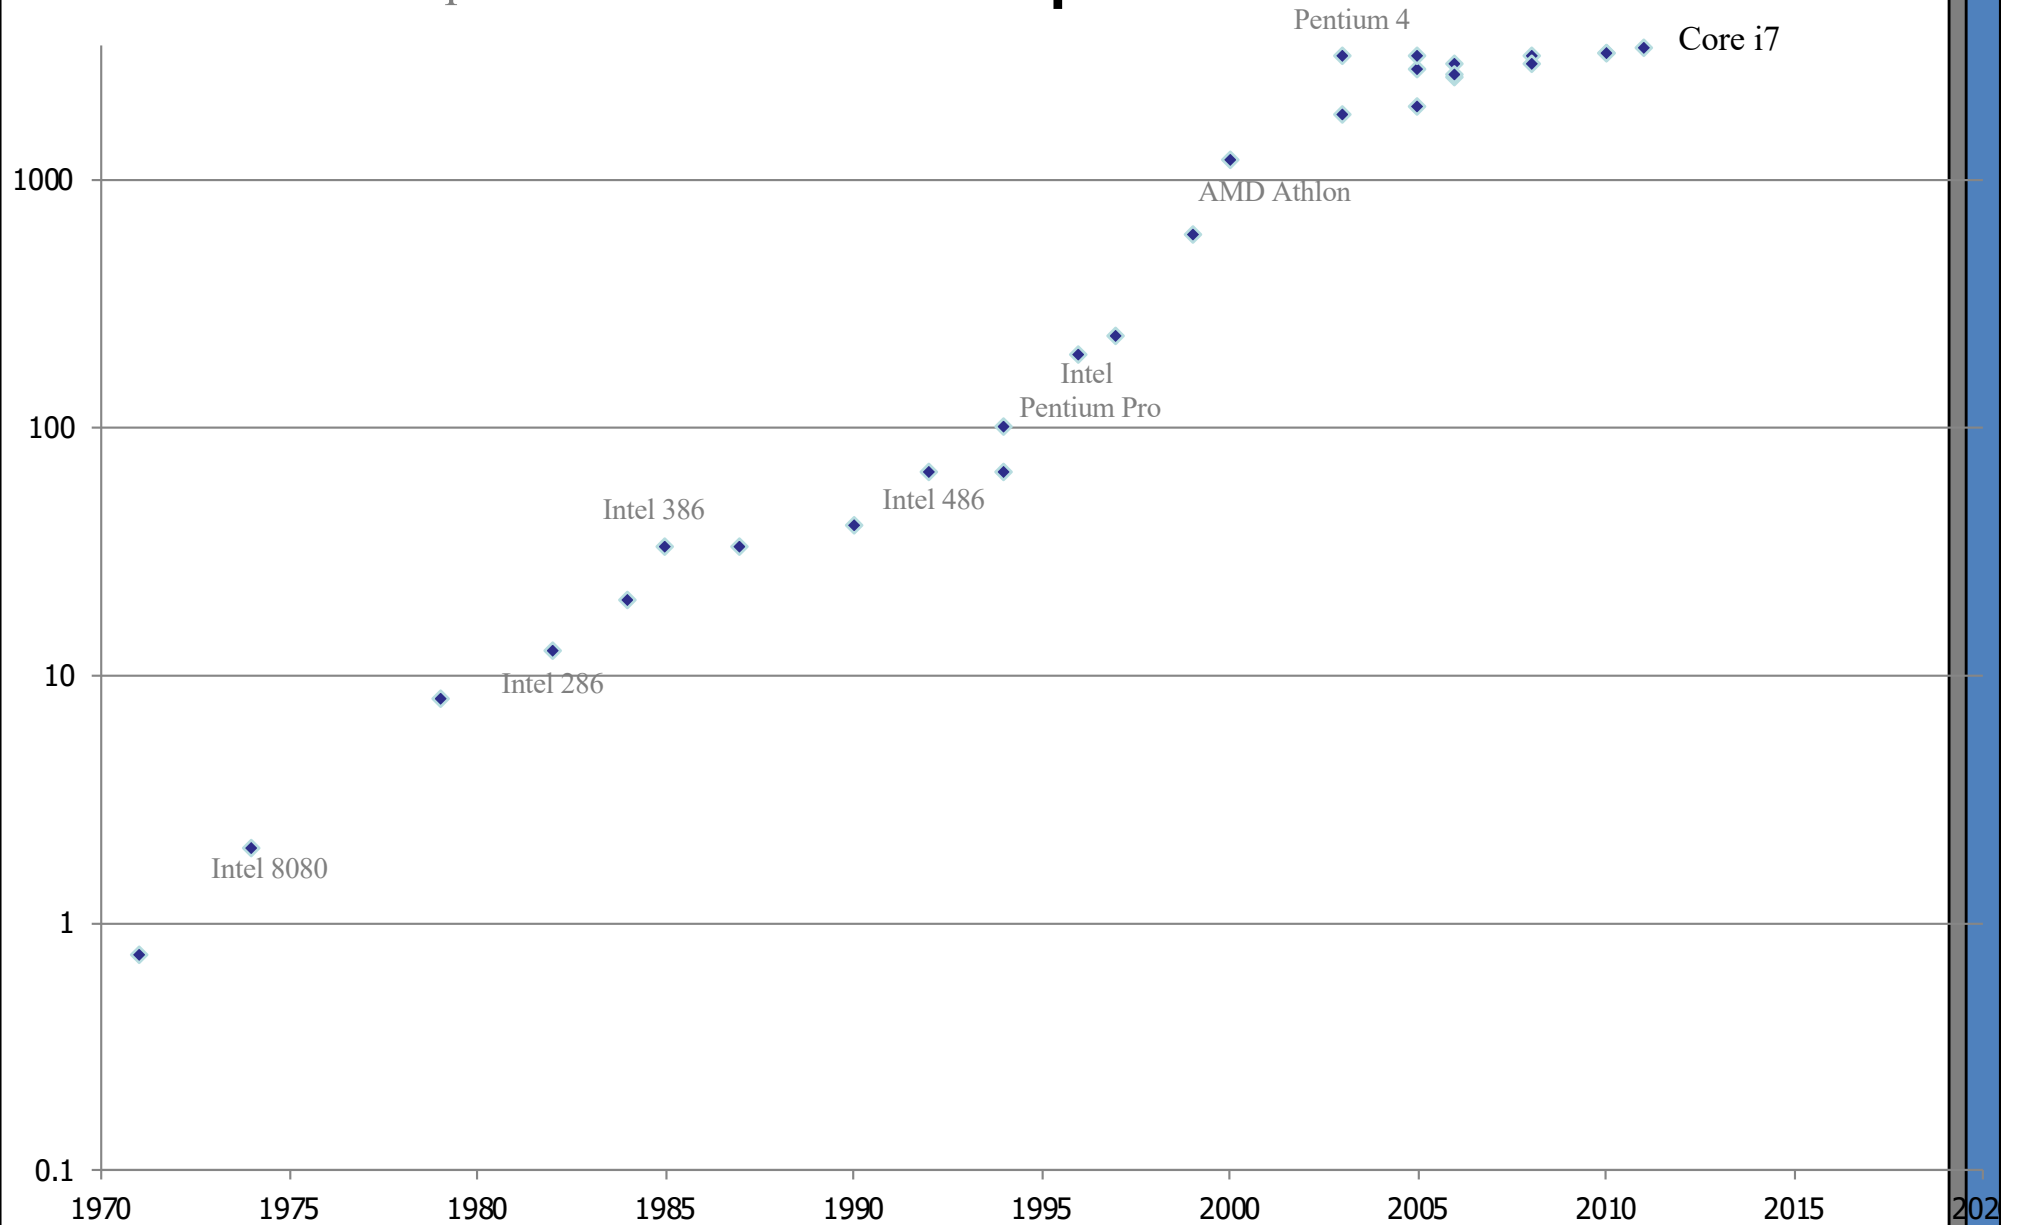
Instructions per Second



Parallel Algorithms

Data source: Wikipedia

Clock Speed



Parallel Algorithms

To make an algorithm run faster:

- Must take advantage of multiple cores.
- Many steps executed at the same time!

Parallel Algorithms

Different types of parallelism:

- multicore
 - on-chip parallelism: synchronized, shared caches, etc.
- multsocket
 - closely coupled, highly synchronized, shared caches
- cluster / data center
 - connected by a high-performance interconnect
- distributed networks
 - slower interconnect, less tightly synchronized

Parallel Algorithms

Map-Reduce:

- Target: high-performance clusters
- Focus: data (not computation)

Inventor: Google

- processing web data

Today: ubiquitous (Amazon, Yahoo, Facebook, etc.,)

- Hadoop, etc.

Map-Reduce Model

Basic round:

1. **Map**: process each (key, value) pair
2. **Shuffle**: group items by key
3. **Reduce**: process items with same key together

Plan:

Load data from disk.

Execute several rounds.

Save (key, value) pairs, sorted by key.

Map-Reduce

Boruvka's Algorithm

Repeat:

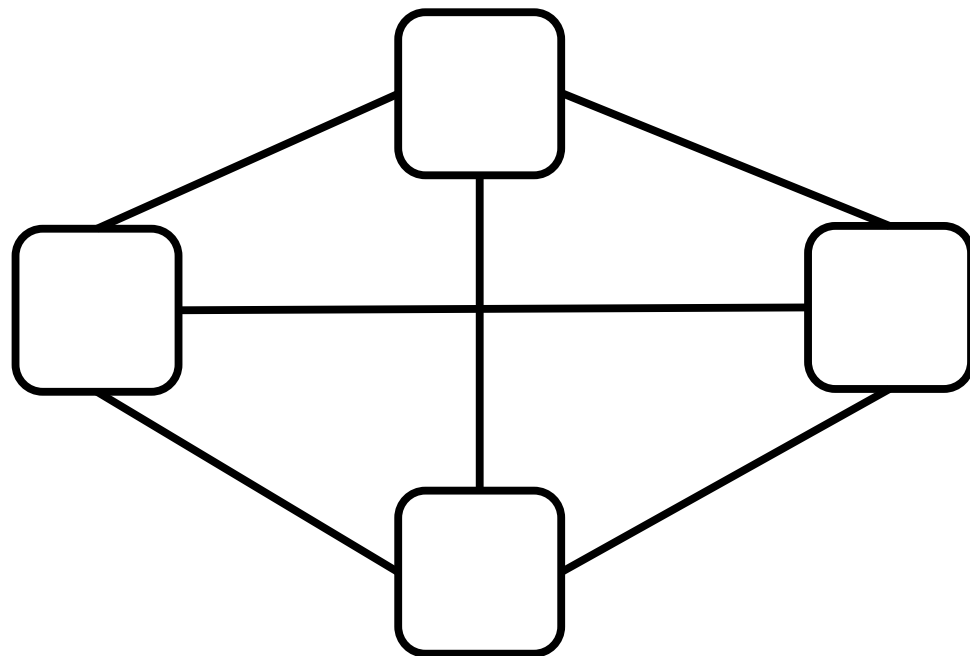
1. Find min-weight-outgoing-edge for each node.
2. Merge components.

Time: $O(\log n)$ rounds of map-reduce

Parallel Algorithms

k-Machine Cluster:

- **k servers**: system is a collection of cores/CPU's/etc.
- **all-to-all communication**: communicate via messages
- **bandwidth limit B**: limited data transfer



k = 4

k-Machine Model

Boruvka's Algorithm

Repeat:

1. Find min-weight-outgoing-edge for each node.
2. Merge components.

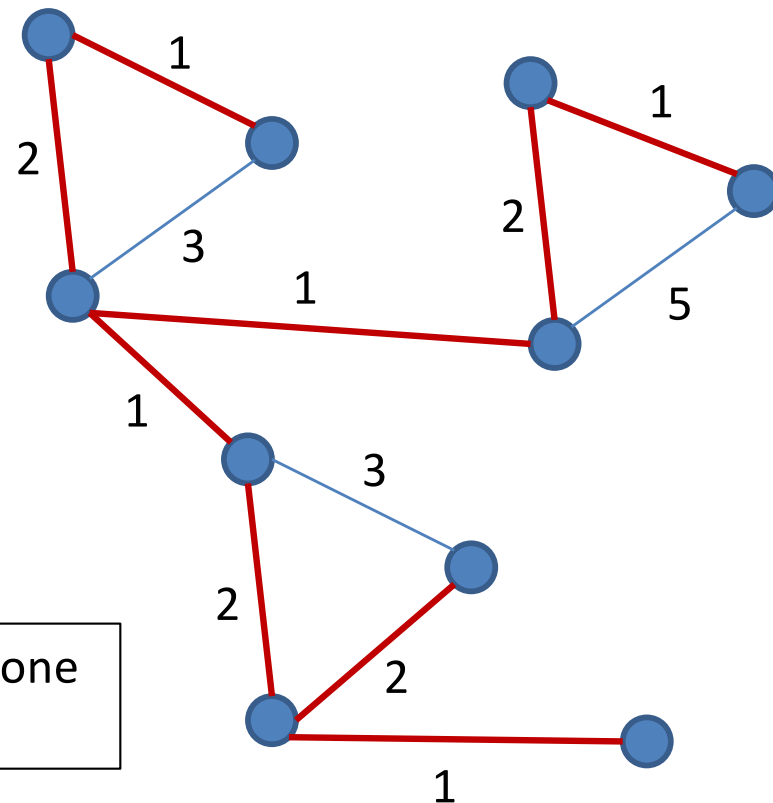
Time: $O(n \log^c(n)/k^2)$

Fully Distributed Model

Network Model

Assume each node in the graph is its own machine.

Each edge in the graph is a real communication edge.



Cannot send message to everyone
like in k-machine model.

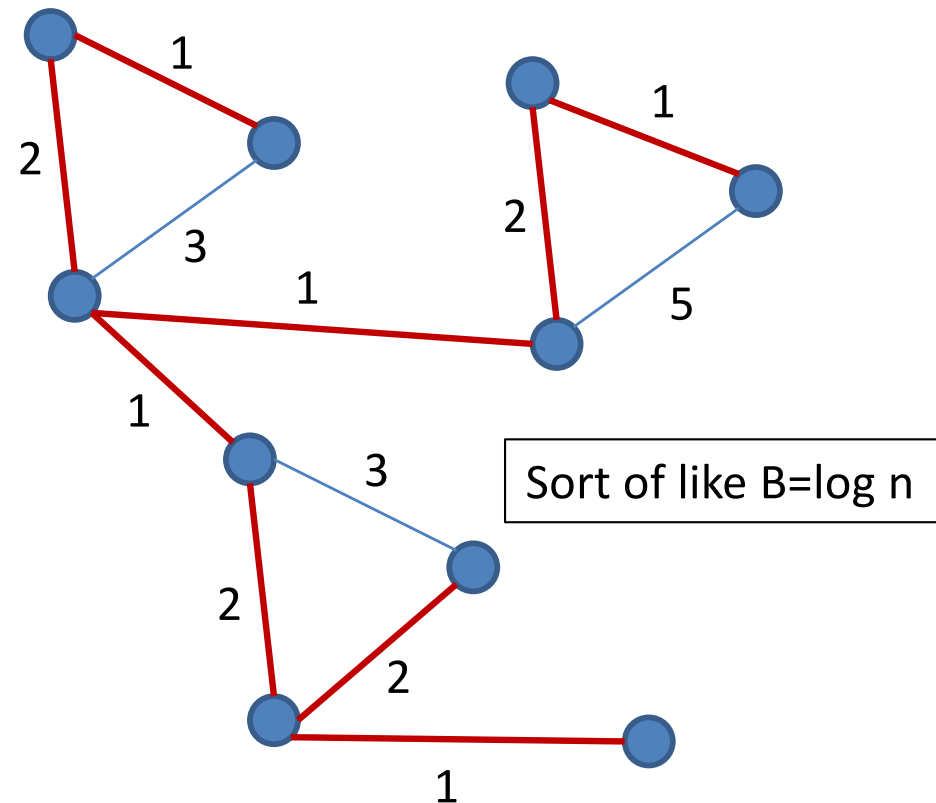
Fully Distributed Model

Network Model

Assume each node in the graph is its own machine.

Each edge in the graph is a real communication edge.

Each edge carries 1 message per round.



Fully Distributed Model

Boruvka's Algorithm

Repeat:

1. Find min-weight-outgoing-edge for each node.
2. If component is $< n^{1/2}$ then aggregate MWOE in component. Otherwise aggregate on BFS tree.
3. Merge components.

Time: $O((D + n^{1/2})\log n)$

Parallel Algorithms

A huge amount of ongoing research on how to process big graphs fast in parallel...

A huge amount of ongoing research on how to efficiently find shortest paths, spanners, MSTs, and more in a network...

What's next?

Topic 1: Searching and Optimization

Topic 2: Sorting

Topic 3: Trees

Topic 4: Hashing

Topic 5: Shortest Paths

Topic 6: Minimum Spanning Trees

What next?

Algorithms modules:

- CS3230: Design and Analysis of Algorithms
- CS3233: Competitive Programming
- CS4231: Parallel and Distributed Algorithms
- CS4234: Optimization Algorithms
- CS5234: Combinatorial and Graph Algorithms
- CS5237: Computational Geometry
- CS5330: Randomized Algorithms

Theory:

- CS4232: Theory of Computation
- CS5230: Computational Complexity

What next?

Software engineering modules:

- CS2103: Software engineering
- CS4211: Formal methods for software engineering
- CS4218: Software testing and debugging

System design and programming modules:

- CS3216: Software development on evolving platforms
- CS3217: Software engineering on modern application platforms

What next?

Specialized modules:

- Distributed Systems
- Computer Security
- Game Design
- Computer Graphics
- Machine Learning
- Computational Biology
- Wireless computing and sensor networks
- Etc...

The End

Goals:

- Apply algorithmic thinking and techniques.
- Describe the structure and operation of different data structures.
- Assess the suitability of different data-structures and algorithms.
- Adapt existing data-structures and algorithms.

The End

Goals:

- Apply algorithmic thinking and techniques.
- Describe the structure and operation of different data structures.
- Assess the suitability of different data-structures and algorithms.
- Adapt existing data-structures and algorithms.

I hope everyone has had fun...

The End

Goals:

- Apply algorithmic thinking and techniques.
- Describe the structure and operation of different data structures.
- Assess the suitability of different data-structures and algorithms.
- Adapt existing data-structures and algorithms.

I hope everyone has had fun...

If you ever want to chat about algorithms, just drop by my ~~office~~ zoom meeting...

The End

Goals:

- Apply algorithmic thinking and techniques.
- Describe the structure and operation of different data structures.
- Assess the suitability of different data-structures and algorithms.
- Adapt existing data-structures and algorithms.

I hope everyone has had fun...

If you ever want to chat about algorithms, just drop by my ~~office~~ zoom meeting...

And it's over... congratulations!