

CS2040S

Data Structures and Algorithms
(e-learning edition)

Welcome!

Announcements

Midterm : Monday March 9

- Location TBA (MPSH)
- < 50 students / room
- 9 rooms



Bring to quiz:

- One sheet of paper with any notes you like.
- Pens/pencils.
- You may not use anything else.

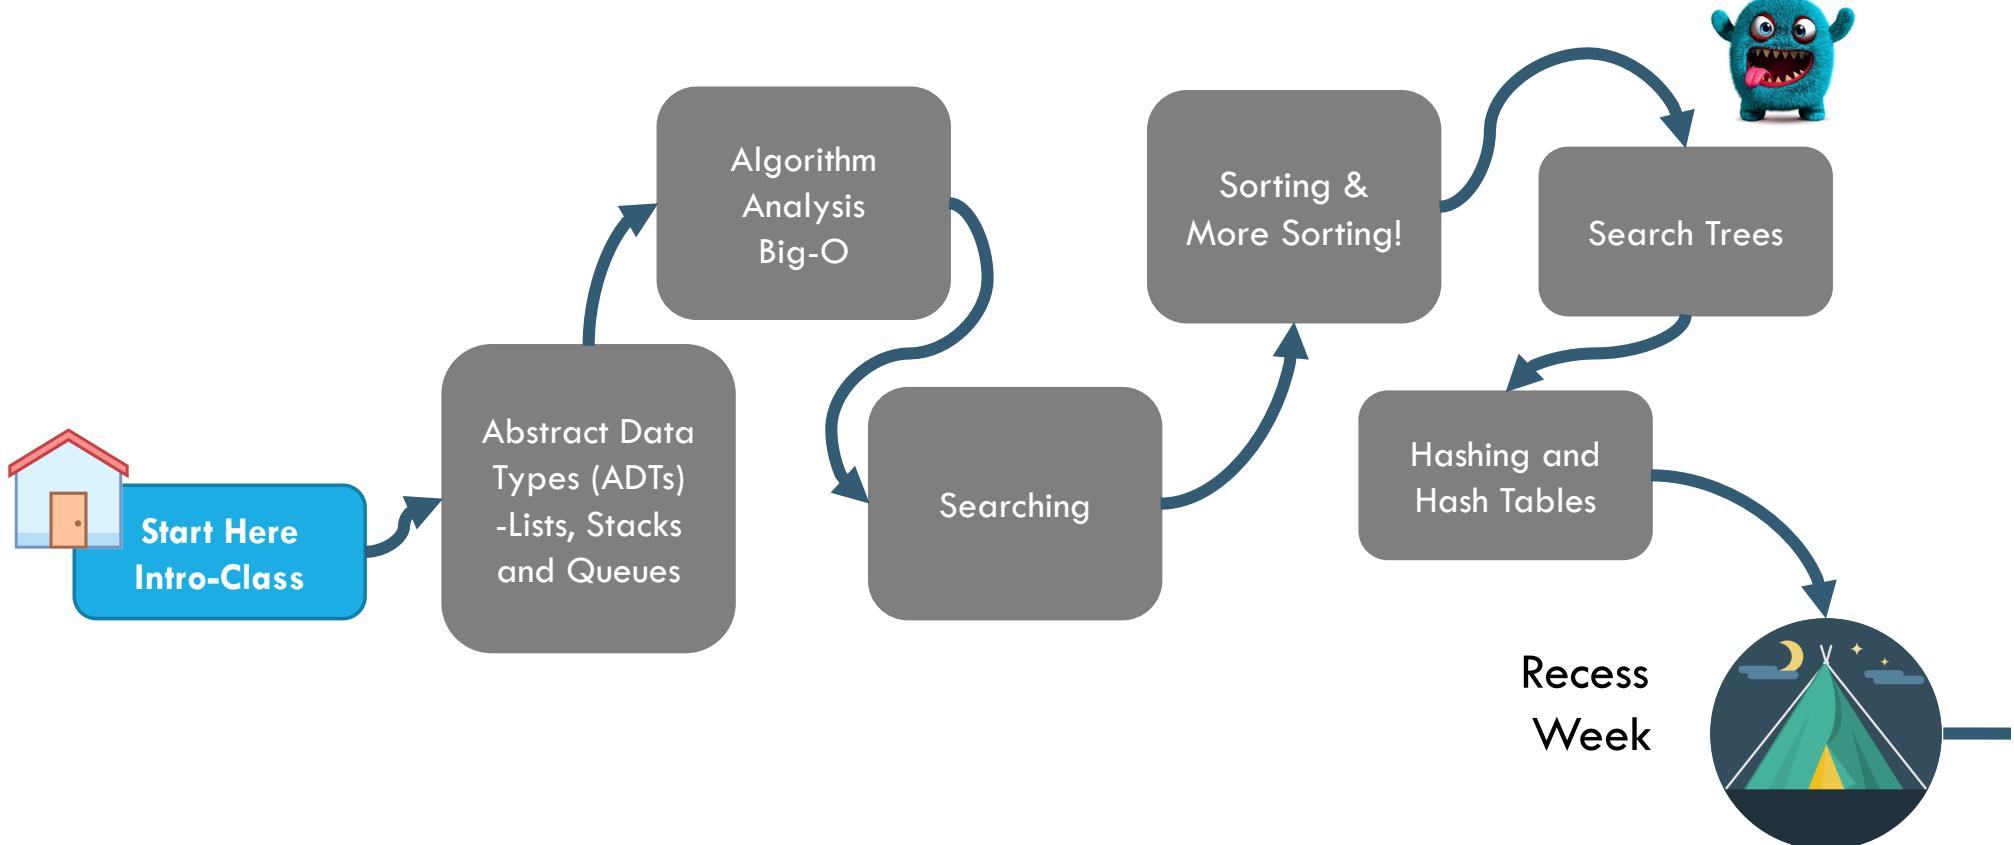
Announcements

Practice Quizzes (from 2015-17):

- Posted on Coursemology
- From CS2020
 - Beware: CS2020 != CS2040S
 - Different subset of material
- Discussed in tutorials



What have we done so far?



Part I: Organizing your Data

So much material?

Range trees
AVL
B-trees
SelectionSort
PancakeSort
MergeSort
Binary tree
SkipLists
Tries
Newton's Method
Interval trees
Chaining
Order statistics
Binary Search
BubbleSort
InsertionSort
ReversalSort
QuickSelect
QuickSort
BogoSort
CardFlipTrees
Hash table
Gradient descent
Random Permutations

So much material?

So much material?

PancakeSort

MergeSort

SkipLists

Binary tree

Tries

QuickSelect

BogoSort
CardFlipTrees
Hash table

QuickSort

Gradient descent

Random Permutations

ReversalSort

InsertionSort

BubbleSort

Order statistics

Binary Search

Chaining

Interval trees

Newton's Method

AVL

B-trees

SelectionSort

Range trees

So much material?

Range trees

Interval trees

Chaining

Order statistics

AVL

Newton's Method

Binary Search

B-trees

BubbleSort

SelectionSort

InsertionSort

PancakeSort

QuickSelect

ReversalSort

MergeSort

QuickSort

Gradient descent

Binary tree

BogoSort

CardFlipTrees

Random Permutations

SkipLists

Tries

Hash table

So much material?

Three Parts

1) Basic Theory / Algorithm Analysis

2) Major Algorithms

3) Overarching Principles

Warning:

Today I am going to review some
of the most important topics.

This will not be comprehensive.

1) Basic Theory / Algorithm Analysis

Key Topics

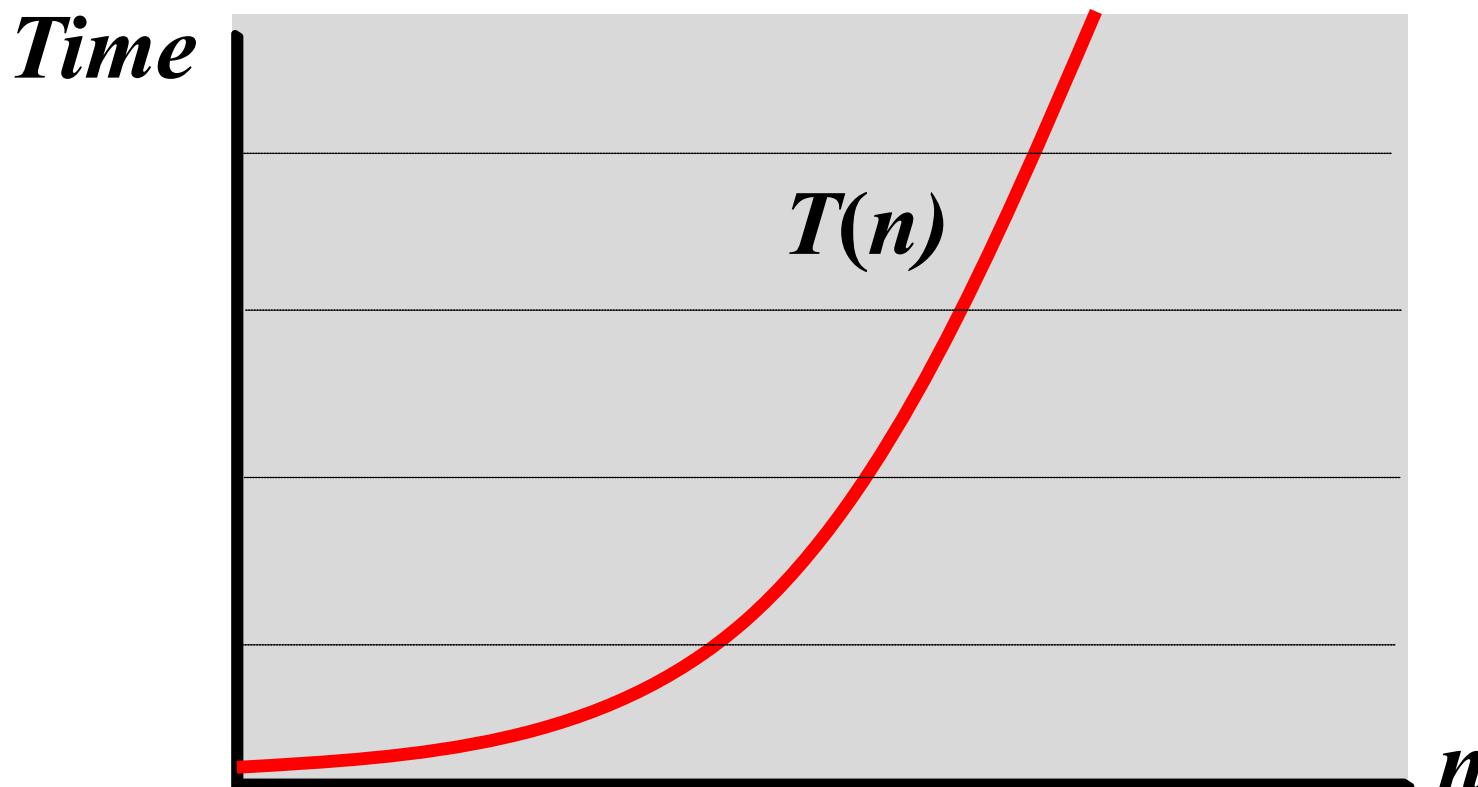


- Asymptotic notation (big-O, etc.)
- Simple recurrences
- Asymptotic analysis
- Basic probability (e.g., CS1231 review)

Big-O Notation

How does an algorithm scale?

- For large inputs, what is the running time?
- $T(n)$ = running time on inputs of size n



Big-O Notation

Definition: $T(n) = O(f(n))$ if T grows no faster than f

$T(n) = O(f(n))$ if:

- there exists a constant $c > 0$
- there exists a constant $n_0 > 0$

such that for all $n > n_0$:

$T(n) \leq c f(n)$

Example

$T(n)$	$f(n)$	big-O
$T(n) = 1000n$	$f(n) = n$	$T(n) = O(n)$
$T(n) = 1000n$	$f(n) = n^2$	$T(n) = O(n^2)$
$T(n) = n^2$	$f(n) = n$	$T(n) \neq O(n)$
$T(n) = 13n^2 + n$	$f(n) = n^2$	$T(n) = O(n^2)$

Simple recurrences

Recurrences we have seen a lot:

$$T(n) = 2T(n/2) + n$$

$$T(n) = 2T(n/2) + 1$$

$$T(n) = T(n/2) + 1$$

$$T(n) = T(n/2) + n$$

Et cetera

Simple recurrences

Recurrences we have seen a lot:

$$T(n) = 2T(n/2) + n$$

$$T(n) = 2T(n/2) + 1$$

$$T(n) = T(n/2) + 1$$

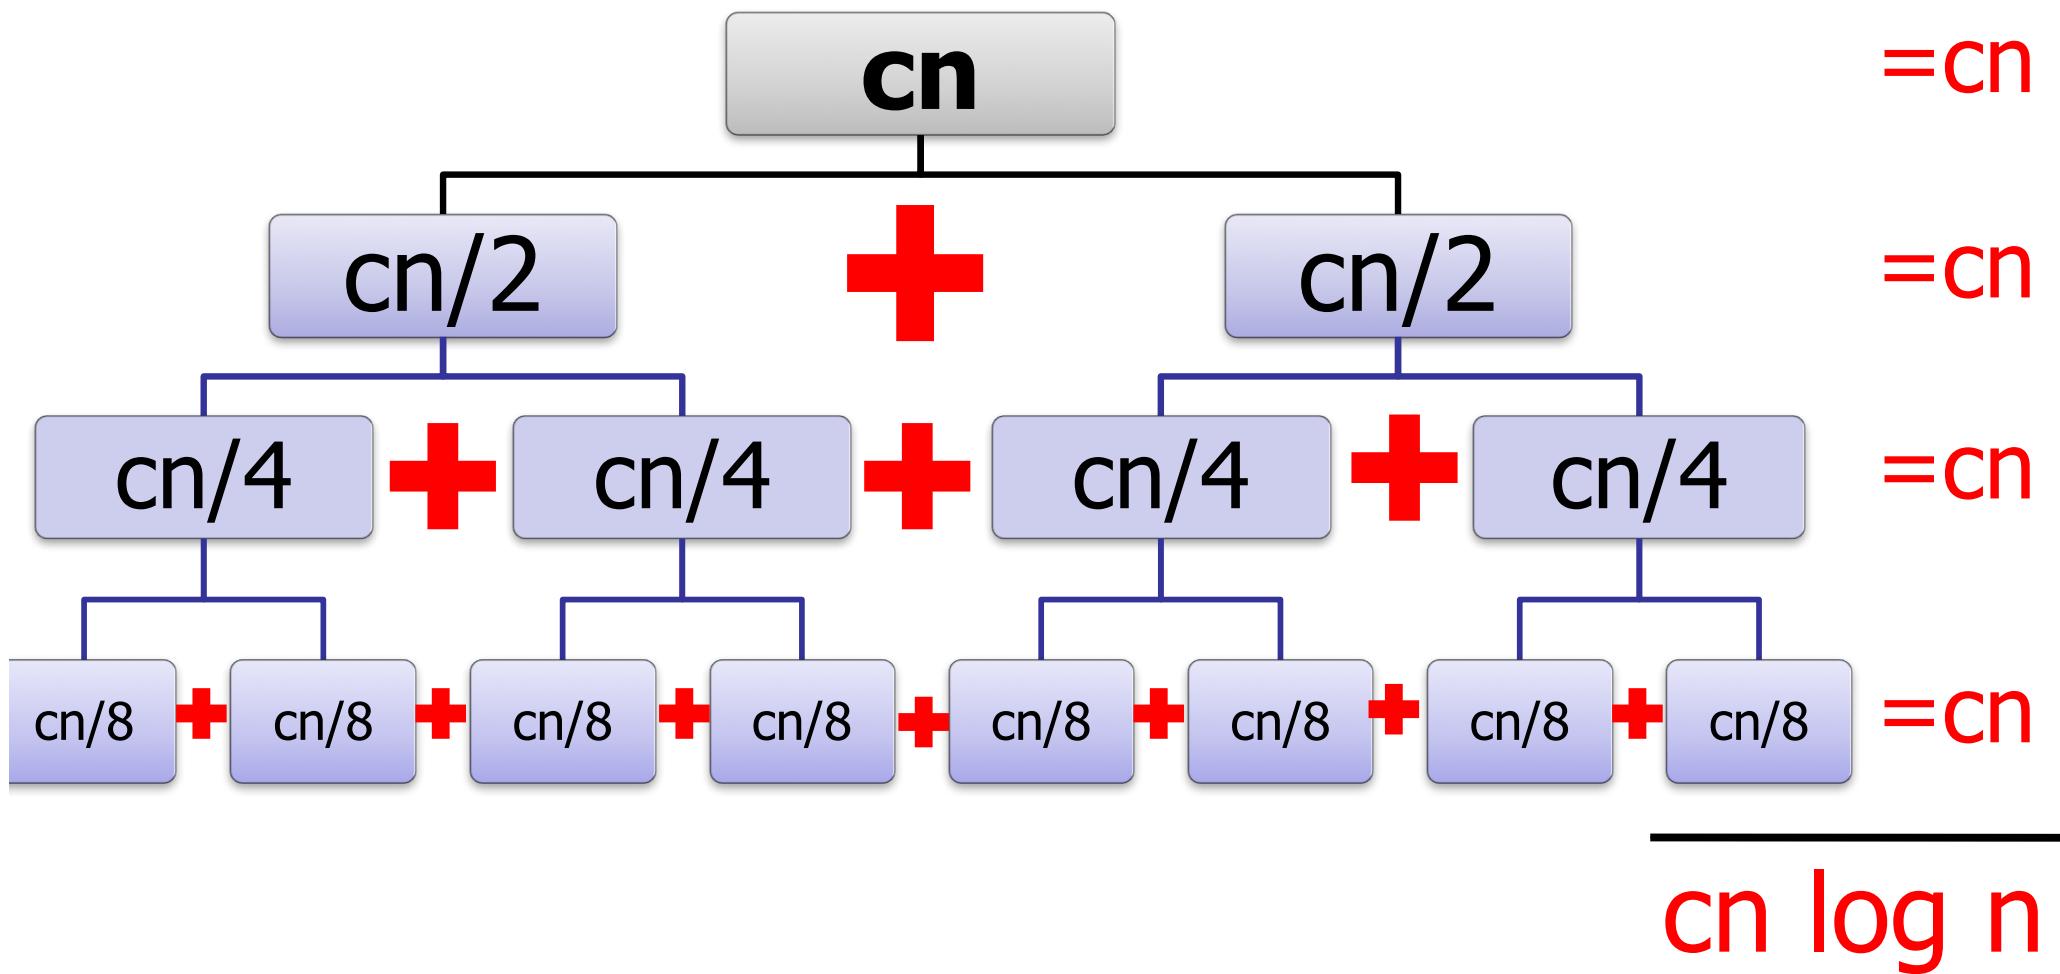
$$T(n) = T(n/2) - 1$$

Et cetera

1. Know the solution.
2. Be able to draw the recurrence tree.
3. Be able to substitute.

Example: MergeSort Analysis

$$T(n) = 2T(n/2) + cn$$



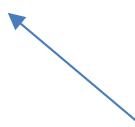
Guess: $T(n) = c \cdot n \log n$

$$T(1) = c$$

$$T(x) = c \cdot x \log x \text{ for all } x < n.$$

$$\begin{aligned} T(n) &= 2T(n/2) + cn \\ &= 2(c(n/2) \log(n/2)) + cn \\ &= cn \log(n/2) + cn \\ &= cn \log(n) - cn \log(2) + cn \\ &= cn \log(n) \end{aligned}$$

Induction:
It works!



Recurrence being analyzed:

$$\begin{aligned} T(n) &= 2T(n/2) + c \cdot n \\ T(1) &= c \end{aligned}$$

Algorithm Analysis

Asymptotics?

```
void pushAll(int k) {  
    for (int i=0;  
         i<= 100*k;  
         i++)  
    {  
        stack.push(i);  
    }  
}
```

$O(k)$

```
void pushAdd(int k) {  
    for (int i=0; i<= k; i++)  
    {  
        for (int j=0; j<= k; j++) {  
            stack.push(i+j);  
        }  
    }  
}
```

$O(k^2)$

Asymptotics on problem sets

Many questions asked:

What is the asymptotic running time of your program?

Many got that wrong (even when their program was right).

Algorithm Analysis

Example:

```
void sum(int k, int[] intArray) {  
    int total=0;  
    for (int i=0; i<= k; i++) {  
        total = total + intArray[i];  
    }  
    return total;  
}
```

$$\text{Total: } 1 + 1 + (k+1) + 3k + 1 = 4k+4 = O(k)$$

Rules

Loops

cost = (# iterations)x(max cost of one iteration)

```
int sum(int k, int[] intArray) {  
    int total=0;  
  
    for (int i=0; i<= k; i++) {  
        total = total + intArray[i];  
    }  
    return total;  
}
```



Rules

Nested Loops

cost = (# iterations)(max cost of one iteration)

```
int sum(int k, int[] intArray) {  
    int total=0;  
    for (int i=0; i<= k; i++) {  
        for (int j=0; j<= k; j++) {  
            total = total + intArray[i];  
        }  
    }  
    return total;
```



Rules

Sequential statements

cost = (cost of first) + (cost of second)

```
int sum(int k, int[] intArray) {  
    for (int i=0; i<= k; i++)  
        intArray[i] = k;  
    for (int j =0; j<= k; j++)  
        total = total + intArray[i];  
    return total;  
}
```

Rules

if / else statements

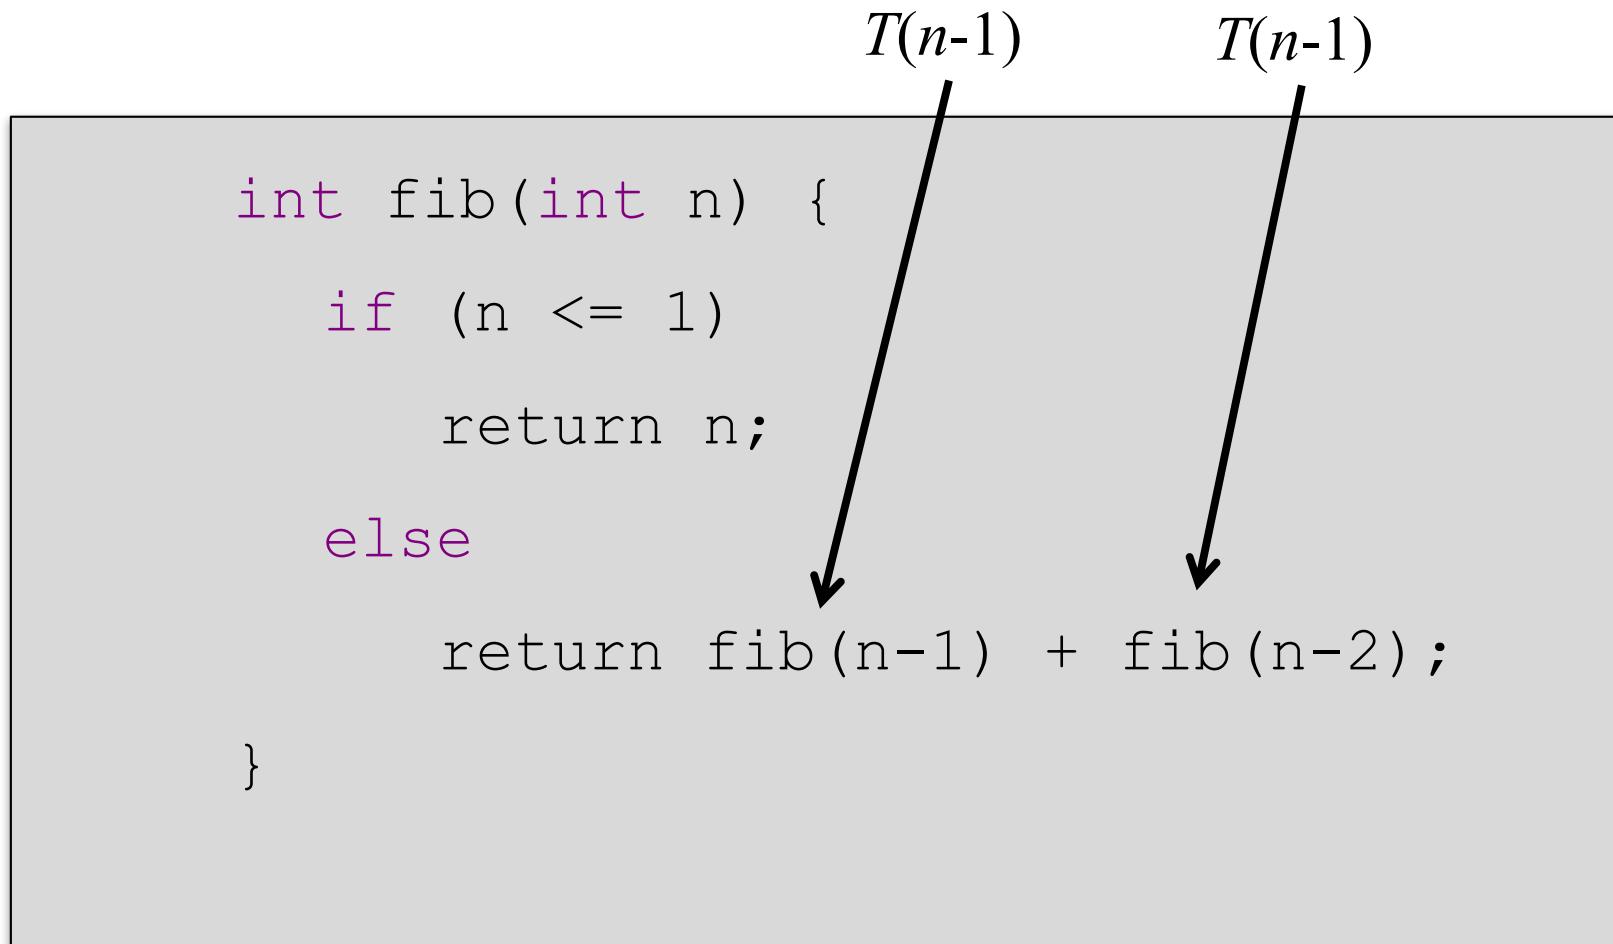
cost = 2max(cost of first, cost of second)

\geq (cost of first) + (cost of second)

```
void sum(int k, int[] intArray) {  
    if (k > 100)  
        doExpensiveOperation();  
    else  
        doCheapOperation();  
    return;  
}
```

Recurrences

$$\begin{aligned}T(n) &= 1 + T(n - 1) + T(n - 2) \\&= O(2^n)\end{aligned}$$



What is the running time?

```
for (int i = 0; i<n; i++)  
    for (int j = 0; j<i; j++)  
        store[i] = i + j;
```

Probability Theory (CS1231)

Probability Theory

Flipping a coin:

- $\Pr(\text{heads}) = \frac{1}{2}$
- $\Pr(\text{tails}) = \frac{1}{2}$

Coin flips are independent:

- $\Pr(\text{heads} \rightarrow \text{heads}) = \frac{1}{2} * \frac{1}{2} = \frac{1}{4}$
- $\Pr(\text{heads} \rightarrow \text{tails} \rightarrow \text{heads}) = \frac{1}{2} * \frac{1}{2} * \frac{1}{2} = \frac{1}{8}$

You flip a coin 8 times. Which is more likely?

- a. 4 heads, followed by 4 tails
- b. 8 heads in a row
- c. Same
- d. Incomparable

Probability Theory

Events **A**, **B**:

- $\Pr(\mathbf{A})$, $\Pr(\mathbf{B})$
- **A** and **B** are independent
(e.g., unrelated random coin flips)

Then:

- $\Pr(\mathbf{A} \text{ and } \mathbf{B}) = \Pr(\mathbf{A})\Pr(\mathbf{B})$

Probability Theory

Set of outcomes for $X = (e_1, e_2, e_3, \dots, e_k)$:

- $\Pr(e_1) = p_1$
- $\Pr(e_2) = p_2$
- ...
- $\Pr(e_k) = p_k$

Expected outcome:

$$E[X] = e_1 p_1 + e_2 p_2 + \dots + e_k p_k$$

Probability Theory

Expected value:

- Weighted average

Example: event **A** has two outcomes:

- $\Pr(\mathbf{A} = 12) = \frac{1}{4}$
- $\Pr(\mathbf{A} = 60) = \frac{3}{4}$

Expected value of A:

$$E[A] = (\frac{1}{4})12 + (\frac{3}{4})60 = 48$$

Probability Theory

Flipping a coin:

- $\Pr(\text{heads}) = \frac{1}{2}$
- $\Pr(\text{tails}) = \frac{1}{2}$

In two coin flips: I expect one heads.

Probability Theory

Define event **A**:

- **A** = number of heads in two coin flips

In two coin flips: I expect one heads.

- $\Pr(\text{heads, heads}) = \frac{1}{4}$ $2 * \frac{1}{4} = \frac{1}{2}$
 - $\Pr(\text{heads, tails}) = \frac{1}{4}$ $1 * \frac{1}{4} = \frac{1}{4}$
 - $\Pr(\text{tails, heads}) = \frac{1}{4}$ $1 * \frac{1}{4} = \frac{1}{4}$
 - $\Pr(\text{tails, tails}) = \frac{1}{4}$ $0 * \frac{1}{4} = 0$
-

Probability Theory

Linearity of Expectation:

- $E[A + B] = E[A] + E[B]$

Example:

- $A = \# \text{ heads in 2 coin flips}$: $E[A] = 1$
- $B = \# \text{ heads in 2 coin flips}$: $E[B] = 1$
- $A + B = \# \text{ heads in 4 coin flips}$

$$E[A+B] = E[A] + E[B] = 1 + 1 = 2$$

Probability Theory

Flipping an (unfair) coin:

- $\Pr(\text{heads}) = p$
- $\Pr(\text{tails}) = (1 - p)$

How many flips to get at least one head?

$$E[X] = (p)(1) + (1 - p) (1 + E[X])$$



How many more flips to get a head?

Idea: If I flip “tails,” the expected number of additional flips to get a “heads” is still $E[X]$!!

Probability Theory

Flipping an (unfair) coin:

- $\Pr(\text{heads}) = p$
- $\Pr(\text{tails}) = (1 - p)$

How many flips to get at least one head?

$$\begin{aligned}\mathbf{E}[X] &= (p)(1) + (1 - p) (1 + \mathbf{E}[X]) \\ &= p + 1 - p + 1\mathbf{E}[X] - p\mathbf{E}[X]\end{aligned}$$

$$p\mathbf{E}[X] = 1$$

$$\mathbf{E}[X] = 1/p$$

Probability Theory

Flipping an (unfair) coin:

- $\Pr(\text{heads}) = p$
- $\Pr(\text{tails}) = (1 - p)$

How many flips to get at least one head?

If $p = \frac{1}{2}$, the expected number of flips to get one head equals:

$$\mathbf{E}[X] = 1/p = 1/\frac{1}{2} = 2$$

Key Topics



- Asymptotic notation (big-O, etc.)
- Simple recurrences
- Asymptotic analysis
- Basic probability (e.g., CS1231 review)

Three Parts

1) Basic Theory / Algorithm Analysis

2) Major Algorithms

3) Overarching Principles

Key Algorithms



- Binary Search
- Sorting
- Balanced binary tree (AVL tree)
- Hash table (with chaining, linear probing)

Searching

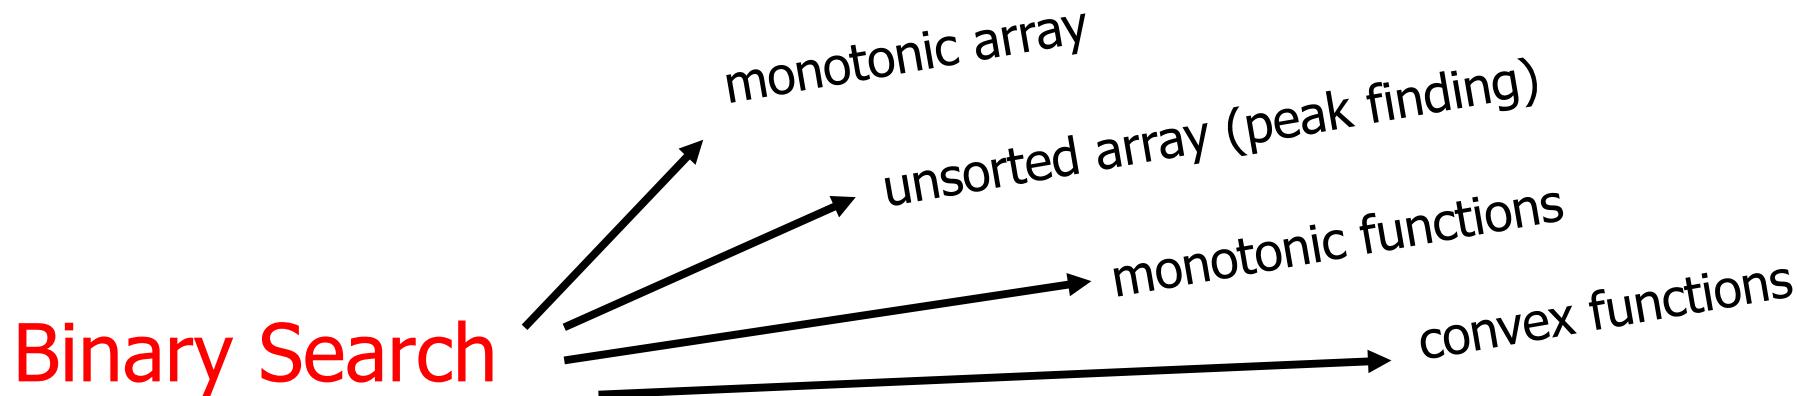
Binary Search

Reduce-and-conquer: make the problem smaller at every step.

At every step, move toward your target.

Key question: how far to move?

Searching

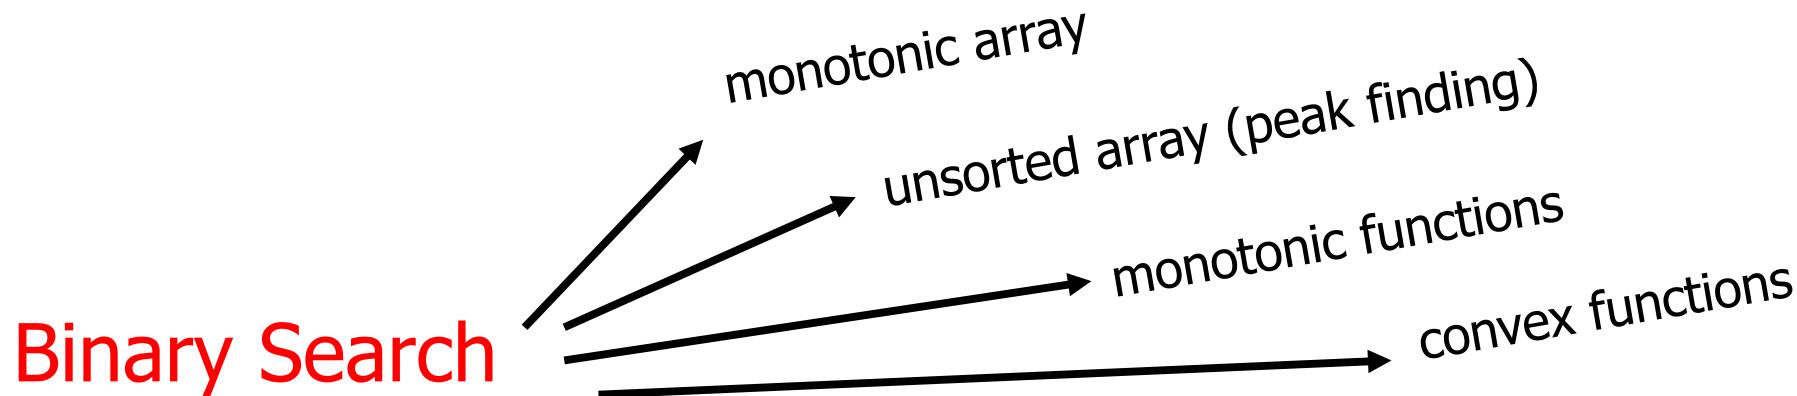


Reduce-and-conquer: make the problem smaller at every step.

At every step, move toward your target.

Key question: how far to move?

Searching



Reduce-and-conquer: make the problem smaller at every step.

At every step, move toward your target.

Key question: how far to move?

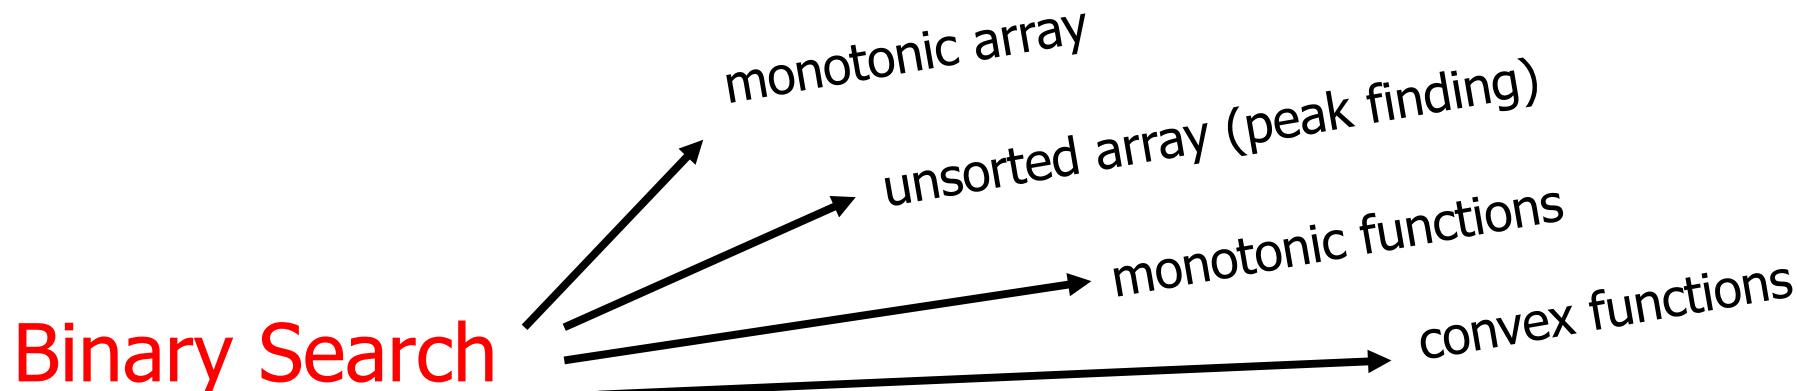
Peak Finding

QuickSelect

Newton's
Method

Gradient
Descent

Searching



Reduce-and-conquer: make the problem smaller at every step.

At every step, move toward your target.

Key question: how far to move?

Peak Finding

Eliminate $\frac{1}{2}$ or $\frac{3}{4}$ of the matrix in every step.

Use matrix bounds to choose how much to eliminate.

QuickSelect

Eliminate approximately $\frac{1}{2}$ the array in each iteration

Use matrix bounds to choose how much to eliminate.

Newton's Method

Always move downhill.

Use first and second derivative to figure out how far to move.

Gradient Descent

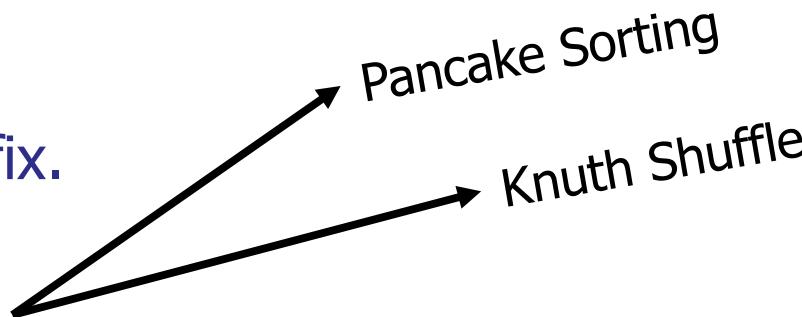
Always move downhill.

Use first derivative to figure out how far to move.

Sorting

BubbleSort

Look for inversions to fix.



SelectionSort

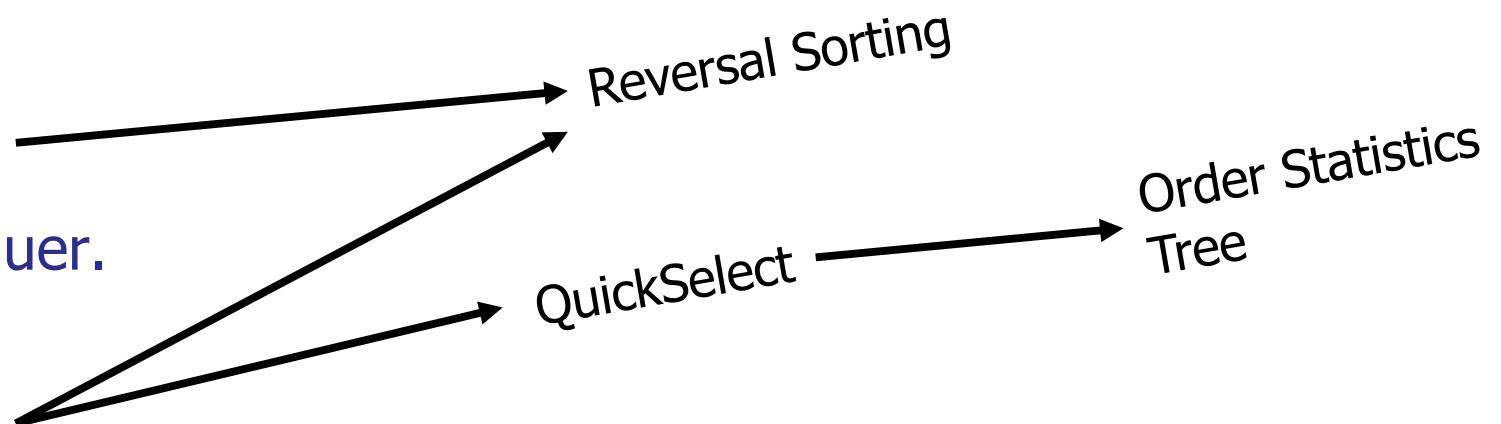
Find the next item to put in place.

InsertionSort

Sort the prefix.

MergeSort

Divide-and-Conquer.



QuickSort

Divide-and-Conquer.

Trees

Binary trees

Unbalanced, in-order, dynamic, traversable

AVL trees

Like binary trees, but height-balanced.

(a,b)-trees

Degree [a,b], balanced, uniform height

SkipLists

Balanced, randomized, list-based

Tries

Unbalanced, string-based

Trees

Binary trees

Unbalanced, in-order, dynamic, traversable

AVL trees

Like binary trees, but height-balanced.



Key idea:
Maintain height
balance using
rotations

(a,b)-trees

Degree [a,b], balanced, uniform height



Key idea:
All leaves
have same
height.

SkipLists

Balanced, randomized, list-based



Key idea:
Random height
makes it balanced.

Tries

Unbalanced, string-based

Trees

Binary trees

Unbalanced, in-order, dynamic, traversable

AVL trees

Like binary trees, but height-balanced.

(a,b)-trees

Degree [a,b], balanced, uniform height

SkipLists

Balanced, randomized, list-based

Tries

Unbalanced, string-based

Design decisions:

Data stored at internal nodes
or only at the leaves?

Which balance method?

- Height-balance
- Weight-balance
- Uniform-height
- Random-height

Trees

Summarizing data:

- Range trees
- Interval trees
- Order statistics trees
- Merkle trees
- Etc...

Hash Tables

Chaining

Resolve collisions using linked lists

Open addressing

Resolve collisions by finding a new slot

Finger print hash tables

Use a hash to identify an item

Bloom filters

Better finger print hash tables

Design decisions:

How do you resolve
collision?

How big is your table?

Do you have false positives?

Key Algorithms



- Binary Search
- Sorting
- Balanced binary tree (AVL tree)
- Hash table (with chaining, linear probing)

Three Parts

1) Basic Theory / Algorithm Analysis

2) Major Algorithms

3) Overarching Principles

Key Ideas

- Basic strategies
- Invariants
- Trade-offs
- Augmenting data structures



Basic strategies for solving problems

Reduce-and-conquer

E.g., binary search

Move in the right direction

E.g., gradient descent

Divide-and-conquer

E.g., MergeSort

Maintaining an invariant

E.g., AVL trees

E.g., keep your data sorted

Augment an existing data structure

E.g., AVL trees

Key Ideas

- Basic strategies
- Invariants
- Trade-offs
- Augmenting data structures



Two important questions:

How do you understand what an algorithm is doing?

How do you show that an algorithm is correct?

Invariants

What is an invariant?

A property or state-relation that is always true.

What is a loop invariant?

A property or state-relation that is true at the beginning or end of every iteration of a loop.

Invariants

Examples:

BinarySearch:

- At the end of iteration j , the active search region between begin and end is size at most $n/2^j$, and the item (if it exists) is in that region.

BubbleSort:

- At the end of iteration j , the last j slots in the array contain the largest j elements in the array in sorted order.

InsertionSort:

- At the end of iteration j , the first $j+1$ slots in the array are sorted.

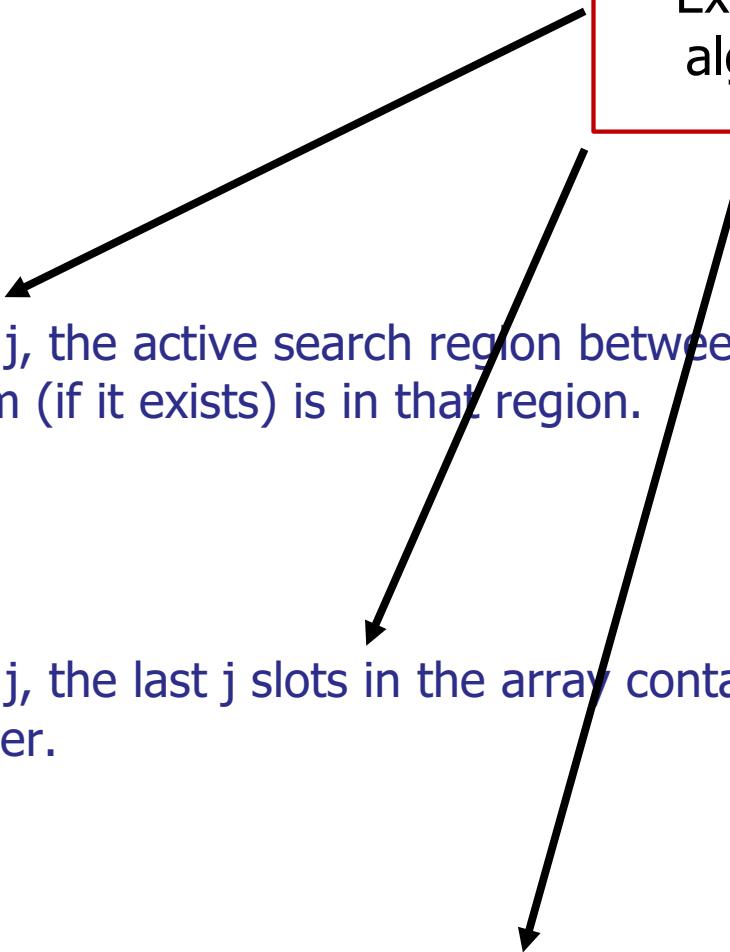
Invariants

Examples:

BinarySearch:

- At the end of iteration j , the active search region between `begin` and `end` is size at most $n/2^j$, and the item (if it exists) is in that region.

Explains how the algorithm works



BubbleSort:

- At the end of iteration j , the last j slots in the array contain the largest j elements in the array in sorted order.

InsertionSort:

- At the end of iteration j , the first $j+1$ slots in the array are sorted.

Invariants

Examples:

BinarySearch: → after $\log(n)$ iterations, search is done and finds items (if it exists).

- At the end of iteration j , the active search region between begin and end is size at most $n/2^j$, and the item (if it exists) is in that region.

BubbleSort: → after n iterations, whole array is sorted.

- At the end of iteration j , the last j slots in the array contain the largest j elements in the array in sorted order.

InsertionSort: → after n iterations, whole array is sorted.

- At the end of iteration j , the first $j+1$ slots in the array are sorted.

Proves that the algorithms works.

Invariants

More examples:

AVL trees:

- After every operation, every node is height balanced.
- The nodes satisfy the BST-order property.

(a,b)-trees:

- After every operation, every leaf has same depth.
- Every node has degree in the range [a,b].
- Nodes satisfy tree-order property.

Skip list:

- Every key is in list 0.
- Lists are in sorted order.

Invariants

More examples:

AVL trees: → tree is height $O(\log n)$ and search works.

- After every operation, every node is height balanced.
- The nodes satisfy the BST-order property.

(a,b)-trees: → tree is height $O(\log_a n)$ and search works.

- After every operation, every leaf has same depth.
- Every node has degree in the range $[a,b]$.
- Nodes satisfy tree-order property.

Skip list: → search works.

- Every key is in list 0.
- Lists are in sorted order.

To understand algorithms:

For every algorithm in the class, identify all of its important invariants.

Key Ideas

- Invariants
- Trade-offs
- Augmenting data structures



Trade-offs

Which data structure should you use?

Unsorted data:

- Storing and searching?
- Inserting and deleting?
- Handling order statistics queries?

Sorted data:

- Searching?
- Inserting and deleting?
- Enumerating items in order?
- Finding the median element?

Trade-offs

What are the strengths and weaknesses of each?

Linked lists:

- Fast insertion

Arrays:

- Efficient scanning
- Binary search
- Space efficient

Balanced Trees:

- Dynamic data structures
- Data summarization
- Predecessor / successor / enumeration

(a,b)-trees:

- All the advantages of balanced trees
- And cache-efficiency.

Hash tables:

- Fast insertion
- Fast lookup
- Space efficient

Fingerprint / Bloom filters:

- Very space efficient
- Fast lookup
- False positives

Simple Problem: Searching

Given:
a collection of **n** items

Find:
the largest item in the collection

Algorithm 1.

1. Sort the collection.
2. Return the largest item.

Simple Problem: Searching

Given:
a collection of n items

Find:
the largest item in the collection

Algorithm 1.

1. Sort the collection.
2. Return the largest item.

Algorithm 2.

1. Iterate through the collection.
2. Return the largest item.

Simple Problem: Searching

Given:
a collection of n items

Find:
the largest item in the collection

Algorithm 1.

1. Sort the collection.
2. Return the largest item.

$O(n \log n)$

Algorithm 2.

1. Iterate through the collection.
2. Return the largest item.

Simple Problem: Searching

Given:
a collection of n items

Find:
the largest item in the collection

Algorithm 1.

1. Sort the collection.
2. Return the largest item.

$O(n \log n)$

Algorithm 2.

1. Iterate through the collection.
2. Return the largest item.

$O(n)$

Simple Problem: Searching

Given:
a collection of n items

Find:
the largest k items
in the collection

Algorithm 1.

1. Sort the collection.
2. Return the largest k items.

Algorithm 2.

1. Iterate through the collection. Store the largest k items seen.
2. Return the largest k items.

Simple Problem: Searching

Extra memory: 1

Time: $O(n \log n + k)$

Algorithm 1.

1. Sort the collection.
2. Return the largest k items.

Extra memory: k

Time: $O(nk)$

Algorithm 2.

1. Iterate through the collection. Store the largest k items seen.
2. Return the largest k items.

Simple Problem: Searching

Extra memory: 1

Time: $O(n \log n + k)$

Algorithm 1.

1. Sort the collection.
2. Return the largest k items.

Extra memory: k

Time: $O(nk)$

Algorithm 2.

1. Iterate through the collection. Store the largest k items seen.
2. Return the largest k items.

Better solution: trees.

When to use a Symbol Table

Example 1: Pilot Scheduling

1. Check to see if feasible to schedule at time t .

No two airplanes can land with 3 minutes of each other.

2. Find schedule of pilot p .

Get a list of all the planes that are being flown by a specified pilot.

Which can be efficiently solved with a symbol table?

1. Both: scheduling and pilots info.
2. Only scheduling.
3. Only pilot info.
4. Neither.

Key Ideas

- Invariants
- Trade-offs
- Augmenting data structures



How to solve new problems?

- 1) Start from scratch and design a new data structure.
- 2) Augment an existing data structure.

How to solve new problems?

- 1) Start from scratch and design a new data structure.
- 2) Augment an existing data structure.



much more common

Augmenting data structures

What do you get when you augment a ... ?

Array:

- Key-value pairs.
- Multidimensional data.
- Cross-linked data structures

Tree:

- Techniques for summarizing data
- Dynamic data structures
- Databases

Order statistics

Interval trees

Range trees

Merkle trees

Trie and Hash Tables:

- All sorts of interesting things... ☺

Three Parts

1) Basic Theory / Algorithm Analysis

2) Major Algorithms

3) Overarching Principles

A bit of advice...

Quiz Advice:

A bit of advice...

Quiz Advice:

Get the maximum number of points you can.

A bit of advice...

Quiz Advice:

Get the maximum number of points you can.

- Do not leave easy questions blank.
- Give a simple solution for partial credit.
- Bypass questions instead of getting stuck.
- Show me what you know.

A bit of advice...

Quiz Advice:

Be as clear as possible.

- Do not be ambiguous.
- Circle your final answer (if it is unclear).
- Cross out incorrect answers.
- Write neatly.

A bit of advice...

Quiz Advice:

State your assumptions.

- If the question is ambiguous, state precisely what you are assuming.
- If your assumptions are reasonable, and your answer is correct subject to those assumptions, you will (most likely) get full credit!

A bit of advice...

Quiz Advice:

Review the basics:

- Know the basic recurrences.
- Review how the algorithms we have studied work.
- Know the running time of the algorithms we have studied.

A bit of advice...

Quiz Advice:

Review problem solving strategies:

- Review problems we have solved on problem sets, in tutorials, in recitation, in class.
- What is the basic strategy used in the question?
 - Binary search? Divide-and-conquer? Sorting?
- What strategy is good for which types of problems?

Three Parts

1) Basic Theory / Algorithm Analysis

2) Major Algorithms

3) Overarching Principles